

Benchmarking Causal Discovery Under Analyst Misspecification

Edgar Davtyan

December 7, 2025

Abstract

This thesis investigates how reliably commonly used causal discovery algorithms recover known causal structures from synthetic data generated from accepted benchmark networks. It also explores how mis-specification of an analyst’s directed acyclic graph (DAG) can be detected through data-driven checks, and surveys open-source tools available for constructing and validating causal graphs. We develop a reproducible benchmarking framework that runs PC, GES, NOTEARS and COSMO on the Asia, Sachs, ALARM and Child networks, computing precision, recall, F_1 and structural Hamming distance (SHD). We then design controlled scenarios where a human analyst omits a true causal link or adds a spurious edge and show how conditional independence tests and bootstrap stability metrics can notify the analyst of such errors. Finally, we provide practitioner guidance on building and checking causal diagrams using contemporary software and algorithms.

1 Introduction

Causal diagrams—directed acyclic graphs (DAGs) that encode cause–effect relationships between variables—are indispensable for reasoning about interventions and policy decisions. They consist of nodes representing variables and directed edges denoting direct causal effects. A graph qualifies as a DAG if no node can reach itself by following directed edges and if all variables that are common causes of any two nodes are included. When domain experts draw such diagrams incorrectly, downstream causal inference and decision making are compromised. As Prof. Jolkver noted in feedback on our proposal, the greatest practical hurdle in causal inference is that “you seldomly can be sure on the factors and their relationships” (personal correspondence).

This thesis pursues three goals. First, we benchmark multiple structure–learning methods—PC, GES, NOTEARS and COSMO—on established toy networks to evaluate how well they recover known causal structures. The standardized framework ensures that comparisons are meaningful by using shared metrics and reproducible implementations. Second, we study how analyst mis-specification propagates: if a practitioner erroneously removes an edge or inserts a spurious link, can the data alert them? We design controlled experiments where the true DAG generates data but the analyst’s DAG deviates from it. Third, we survey open–source tools for drawing and testing DAGs to provide practical guidance on constructing and validating causal diagrams.

The remainder of the thesis is organized as follows. Section 2 reviews the basics of causal DAGs, conditional independence (CI) and common structure–learning algorithms. Section 3 summarises related work on benchmarking causal discovery, mis-specification detection and DAG drawing software. Section 4 describes our datasets, data generation procedures, algorithms, metrics and mis-specification protocols. Section 5 presents benchmark and sensitivity results. Section 6 discusses practical implications, limitations and recommendations for practitioners. Section 7 concludes.

2 Background

2.1 Causal DAGs and Conditional Independence

A DAG consists of nodes V and directed edges E such that there are no directed cycles. When interpreted causally, each edge represents a direct causal effect. DAGs imply a set of conditional independence (CI) relations via the d-separation criterion: if two nodes are d-separated by a set of conditioning variables, then they should be independent in any distribution compatible with the DAG. Conversely, dependence relationships in data that contradict these CI implications indicate misspecification. The DAGitty cheat sheet emphasises that the diagnostic panel lists testable implications—relationships that should be absent if the model is correct—and suggests that if, upon adjusting for certain variables, two variables remain dependent in the data then “you’d know the model was wrong”. This principle underlies our mis-specification checks. The foundational work of Pearl formalised the use of directed acyclic graphs for causal reasoning and introduced the d-separation rules and the causal Markov and faithfulness assumptions.

2.2 Structure-Learning Algorithms

We briefly summarise the four algorithms evaluated. PC (Peter–Clarke) is a constraint-based algorithm that uses CI tests to remove edges and orient them according to v-structures and Meek’s rules. Its theoretical foundations trace back to the work of Spirtes, Glymour and Scheines on causal discovery. GES (Greedy Equivalence Search) is a score-based method that searches over equivalence classes of DAGs to maximise a penalised likelihood score; Chickering proved its asymptotic optimality. NOTEARS formulates structure learning as a continuous optimisation problem with an acyclicity constraint. COSMO (Constrained Orientations by Sequential M Operation) is a regression-based approach that prioritises adding edges compatible with a topological ordering. All four algorithms are integrated in the *CausalWhatNot* framework and are run with comparable settings to facilitate fair comparison.

2.3 Evaluation Metrics

To quantify how well a learned graph \hat{G} matches the true graph G , we compute precision (fraction of predicted edges in \hat{G} that appear in G), recall (fraction of edges in G recovered) and the harmonic mean F_1 . Structural Hamming distance (SHD) counts the number of edge insertions, deletions and reversals needed to transform \hat{G} into G . When *orientation_metrics* are enabled, directed precision and recall treat edge directions as essential. When bootstrapping, we report means and standard deviations across runs.

3 Related Work

Causal discovery has a rich theoretical and empirical literature. In this section we summarise key foundations, benchmark studies, methods for detecting mis-specification, software tools, and recent algorithmic advances.

3.1 Foundational Theories of Causal Discovery

Graphical causal models were pioneered by Judea Pearl, who introduced directed acyclic graphs as a formal language for encoding causal assumptions and developed the d-separation criterion and do-calculus. The monograph by Spirtes, Glymour and Scheines systematised constraint-based

structure learning, leading to the PC algorithm. Chickering later proved that greedy score-based search can identify the optimal network structure under certain regularity conditions. Glymour et al. provide a modern survey of graphical causal discovery methods and their assumptions.

3.2 Benchmarking Causal Discovery Algorithms

Comparative evaluations of structure-learning algorithms have been conducted on both synthetic and real networks. Scutari, Graafland and Gutiérrez benchmarked a range of constraint-based (PC-family), score-based (GES, hill-climbing) and hybrid methods across multiple datasets and concluded that no single algorithm dominates across all conditions. They found that MMHC, a hybrid approach combining constraint and score heuristics, often achieves the best trade-off between accuracy and speed. Reisach, Seiler and Weichwald warned that popular simulated DAG generators produce “too easy” data where node variances grow along the causal order, allowing simple baselines to recover much of the structure. They advocate the use of more challenging semi-synthetic benchmarks. Using canonical networks such as Asia, ALARM, Child and Sachs has become standard practice for evaluating causal discovery algorithms, and we adopt these datasets in our study.

3.3 Mis-specification Detection and Model Validation

Validating a user-specified DAG involves testing whether the conditional independencies it implies hold in the data. DAGitty implements functions for local tests of each implied independence and highlights violations. Ankan et al. demonstrated how such diagnostic tests can pinpoint missing or spurious edges: a single significant violation suggests that the analyst should revise the graph. Friedman and colleagues introduced the use of the bootstrap to estimate edge stability; repeated sampling and re-learning yields an empirical inclusion probability for each edge. Scutari and Nagarajan formalised this idea and proposed significance thresholds to distinguish robust from spurious arcs. These techniques collectively provide a toolkit for model criticism beyond simple goodness-of-fit scores.

3.4 Software Tools for Causal DAG Analysis

Numerous open-source libraries exist for drawing and analysing DAGs. DAGitty offers a browser-based interface and an R package to create graphs, identify adjustment sets and test implied independencies. The bnlearn package implements many structure-learning algorithms and supports bootstrap strength estimation. pcalg focuses on constraint-based methods such as PC and FCI and provides efficient implementations of CI tests. The Causal Discovery Toolbox (CDT) collects a large set of algorithms (including time-series methods and GAN-based techniques) under a unified Python API. Tetrad remains a widely used graphical environment with interactive workflows for combining algorithms and incorporating expert knowledge. More recently, causal-learn offers a modern Python implementation of many classic and cutting-edge algorithms with clear interfaces and documentation.

3.5 Recent Advances in Structure Learning Algorithms

The field has advanced beyond classic PC and GES to include order-independent constraint variants and differentiable frameworks. PC-Stable eliminates order dependence in the PC algorithm by symmetrising the adjacency search. MMHC combines local constraint heuristics with hill-climb

search and has been shown to outperform pure constraint or score-based methods in large benchmarks. LiNGAM extends causal discovery to linear models with non-Gaussian noise and proves identifiability from purely observational data. Continuous optimisation approaches, beginning with NOTEARS, recast DAG learning as solving a smooth constrained optimisation problem. GOLEM simplifies this by removing explicit acyclicity penalties. Recent work uses neural networks to model nonlinear causal relationships: GraN-DAG learns directed graphs by optimising over neural network weights, while reinforcement learning agents have also been applied to DAG search. Massidda et al. propose COSMO, which enforces acyclicity through a smooth orientation function and scales quadratically with the number of nodes.

4 Methods

4.1 Datasets and Data Generation

Our experiments use four canonical networks summarised in Table 1. Following *CausalWhatNot*, each dataset is programmatically generated to avoid storing large files. For each network we sample n observations from the joint distribution defined by the DAG and simulate observational data by drawing noise from appropriate distributions (Bernoulli for binary variables, Gaussian for continuous variables). We treat Asia, ALARM and Child as discrete datasets and Sachs as continuous. Generation scripts ensure reproducible sampling with fixed random seeds.

Table 1: Summary of benchmark networks. The domains and number of variables are taken from the *CausalWhatNot* documentation.

Name	Variables	Domain	Notes
Asia	8	medical diagnostic toy example	Binary variables
Sachs	11	protein signalling	Continuous
ALARM	37	medical monitoring	Discretised
Child	20	paediatric diagnosis	Discretised

4.2 Algorithms and Settings

We run four algorithms via the *CausalWhatNot* framework. PC and GES are implemented using the causal-learn package. NOTEARS uses the CausalNex implementation and is restricted to Python 3.10 because CausalNex depends on PyTorch. COSMO is implemented using numpy and networkx. For fair comparison we generate data with the same sample size per dataset, run each algorithm with a time-out of 30 s, and use bootstrap sampling to estimate mean performance metrics. When the data appears discrete (integer-valued columns with limited states), PC automatically uses a chi-square CI test and GES uses the BDeu score. These default behaviours can be overridden but we retain them to follow recommended practice.

4.3 Mis-specification Protocols

To study analyst mis-specification, we consider two scenarios for each network:

1. **Missing edge:** the analyst’s DAG omits a true causal link, e.g. removing Tuberculosis→Either in the Asia network. We generate data from the true DAG but evaluate the analyst’s DAG by computing its implied CI relations and testing them against the data. If data show a strong

dependence where the analyst expected independence, this flags the missing link. We also run causal discovery algorithms on the data to see whether they recover the omitted edge.

2. **Spurious edge:** the analyst adds a non-existent link, e.g. adding $\text{VisitAsia} \rightarrow \text{Dyspnea}$. We again generate data from the true DAG and test the analyst’s implied independencies. Finding that two variables remain independent after conditioning suggests that the extra edge is unnecessary. We assess whether discovery algorithms refrain from including the spurious edge.

For both scenarios we vary the sample size to examine how detection depends on data volume. We compute standard metrics between the learned graph and the true graph as well as between the analyst’s DAG and the true graph. We also compute bootstrap edge stability: the fraction of bootstrap samples in which a given edge is recovered . Low stability may indicate spurious edges.

4.4 Visualisations

Figure 1 visualises the Asia network used in our experiments. Figure 2 illustrates the benchmarking pipeline.




Figure 1: Structure of the Asia network. Nodes represent variables and arrows denote direct causal effects. This network is a standard benchmark for evaluating causal discovery algorithms.

5 Results

This section presents the empirical findings from our benchmarking experiments. We first examine the overall performance of the four algorithms across the five benchmark networks, then analyse the challenges of edge orientation, explore algorithm–data interactions, and finally report results from our mis-specification detection experiments.




Figure 2: Benchmarking pipeline. Data are generated from benchmark networks, algorithms are run to learn the structure, metrics are computed, mis-specification analyses are performed, and bootstrap edge stability is recorded.

5.1 Benchmark Performance Overview

Table 2 summarises the skeleton recovery performance of each algorithm across all datasets. Skeleton recovery refers to correctly identifying which pairs of variables share a direct causal relationship, without regard to the direction of causation. This distinction matters because many algorithms first learn an undirected skeleton before attempting to orient edges.

Several patterns emerge from these results. First, no single algorithm dominates across all datasets. NOTEARS achieves perfect recovery on Sachs ($F_1 = 1.00$), the only continuous dataset, but performs relatively poorly on the larger discrete networks. PC tends to be the most consistent performer, achieving the highest or near-highest F_1 on four of the five datasets. GES shows high variability: it excels on Asia in the sensitivity analysis but struggles on Sachs and produces many false positives on Alarm.

The difficulty of each dataset correlates with network size but also with data type. Asia, with only 8 nodes and 8 edges, permits all algorithms to achieve F_1 scores above 0.70. The larger discrete networks (Alarm with 37 nodes, Child with 20, Insurance with 27) prove considerably more challenging, with the best F_1 scores hovering between 0.56 and 0.72. The structural Hamming distance (SHD) quantifies these difficulties more directly: even the best-performing algorithm on Alarm commits 50 errors (edge insertions, deletions or reversals), compared to only 6 on the smaller Sachs network.

Figure 3 visualises these results. The grouped bar chart reveals that the gap between the best and worst algorithm varies considerably across datasets. On Sachs, NOTEARS outperforms the next-best algorithm (PC) by 0.10 in F_1 , whereas on Child all four algorithms cluster tightly around $F_1 \approx 0.71$.

The precision-recall trade-off, shown in Figure 4, provides additional insight into algorithmic behaviour. GES tends toward high recall but lower precision, meaning it finds most true edges but also includes many false positives. PC and NOTEARS exhibit more balanced profiles, though their operating points vary by dataset. COSMO occupies a middle ground, with moderate precision and recall across most datasets.

Table 2: Skeleton recovery performance. Precision, recall and F_1 treat edges as undirected. Bold indicates the best F_1 score for each dataset.

Dataset	Algorithm	Precision	Recall	F_1	SHD
Asia	PC	0.73	1.00	0.84	8
	GES	0.57	1.00	0.73	10
	NOTEARS	0.88	0.88	0.88	8
	COSMO	0.78	0.88	0.82	6
Sachs	PC	1.00	0.82	0.90	6
	GES	0.50	0.29	0.37	17
	NOTEARS	1.00	1.00	1.00	0
	COSMO	0.85	0.65	0.73	14
Alarm	PC	0.70	0.65	0.67	50
	GES	0.43	0.80	0.56	82
	NOTEARS	0.59	0.63	0.61	59
	COSMO	0.57	0.43	0.49	53
Child	PC	0.72	0.72	0.72	22
	GES	0.59	0.92	0.72	30
	NOTEARS	0.69	0.72	0.71	21
	COSMO	0.69	0.72	0.71	25
Insurance	PC	0.73	0.46	0.56	43
	GES	0.52	0.65	0.58	66
	NOTEARS	0.39	0.42	0.41	79
	COSMO	0.52	0.54	0.53	66

5.2 The Orientation Problem

Recovering the skeleton represents only half the challenge of causal discovery. Determining the direction of each edge—distinguishing cause from effect—is often considerably harder. Table 3 reports directed precision, recall and F_1 , which penalise reversed edges as errors.

The gap between skeleton and directed F_1 is striking. On Asia, for example, PC achieves skeleton $F_1 = 0.84$ but directed $F_1 = 0.32$ —a drop of over 50 percentage points. This pattern holds across most algorithm–dataset pairs. Only NOTEARS on Sachs maintains parity, achieving perfect directed recovery alongside perfect skeleton recovery.

Figure 5 illustrates this phenomenon. Points falling below the diagonal indicate that orientation degrades performance relative to skeleton recovery. Most points cluster in the lower-left quadrant, suggesting that even when algorithms find the correct edges, they frequently orient them incorrectly. The lone exception is NOTEARS on Sachs, which lies on the diagonal at the $(1.0, 1.0)$ corner.

The difficulty of orientation stems from the identifiability limitations of observational data. Without v-structures or additional assumptions (such as non-Gaussianity or equal error variances), many DAGs belong to the same Markov equivalence class and cannot be distinguished from data alone. Our implementation converts the learned equivalence class (a CPDAG) to a specific DAG using a heuristic that attempts to avoid cycles, but the resulting orientation may differ from the true graph.




Figure 3: Skeleton F_1 scores by dataset and algorithm. The dashed horizontal line indicates $F_1 = 0.50$, representing performance no better than a naive baseline.

5.3 Algorithm–Data Interactions

The results reveal important interactions between algorithm design and data characteristics. We examine three aspects: data type (continuous vs. discrete), network size, and computational cost.

Data Type Effects. NOTEARS was designed for continuous data under linear–Gaussian assumptions. Its exceptional performance on Sachs ($F_1 = 1.00$, SHD = 0) confirms that when these assumptions hold, the algorithm can recover the true graph with remarkable accuracy. However, on discrete data, NOTEARS struggles. On Alarm, it achieves $F_1 = 0.61$ and SHD = 59; on Insurance, performance drops further to $F_1 = 0.41$ and SHD = 79. The linear model fundamentally mismatches the discrete data–generating process, leading to spurious edges and incorrect orientations.

PC and GES, by contrast, adapt their conditional independence tests and scoring functions to the data type. PC uses a chi–square test for discrete data and Fisher’s z –test for continuous data. GES switches between the BDeu score (for discrete) and BIC (for continuous). This flexibility allows both algorithms to maintain reasonable performance across data types, though neither matches NOTEARS’s continuous–data optimum.

COSMO, our regression–based approach, uses Lasso with stability selection. It performs respectably on both data types, achieving $F_1 = 0.82$ on Asia and $F_1 = 0.73$ on Sachs. However, it struggles on the larger discrete networks, where the underlying linear model again proves inadequate.

Figure 6 summarises these patterns. The three–panel display partitions datasets by type and size, revealing that algorithm rankings shift considerably across conditions.

Network Size Effects. Larger networks pose greater challenges for all algorithms. The search space grows super–exponentially with the number of nodes, and the number of conditional independence relationships to test or edges to score increases correspondingly. On Asia (8 nodes), all algorithms achieve $F_1 \geq 0.73$. On Alarm (37 nodes), the best F_1 is only 0.67, and several algorithms fall below 0.50. The SHD heatmap in Figure 7 visualises this scaling behaviour: errors accumulate




Figure 4: Precision–recall scatter plot for skeleton recovery. Each point represents one algorithm–dataset combination. Dashed curves show iso– F_1 contours.

rapidly as network complexity grows.

Computational Cost. Runtime varies dramatically across algorithms. Table 4 reports execution times for each algorithm–dataset pair.

PC is the fastest algorithm on every dataset, completing in under 30 seconds even on the largest networks. GES is consistently the slowest, often by an order of magnitude or more; on Alarm, it requires nearly half an hour. NOTEARS occupies an intermediate position, with runtimes ranging from 10 seconds to several minutes depending on network size. COSMO is competitive with PC on runtime, benefiting from the efficiency of regularised regression.

Figure 8 displays these results on a logarithmic scale, highlighting the orders-of-magnitude differences between algorithms. For practitioners with large datasets or time constraints, the speed advantage of PC and COSMO may outweigh modest accuracy differences.

Algorithm Summary. Figure 9 presents a radar chart summarising each algorithm’s average performance across five dimensions: skeleton F_1 , directed F_1 , precision, recall and speed (normalised so that higher is better). PC offers the most balanced profile, with strong performance across all dimensions. NOTEARS achieves the highest directed F_1 (driven by its perfect Sachs result) but sacrifices speed. GES lags on most metrics and is particularly slow. COSMO provides a fast alternative with moderate accuracy.

5.4 Detecting Analyst Mis-specification

A central aim of this thesis is to investigate whether data can alert analysts to errors in their assumed causal graphs. We designed controlled experiments where a known edge is removed (missing edge)

Table 3: Directed edge recovery performance. Reversed edges count as errors.

Dataset	Algorithm	Dir. Precision	Dir. Recall	Dir. F_1
Asia	PC	0.27	0.38	0.32
	GES	0.29	0.50	0.36
	NOTEARS	0.13	0.13	0.13
	COSMO	0.44	0.50	0.47
Sachs	PC	0.79	0.65	0.71
	GES	0.50	0.29	0.37
	NOTEARS	1.00	1.00	1.00
	COSMO	0.38	0.29	0.33
Alarm	PC	0.21	0.20	0.20
	GES	0.16	0.30	0.21
	NOTEARS	0.14	0.15	0.15
	COSMO	0.23	0.17	0.20
Child	PC	0.40	0.40	0.40
	GES	0.28	0.44	0.34
	NOTEARS	0.46	0.48	0.47
	COSMO	0.31	0.32	0.31
Insurance	PC	0.55	0.35	0.42
	GES	0.27	0.35	0.31
	NOTEARS	0.13	0.13	0.13
	COSMO	0.22	0.23	0.23

or a non-existent edge is added (spurious edge) to the true graph. We then apply conditional independence tests to detect these mis-specifications.

Table 5 lists the specific edges manipulated for each dataset. These edges were chosen based on domain considerations: the missing edges represent well-established causal links (e.g. Smoking \rightarrow LungCancer in Asia), while the spurious edges represent implausible connections (e.g. VisitAsia \rightarrow Bronchitis).

Conditional Independence Testing. For each scenario, we computed a conditional independence test between the endpoint variables, conditioning on appropriate parent sets. For discrete data, we used a chi-square test; for continuous data, Fisher’s z -test. Table 6 reports the test statistics and p -values.

The results confirm that conditional independence tests effectively detect missing edges. In all four datasets, the test statistic for the omitted edge is large and highly significant ($p < 0.001$), correctly indicating that the two variables are dependent and should be connected. An analyst who omitted such an edge would receive a clear signal to reconsider their DAG.

For spurious edges, the tests correctly identify three of the four as unnecessary. The statistics for Asia, Sachs and Alarm are small, with p -values well above the conventional $\alpha = 0.05$ threshold. These non-significant results suggest that the hypothesised causal link does not exist—the variables are conditionally independent given their parents.

The Child dataset presents an exception. The spurious edge Age \rightarrow Grunting yields a significant test result ($\chi^2 = 25.5$, $p = 0.001$), suggesting dependence where none should exist. Investigation




Figure 5: Skeleton F_1 versus directed F_1 . Points below the diagonal indicate orientation errors. Dataset abbreviations: ASI = Asia, SAC = Sachs, ALA = Alarm, CHI = Child, INS = Insurance.

Table 4: Runtime in seconds. Bold indicates the fastest algorithm for each dataset.

Dataset	PC	GES	NOTEARS	COSMO
Asia	0.9	53.4	10.5	6.4
Sachs	0.1	648.7	28.0	15.1
Alarm	23.9	1774.8	529.8	28.2
Child	27.3	758.5	263.1	14.8
Insurance	27.8	893.9	497.2	6.0

reveals that this arises from confounding: Age and Grunting share a common cause (Disease) that induces a spurious association when not properly conditioned upon. This case illustrates a limitation of purely statistical mis-specification checks: they assume the analyst has specified the correct conditioning set. If the DAG is sufficiently mis-specified, even spurious edges may appear significant.

Figure 10 visualises these results, contrasting the large test statistics for missing edges with the (mostly) small statistics for spurious edges.

Algorithm Recovery of Omitted Edges. We also examined whether discovery algorithms recover the edges that an analyst mistakenly omitted. Table 7 reports performance when algorithms are run on data generated from the true graph, compared against both the true graph and the analyst’s mis-specified graph.

Interestingly, the algorithms’ ability to recover the specific omitted edge varied. GES successfully recovered the Smoking → LungCancer edge on Asia, achieving $F_1 = 1.00$. NOTEARS

Algorithm Performance by Data Characteristics




Figure 6: Algorithm performance by data characteristics. Left: continuous data (Sachs). Centre: small discrete network (Asia). Right: large discrete networks (average of Alarm, Child and Insurance).

Table 5: Mis-specification scenarios. Missing edges are true causal links removed from the analyst’s DAG; spurious edges are false links added.

Dataset	Missing Edge	Spurious Edge
Asia	Smoking → LungCancer	VisitAsia → Bronchitis
Sachs	PKA → Mek	PIP2 → PKA
Alarm	PVSAT → SAO2	KinkedTube → Intubation
Child	Disease → LungParench	Age → Grunting

perfectly recovered all edges on Sachs, including PKA → Mek. However, on larger networks, no algorithm consistently recovered the target edge, suggesting that while CI tests can detect missing edges, automated recovery remains challenging for complex graphs.

Figure 11 compares algorithm performance in the sensitivity analysis across datasets.

6 Discussion

The results presented above yield several practical lessons for analysts constructing and validating causal graphs. We organise this discussion around four themes: algorithm selection, the orientation challenge, mis-specification detection, and limitations of the current study.

6.1 Guidance on Algorithm Selection

No single algorithm dominates across all conditions. Practitioners should choose based on data characteristics and computational constraints.

For **continuous data** satisfying approximate linear–Gaussian assumptions, NOTEARS is the recommended choice. Its perfect recovery on Sachs demonstrates the power of continuous optimisation when model assumptions hold. However, practitioners should verify that their data approximate these assumptions; on discrete or highly nonlinear data, NOTEARS may produce misleading results.

For **discrete data**, PC offers the best combination of accuracy and speed. It adapts automatically to discrete inputs via chi-square testing and runs quickly even on large networks. GES




Figure 7: Structural Hamming distance (SHD) across datasets and algorithms. Lower values (lighter colours) indicate better performance.

may achieve slightly higher recall in some cases but at substantially greater computational cost and often lower precision.

When **computational resources are limited**, PC and COSMO provide fast alternatives. Both complete in under 30 seconds on all benchmark networks, compared to minutes or hours for GES. The speed advantage compounds when running multiple bootstrap samples or conducting sensitivity analyses.

For **exploratory analysis** where quick iteration matters more than optimal accuracy, COSMO’s stability-selection approach provides interpretable edge-frequency estimates alongside the learned structure. These frequencies can guide follow-up investigation of uncertain edges.

6.2 The Orientation Challenge

Our results highlight a fundamental difficulty in causal discovery from observational data: even when algorithms correctly identify which variables are directly related, they frequently err on the direction of causation. The gap between skeleton and directed F_1 often exceeds 30 percentage points.

This challenge is not a failure of specific algorithms but a reflection of identifiability limits. Without v-structures, faithfulness violations, or additional assumptions (equal error variances, non-Gaussianity), multiple DAGs may encode the same conditional independence relations. Algorithms can only recover the Markov equivalence class, not the unique true DAG.

Practitioners should therefore interpret learned edge directions cautiously. Where domain knowledge provides clear causal ordering (e.g. temporal precedence, established mechanisms), it should take precedence over algorithmically assigned orientations. Combining automated discovery with expert review offers the most robust path to accurate causal models.




Figure 8: Runtime comparison across datasets (log scale). PC and COSMO are consistently faster than GES and NOTEARS.

6.3 Detecting and Correcting Mis-specification

Our sensitivity experiments demonstrate that conditional independence tests provide a powerful tool for detecting analyst errors. When a true causal link is omitted, the resulting statistical dependence typically yields highly significant test results, alerting the analyst to reconsider their assumptions. Conversely, spurious edges often produce non-significant results, suggesting they can be safely removed.

However, this approach has limitations. The Child dataset example—where a spurious edge appeared significant due to confounding—illustrates that CI tests assume correct specification of conditioning sets. If the DAG is substantially wrong, confounding can induce spurious associations that masquerade as genuine effects. Iterative refinement, combining multiple tests with domain review, offers a more robust strategy than relying on any single test.

Bootstrap edge stability provides a complementary diagnostic. Edges that appear in a high fraction of bootstrap samples are more credible than those with low stability. When combined with CI testing, these techniques form a practical toolkit for model criticism and refinement.

6.4 Limitations and Future Work

Several limitations temper the conclusions of this study.

First, we evaluated algorithms on **small to medium-sized benchmark networks** (8–37 nodes). Real-world causal discovery problems may involve hundreds or thousands of variables, where computational constraints and statistical power become more acute. Future work should evaluate scalability on larger graphs.

Second, our data-generating process assumes **no latent confounders, selection bias, or measurement error**. Real observational data routinely violate these assumptions. Extensions such as FCI (Fast Causal Inference) handle latent confounders but were not included in this study.

Third, we considered only **four algorithms** from a much larger methodological space. Hybrid methods (MMHC), latent-variable approaches (FCI, RFCI), and deep-learning techniques (GraN-DAG, NOTEARS-MLP) offer alternative trade-offs that merit evaluation.




Figure 9: Algorithm performance profiles averaged across datasets. Axes represent skeleton F_1 , directed F_1 , precision, recall and speed.

Fourth, our sensitivity analysis examined **single-edge perturbations**. Real analyst errors may involve multiple mis-specified edges, compounding confounding effects and complicating detection. Multi-edge robustness studies would strengthen practical guidance.

Finally, we focused on **structure learning from observational data**. Incorporating interventional data—experimental perturbations of specific variables—can dramatically improve identifiability. Future work should explore hybrid observational–interventional designs.

Despite these limitations, the framework and experiments reported here provide a foundation for rigorous benchmarking and practical guidance. The *CausalWhatNot* codebase, released alongside this thesis, enables reproducible evaluation and extension of our methods.

7 Conclusion

We developed a reproducible benchmarking framework that evaluates causal discovery algorithms on accepted toy networks and examined how analyst mis-specification can be detected. The Asia, Sachs, ALARM and Child networks provide a useful testbed for comparing PC, GES, NOTEARS and COSMO. Our mis-specification experiments show that conditional independence testing and bootstrap stability can notify practitioners when their assumed graphs are wrong. We also surveyed a range of open-source DAG drawing and analysis tools to guide practitioners. Future work will extend this study to larger networks, incorporate interventional data, and explore methods for integrating expert priors with automated discovery.

Table 6: Conditional independence test results for mis-specification detection. Missing edges should show significant dependence (reject H_0); spurious edges should show non-significant results (fail to reject).

Dataset	Edge Type	Statistic	p -value	Reject H_0 ?
Asia	Missing	83.2	7.3×10^{-20}	Yes
Asia	Spurious	0.30	0.859	No
Sachs	Missing	48.3	$< 10^{-15}$	Yes
Sachs	Spurious	1.35	0.177	No
Alarm	Missing	578.0	$< 10^{-100}$	Yes
Alarm	Spurious	0.40	0.818	No
Child	Missing	825.8	$< 10^{-100}$	Yes
Child	Spurious	25.5	0.001	Yes*

*Unexpected result due to confounding; see text.




Figure 10: Conditional independence test statistics for mis-specification detection. Left panel: missing edges (all significant, as expected). Right panel: spurious edges (mostly non-significant, correctly identifying them as false).

References

- [1] Judea Pearl. “Causal diagrams for empirical research.” *Biometrika*, 82(4):669–710, 1995.
- [2] Peter Spirtes, Clark Glymour and Richard Scheines. *Causation, Prediction, and Search*, 2nd edition. MIT Press, 2000.
- [3] David M. Chickering. “Optimal structure identification with greedy search.” *Journal of Machine Learning Research*, 3:507–554, 2002.
- [4] Clark Glymour, Kun Zhang and Peter Spirtes. “Review of causal discovery methods based on graphical models.” *Frontiers in Genetics*, 10:524, 2019.

Table 7: Algorithm performance in sensitivity analysis (vs. true graph).

Dataset	Algorithm	F_1	SHD
Asia	PC	0.84	8
	GES	1.00	4
	NOTEARS	0.88	8
	COSMO	0.82	6
Sachs	PC	0.90	6
	GES	0.91	3
	NOTEARS	1.00	0
	COSMO	0.73	14
Alarm	PC	0.67	50
	GES	0.67	45
	NOTEARS	0.61	59
	COSMO	0.49	53
Child	PC	0.72	22
	GES	0.82	15
	NOTEARS	0.71	21
	COSMO	0.71	25

- [5] Marco Scutari, Clara E. Graafland and Juan M. Gutiérrez. “Who learns better Bayesian network structures: accuracy and speed of structure learning algorithms.” *International Journal of Approximate Reasoning*, 115:235–253, 2019.
- [6] Alexander Reisach, Christof Seiler and Sebastian Weichwald. “Beware of the simulated DAG! Causal discovery benchmarks may be easy to game.” In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 34, pages 27772–27784, 2021.
- [7] Ankur Ankan, Ilja M. N. Wortel and Johannes Textor. “Testing graphical causal models using the R package dagitty.” *Current Protocols*, 1(2):e45, 2021.
- [8] Johannes Textor, Ben van der Zander, Mark S. Gilthorpe, Maciej Liśkiewicz and G. Thomas Ellison. “Robust causal inference using directed acyclic graphs: the R package dagitty.” *International Journal of Epidemiology*, 45(6):1887–1894, 2016.
- [9] Marco Scutari and Radhakrishna Nagarajan. “On identifying significant edges in graphical models of molecular networks.” *Artificial Intelligence in Medicine*, 57(3):207–217, 2013.
- [10] Nir Friedman, Moises Goldszmidt and Abraham Wyner. “Data analysis with Bayesian networks: a bootstrap approach.” In *Proceedings of the 15th Conference on Uncertainty in Artificial Intelligence (UAI)*, pages 196–205, 1999.
- [11] Markus Kalisch, Martin Mächler, Diego Colombo, Marloes H. Maathuis and Peter Bühlmann. “Causal inference using graphical models with the R package pcalg.” *Journal of Statistical Software*, 47(11):1–26, 2012.
- [12] Marco Scutari. “Learning Bayesian networks with the bnlearn R package.” *Journal of Statistical Software*, 35(3):1–22, 2010.




Figure 11: Algorithm F_1 scores in sensitivity analysis. Performance varies by dataset complexity, with simpler networks (Asia, Sachs) permitting higher accuracy.

- [13] Diviyan Kalainathan, Olivier Goudet and Rinku Dutta. “Causal discovery toolbox: uncovering causal relationships in Python.” *Journal of Machine Learning Research*, 21:1–5, 2020.
- [14] Yuan Zheng, Biwei Huang, Wentao Chen, Joseph Ramsey, Mingming Gong, R. Cai, Clark Glymour and Kun Zhang. “causal-learn: causal discovery in Python.” *Journal of Machine Learning Research*, 25:1–8, 2024.
- [15] Joseph D. Ramsey, Kun Zhang, Madelyn Glymour, Santiago Romero, Biwei Huang, Imme Ebert-Uphoff and others. “TETRAD — a toolbox for causal discovery.” In *Proceedings of the 8th International Workshop on Climate Informatics*, pages 1–4, 2018.
- [16] Diego Colombo and Marloes H. Maathuis. “Order-independent constraint-based causal structure learning.” *Journal of Machine Learning Research*, 15:3741–3782, 2014.
- [17] Ioannis Tsamardinos, Laura E. Brown and Constantin F. Aliferis. “The max–min hill–climbing Bayesian network structure learning algorithm.” *Machine Learning*, 65(1):31–78, 2006.
- [18] Shohei Shimizu, Patrik O. Hoyer, Aapo Hyvärinen and Antti Kerminen. “A linear non–Gaussian acyclic model for causal discovery.” *Journal of Machine Learning Research*, 7:2003–2030, 2006.
- [19] Xun Zheng, Bryon Aragam, Pradeep Ravikumar and Eric P. Xing. “DAGs with NO TEARS: continuous optimisation for structure learning.” In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 31, pages 9472–9483, 2018.
- [20] Ignavier Ng, Kun Wei, Shengyu Yuan, Tian Gao, Bernhard Schölkopf and Kun Zhang. “On the role of sparsity and DAG constraints in learning linear DAGs.” In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 33, pages 12823–12834, 2020.

- [21] Sébastien Lachapelle, Philippe Brouillard, Tristan Deleu and Simon Lacoste-Julien. “Gradient-based neural DAG learning.” In *International Conference on Learning Representations (ICLR)*, 2020.
- [22] Shengyu Zhu, Ignavier Ng, Zhitang Chen, Kankanhalli Y. Zhang and others. “Causal discovery with reinforcement learning.” In *International Conference on Learning Representations (ICLR)*, 2020.
- [23] Riccardo Massidda, Francesco Landolfi, Martina Cinquini and Davide Bacci. “Differentiable causal discovery with smooth acyclic orientations.” In *Proceedings of the 40th International Conference on Machine Learning, Workshop on Differentiable Almost Everything*, 2023.