

Benchmarking Causal Discovery Under Analyst Misspecification

Edgar Davtyan

January 6, 2026

Abstract

This thesis investigates the reliability with which common causal discovery algorithms recover known causal structures from synthetic data generated by established benchmark networks. It further explores detecting mis-specification in an analyst's directed acyclic graph (DAG) through data-driven checks, alongside a survey of open-source tools for constructing and validating causal graphs. To this end, we develop a reproducible benchmarking framework applying PC, GES, NOTEARS, and COSMO to five standard networks—Asia, Sachs, ALARM, Child, and Insurance—evaluating performance via precision, recall, F_1 score, and structural Hamming distance (SHD). In controlled experiments simulating analyst error, where a true causal link is omitted or a spurious edge included, we demonstrate how conditional independence tests and bootstrap stability metrics can alert practitioners to such discrepancies. Finally, the work provides guidance on building and validating causal diagrams using contemporary software and algorithms.

Contents

Abstract	I
List of Figures	VII
List of Tables	IX
List of Abbreviations	X
1 Introduction	1
2 Theoretical Framework	2
2.1 Probabilistic Graphical Models	2
2.2 Constraint-Based Discovery: The PC Algorithm	5
2.3 Score-Based Discovery: GES	6
2.4 Continuous Optimisation: NOTEARS	7
2.5 Regression-Based Discovery: COSMO	8
2.6 Evaluation Metrics	9
2.7 Structural Causal Models and Interventions	10
2.8 Assumptions, Identifiability, and Scope	10
2.9 Conditional Independence Testing	11
2.10 Statistical Tests for Comparing Algorithms	11
3 Related Work	12
3.1 Foundations of Causal Discovery	12
3.1.1 The Graphical Viewpoint	12
3.1.2 The Search and Score Viewpoint	12
3.1.3 A brief historical arc of structure learning paradigms	13
3.1.4 Algorithm families, assumptions, and outputs	13
3.1.5 Positioning of this thesis within the literature	15
3.2 The Benchmarking Crisis in Causal Discovery	15
3.2.1 Standardisation Efforts	15
3.2.2 The "Scale-Free" Trap	15
3.3 Model Criticism and Mis-specification	15
3.3.1 Global vs. Local Fit	16

3.3.2	Refutation and Stability	16
3.4	The Differentiable Discovery Revolution	16
3.4.1	Extensions and Limitations	16
3.5	Software Ecosystem	17
3.6	Positioning of this Thesis	17
4	Methods	17
4.1	Datasets and Data Generation	17
4.1.1	Implemented synthetic generation: linear SEM with quantile discretisation	18
4.1.2	Determinism and dataset reuse	19
4.2	Algorithms and Settings	20
4.2.1	Hyperparameter Selection Protocol	20
4.2.2	Computational Environment	21
4.2.3	Mapping algorithm settings to the implementation	21
4.2.4	Runtime measurement and parallel execution	22
4.2.5	Reproducibility artefacts	22
4.2.6	PC Algorithm	22
4.2.7	Greedy Equivalence Search (GES)	23
4.2.8	NOTEARS	24
4.2.9	COSMO	25
4.3	Benchmarking pipeline and post-processing	25
4.3.1	Detailed Pipeline Walkthrough	25
4.4	Mis-specification Protocols	27
4.5	Visualisations	27
5	Results	28
5.1	Benchmark Performance Overview	28
5.2	Dataset-by-Dataset Analysis	30
5.2.1	Asia (Discrete, 8 nodes, 8 edges)	30
5.2.2	Sachs (Continuous, 11 nodes, 17 edges)	31
5.2.3	Alarm (Discrete, 37 nodes, 46 edges)	31
5.2.4	Child (Discrete, 20 nodes, 25 edges)	32
5.2.5	Insurance (Discrete, 27 nodes, 52 edges)	32
5.3	Hyperparameter Sensitivity	32

5.4	Cross-Dataset Patterns	33
5.4.1	Data Type Effects on Algorithm Performance	33
5.4.2	The Skeleton vs Directed Gap	33
5.4.3	Skeleton–directed gap: quantifying the orientation difficulty	33
5.4.4	Error-type decomposition: extra, missing, and reversed edges	35
5.4.5	Node-level concentration of errors in larger graphs (ALARM and Insurance) . .	37
5.4.6	Cross-Algorithm Agreement Analysis	38
5.4.7	Runtime vs Accuracy	39
5.4.8	Algorithm Summary	40
5.4.9	Statistical Significance	40
5.5	Edge-Level Case Studies	42
5.5.1	Case Study 1: Asia - Smoking → LungCancer	42
5.5.2	Case Study 2: Sachs - PKA → Mek	43
5.6	Detecting Analyst Mis-specification	43
5.6.1	Conditional Independence Testing Results	43
5.6.2	Algorithm recovery of omitted edges	45
5.6.3	Key-edge bootstrap stability across algorithms	47
6	Discussion	48
6.1	Answers to Research Questions	48
6.1.1	RQ1: Benchmark accuracy across data types and network sizes	48
6.1.2	RQ2: Detecting omitted and spurious edges with CI tests	49
6.1.3	RQ3: Practitioner guidance	49
6.2	The Analyst-in-the-Loop Paradigm	50
6.2.1	A practitioner decision guide for method selection	50
6.2.2	Software ecosystem considerations	50
6.2.3	Common pitfalls in analyst-facing causal discovery	51
6.2.4	Deployment checklist for analyst-in-the-loop use	51
6.3	Robustness to Assumption Violations	52
6.4	Case Study: End-to-End Analyst Workflow	52
6.5	Threats to Validity	53
6.5.1	Internal Validity	53
6.5.2	Construct Validity	53

6.5.3 External Validity	53
6.6 Limitations	54
6.7 Ethical and Responsible Use	54
6.7.1 The Risk of False Confidence	54
6.7.2 Algorithmic Bias	54
6.7.3 Dual Use	54
6.8 Lessons Learned	55
6.9 Future Research Directions	55
7 Conclusion	55
7.1 Reproducibility and artifact availability	56

List of Figures

1	Pseudocode for the semi-synthetic data generation process.	19
2	Pseudocode for the PC algorithm.	23
3	Pseudocode for the GES algorithm.	24
4	Structure of the Asia network. Nodes represent variables and arrows denote direct causal effects. This network is a standard benchmark for evaluating causal discovery algorithms.	27
5	Benchmarking pipeline. Data are generated from benchmark networks, algorithms are run to learn the structure, metrics are computed, mis-specification analyses are performed, and bootstrap edge stability is recorded.	28
6	Skeleton F_1 scores by dataset and algorithm. The dashed horizontal line indicates $F_1 = 0.50$, representing performance no better than a naive baseline.	29
7	Precision–recall scatter plot for skeleton recovery. Each point represents one algorithm–dataset combination. Dashed curves show iso– F_1 contours.	30
8	Algorithm performance stratified by data type. Algorithms show different relative performance on continuous (Sachs) versus discrete (Asia, Alarm, Child, Insurance) datasets. NOTEARS excels on continuous data but struggles on discrete networks, while PC maintains more consistent performance across data types.	33
9	Skeleton F_1 versus directed F_1 . Points below the diagonal indicate orientation errors. Dataset abbreviations: ASI = Asia, SAC = Sachs, ALA = Alarm, CHI = Child, INS = Insurance.	35
10	Breakdown of structural errors by type: False Positives (Extra), False Negatives (Missing), and Reversals. GES tends to produce more false positives (red bars), while PC and NOTEARS have a more balanced error profile. Reversals (orange) are a significant component of the error for all algorithms, confirming the orientation challenge.	36
11	Heatmap of Structural Hamming Distance across algorithms and datasets. Darker colors indicate higher SHD (more errors). The matrix reveals that Alarm and Insurance are universally challenging (darker columns), while algorithm performance varies considerably by dataset. NOTEARS shows the highest variance, with near-zero SHD on Sachs but high SHD on Insurance.	38
12	Runtime comparison across datasets (log scale). PC and COSMO are consistently faster than GES and NOTEARS. Note the logarithmic y-axis; GES is orders of magnitude slower on larger networks.	40
13	Direct runtime comparison across algorithms and datasets. Bar heights represent execution time in seconds. GES consistently requires the most time, while PC and COSMO are fastest across all networks.	41
14	Algorithm performance profiles averaged across datasets. Axes represent skeleton F_1 , directed F_1 , precision, recall and speed.	41

15	Critical difference diagram for skeleton F_1 . Algorithms are ranked by average performance across datasets, with rank 1 being best. The horizontal bar shows the Nemenyi critical difference (CD) at $\alpha = 0.05$. Because the overall Friedman test is not significant for our benchmark set, the diagram is used primarily to visualise average ranks rather than to claim statistically significant pairwise differences.	42
16	Conditional independence test results for sensitivity analysis. Compares test statistics for missing edge detection (where high values indicate correct detection) versus spurious edge detection (where low values indicate correct identification of unnecessary edges). Error bars represent confidence intervals where applicable.	45
17	Signal strength of conditional independence tests for mis-specification detection. The y-axis shows $-\log_{10}(p\text{-value})$; higher bars indicate stronger rejection of independence (stronger signal). The dashed line represents the $\alpha = 0.05$ threshold. Missing edges (green) consistently produce strong signals. Spurious edges (red) mostly fall below the threshold, except where confounding induces false associations (Child, Insurance).	46
18	Algorithm performance comparison across sensitivity analysis scenarios. Shows how well each algorithm performs when data is generated from the true graph but evaluated against different misspecification types (missing edges vs spurious edges). Performance metrics indicate algorithms' robustness to analyst errors.	47
19	Recovery status of the specific omitted edge by each algorithm. Green indicates the edge was correctly recovered and oriented. Yellow indicates the edge was found but reversed. Red indicates the edge was missed entirely. Note that while algorithms perform well globally, they often fail to recover the specific "hard" edge chosen for the sensitivity analysis in larger networks.	48

List of Tables

2	Worked examples of d-separation in the Asia network (Figure 4)	4
3	A sketch timeline of representative developments in causal structure learning. This is not exhaustive; it highlights the conceptual shift from combinatorial search and CI testing to differentiable and learning-based approaches.	13
4	Representative causal discovery methods and the assumptions most relevant to mis-specification robustness. “Output” refers to the standard graphical object returned when run on observational data.	14
5	Summary of benchmark networks. Edges refers to the number of directed edges in the ground-truth DAG. Density is $ E /(V (V - 1)/2)$, the fraction of possible edges present. Sample size n indicates the number of observations generated for each experiment.	20
6	Algorithm implementations and hyperparameter settings. CI = conditional independence test. All algorithms were allowed to run to completion.	20
7	Skeleton recovery performance. Precision, recall and F_1 treat edges as undirected. Bold indicates the best F_1 score for each dataset.	29
8	Edge-by-edge recovery analysis for the Asia network. ✓ indicates correct recovery (including direction), ✗ indicates missing edge, ↳ indicates reversed edge.	31
9	PC Algorithm performance on Alarm with varying α	32
10	Directed edge recovery performance. Reversed edges count as errors.	34
11	Skeleton–directed gap $\Delta F_1 = F_1^{\text{skel}} - F_1^{\text{dir}}$ computed from Tables 7 and 10.	34
12	Edge-error decomposition from per-run diff logs. “Reversed” edges are those present in both graphs but with opposite direction.	37
13	Jaccard similarity of predicted skeletons on the Alarm dataset. Values close to 1 indicate high agreement.	38
14	Runtime in seconds. Bold indicates the fastest algorithm for each dataset.	39
15	Mis-specification scenarios. Missing edges are true causal links removed from the analyst’s DAG; spurious edges are false links added.	43
16	Conditional independence test results for mis-specification detection. Missing edges should show significant dependence (reject H_0); spurious edges should show non-significant results (fail to reject).	44
17	Algorithm performance in sensitivity analysis (vs. true graph).	46
18	Bootstrap stability of the key missing edge: frequency with which the removed true edge appears in the learned graph across bootstrap resamples.	47
19	Decision guide for choosing a discovery approach in analyst-in-the-loop settings. Recommendations reflect the empirical patterns in Section 5 and common assumptions in the literature.	50

List of Abbreviations

Abbrev.	Meaning
CI	Conditional independence
CPDAG	Completed partially directed acyclic graph
COSMO	Constrained Orientations by Sequential M Operation
DAG	Directed acyclic graph
FDR	False discovery rate
GES	Greedy Equivalence Search
IID	Independent and identically distributed
NOTEARS	Non-combinatorial optimisation via trace exponential and augmented lagrangian for structure learning
PC	Peter–Clark algorithm
SCM	Structural causal model
SEM	Structural equation model
SHD	Structural Hamming distance
SID	Structural intervention distance

1 Introduction

Causal diagrams—directed acyclic graphs (DAGs) that encode cause-effect relationships between variables—are indispensable for reasoning about interventions and policy decisions. These graphs consist of nodes representing variables and directed edges denoting direct causal effects. A directed graph qualifies as a DAG if it contains no directed cycles (no node can reach itself by repeatedly following arrows). When DAGs are interpreted causally, analysts typically make additional assumptions—most notably *causal sufficiency* (no unmeasured common causes of included variables) and *faithfulness* (conditional independences in the data correspond to graph separations)—that are convenient but rarely guaranteed in practice [1, 2]. When practitioners draw such diagrams incorrectly, downstream causal inference and decision-making can be compromised, motivating systematic methods to benchmark discovery algorithms and to diagnose analyst mis-specification. We operationalize analyst mis-specification as an analyst-proposed DAG that omits a true causal link or includes a non-existent link.

This thesis pursues three primary objectives. First, we benchmark multiple structure-learning methods—PC, GES, NOTEARS, and COSMO—on established benchmark networks to evaluate how well they recover known causal structures. The standardized framework ensures that comparisons are meaningful by using shared metrics and reproducible implementations. Second, we study how analyst mis-specification propagates: if a practitioner erroneously removes an edge or inserts a spurious link, can the data alert them? We design controlled experiments where the true DAG generates data but the analyst’s DAG deviates from it. Third, we survey open-source tools for drawing and testing DAGs to provide practical guidance on constructing and validating causal diagrams.

These goals translate into three research questions:

RQ1 *How do causal discovery algorithms compare in recovering ground-truth network structures across different data types and network sizes?* We hypothesise that algorithm performance depends strongly on the match between algorithmic assumptions (e.g., linearity, Gaussianity) and data characteristics.

RQ2 *Can conditional independence tests reliably detect when an analyst’s DAG omits a true edge or includes a spurious one?* We hypothesise that omitted edges produce significant dependence signals, while spurious edges yield non-significant results, enabling data-driven model criticism.

RQ3 *What practical guidance can we offer practitioners for selecting algorithms and validating their causal assumptions?* We synthesise our empirical findings into actionable recommendations.

The remainder of the thesis is structured as follows. Section 2 reviews the basics of causal DAGs, conditional independence (CI), and common structure-learning algorithms. Section 3 summarises related work on benchmarking causal discovery, mis-specification detection, and DAG drawing software. Section 4 details our datasets, data generation procedures, algorithms, metrics, and mis-specification protocols. Section 5 presents benchmark and sensitivity results. Finally, Section 6 discusses practical implications, limitations, and recommendations for practitioners, before Section ?? concludes.

2 Theoretical Framework

Before algorithms can be benchmarked or mis-specification detected, a rigorous mathematical framework is required. This chapter lays the probabilistic and graphical foundations for causal discovery from observational data. The concepts introduced—d-separation, the Causal Markov Condition, and Faithfulness—are not merely theoretical niceties; they directly determine what can and cannot be inferred from conditional independence patterns in data. Following these preliminaries, we examine the four algorithms central to this thesis: PC, GES, NOTEARS, and COSMO. Each embodies a different philosophy (constraint-based testing, score-based optimisation, or continuous differentiable search) and carries distinct assumptions about data and structure. We conclude by formalising the evaluation metrics—SHD, SID, precision, recall, and F_1 —that will be used throughout the empirical chapters to quantify algorithm performance and assess the severity of analyst errors.

2.1 Probabilistic Graphical Models

Probabilistic graphical models (PGMs) provide a language for encoding complex multivariate dependencies via graph structures. Rather than storing a full joint distribution—exponential in the number of variables—a PGM exploits conditional independence to represent the same information compactly. The graph topology itself becomes a hypothesis about which variables directly influence which others.

Directed acyclic graphs. Formally, a graph $G = (V, E)$ comprises a vertex set $V = \{X_1, \dots, X_d\}$ (corresponding to random variables) and a set of directed edges $E \subseteq V \times V$. An edge $(X_i, X_j) \in E$, denoted $X_i \rightarrow X_j$, asserts a direct causal or predictive relationship from X_i to X_j . Within the graph, a *path* is any sequence of distinct nodes connected by edges (ignoring direction), while a *directed path* follows edges strictly in the direction of the arrows. A *cycle* is a directed path that returns to its starting node. A graph is a Directed Acyclic Graph (DAG) if it contains no such cycles—equivalently, if nodes can be topologically ordered so that all edges point forward.

The Causal Markov condition. The Causal Markov Condition is the bridge from graph topology to probability. It states that once we condition on a variable's direct causes (its parents in the DAG), that variable becomes independent of all variables that are not its descendants. More formally, the *Local Markov Property* asserts that each X_i is conditionally independent of its non-descendants given $\text{Pa}(X_i)$. This local independence statement, applied to every node, implies a global factorisation of the joint distribution:

$$P(X_1, \dots, X_d) = \prod_{i=1}^d P(X_i | \text{Pa}(X_i)) \quad (1)$$

Rather than specifying all $2^d - 1$ marginal and conditional distributions, we need only d local conditional distributions—one per node—each depending solely on that node's parents. This exponential compression is what makes PGMs computationally tractable.

d-Separation. While the Markov condition tells us how to compute probabilities given a graph, *d-separation* (directional separation) provides the inverse tool: reading conditional independencies

directly from the graph's topology. A path p between nodes X and Y is *blocked* by a conditioning set Z under two scenarios: first, if the path contains a chain ($i \rightarrow m \rightarrow j$) or fork ($i \leftarrow m \rightarrow j$) where the middle node m belongs to Z —informally, conditioning on a mediator or common cause “closes” the information flow; second, if the path contains a collider ($i \rightarrow m \leftarrow j$) where neither m nor any of its descendants appear in Z —colliders naturally block paths unless we condition on them or their descendants, which paradoxically “opens” the path via explaining-away reasoning. If every path between X and Y is blocked by Z , the variables are d-separated: $X \perp_G Y | Z$. Under the Causal Markov assumption, this graphical separation implies probabilistic independence $X \perp_P Y | Z$.

Faithfulness. The Markov condition is one-directional: graph separation implies statistical independence. However, the converse need not hold. A distribution might exhibit an independence that is *not* encoded in the graph—for instance, if parameters conspire so that two causal paths cancel exactly (a “knife-edge” condition). Such accidental independencies are possible but fragile: small perturbations restore the dependence. The *Faithfulness Assumption* rules out these measure-zero coincidences, asserting that the distribution P contains *only* those independencies implied by the d-separation structure of G :

$$X \perp_P Y | Z \iff X \perp_G Y | Z \quad (2)$$

This bi-directional equivalence is pivotal for causal discovery: it licenses the inference of edges (or their absence) from patterns of conditional independence observed in data. Without faithfulness, a statistical test rejecting independence would not imply the presence of an edge, undermining constraint-based discovery entirely.

Structural causal models. The *structural causal model* (SCM) framework formalises how a DAG generates data via functional relationships. Given a DAG $G = (V, E)$, an SCM specifies for each variable X_i a deterministic assignment $X_i := f_i(X_{\text{Pa}(i)}, U_i)$, where f_i is an arbitrary measurable function and U_i represents unobserved exogenous noise. The variables X emerge as solutions to this system of equations, with randomness introduced only through the exogenous terms U . Provided the noise terms are jointly independent and the DAG encodes the recursive causal ordering, the Causal Markov condition follows automatically [1, 2]. This construction yields the Markov factorisation $p(x_1, \dots, x_d) = \prod_{i=1}^d p(x_i | x_{\text{Pa}(i)})$ (Equation 5), which serves as the computational backbone for both score-based and constraint-based discovery: score-based methods maximise factorised likelihoods, while constraint-based methods test the conditional independencies that the factorisation entails.

From the SCM to a factorised likelihood (proof sketch). Assume G admits a topological ordering of variables such that parents precede children. By the chain rule, any joint distribution can be written as

$$p(x_1, \dots, x_d) = \prod_{i=1}^d p(x_i | x_{1:i-1}). \quad (3)$$

Under the Markov property implied by the SCM graph, each X_i is independent of its non-descendants given its parents. In particular, conditioning on all previous variables $X_{1:i-1}$ is equivalent to conditioning only on $X_{\text{Pa}(i)}$, because the remaining variables in $\{1, \dots, i-1\} \setminus \text{Pa}(i)$ are non-descendants of i .

Therefore,

$$p(x_i | x_{1:i-1}) = p(x_i | x_{\text{Pa}(i)}), \quad (4)$$

and we obtain the *Markov factorisation*

$$p(x_1, \dots, x_d) = \prod_{i=1}^d p(x_i | x_{\text{Pa}(i)}). \quad (5)$$

This identity is the key bridge between graphs and data: (5) is exactly the decomposition exploited by both score-based methods (which optimise a factorised likelihood or a penalised variant) and constraint-based methods (which test conditional independencies implied by the factorisation) [2, 3].

d-Separation in practice. The preceding definitions become concrete when applied to real benchmark networks. Understanding how d-separation works in practice is essential for interpreting algorithm behaviour: constraint-based methods like PC systematically query these independencies, while score-based methods implicitly encode them through the factorised likelihood. Table 2 demonstrates several representative queries in the Asia network. The first row illustrates collider blocking: *Smoking* and *Tuberculosis* are marginally independent because their only connecting path is blocked by the unobserved collider *Either*. The second row shows the paradoxical “explaining away” effect—conditioning on *Either* *opens* the blocked path, inducing dependence between two variables that were previously independent.

Table 2: Worked examples of d-separation in the Asia network (Figure 4).

Query (graph statement)	d-separated?	Reasoning (path-wise)
<i>Smoking</i> \perp <i>Tuberculosis</i>	Yes	The main connecting path <i>Smoking</i> \rightarrow <i>LungCancer</i> \rightarrow <i>Either</i> \leftarrow <i>Tuberculosis</i> is blocked by the collider at <i>Either</i> . No other open path exists without conditioning.
<i>Smoking</i> \perp <i>Tuberculosis</i> <i>Either</i>	No	Conditioning on the collider <i>Either</i> <i>opens</i> the previously blocked path, producing the classic “explaining-away” dependence.
<i>Smoking</i> \perp <i>Xray</i> <i>Either</i>	Yes	<i>Xray</i> is a child of <i>Either</i> . Once <i>Either</i> is conditioned on, the path <i>Smoking</i> \rightarrow <i>LungCancer</i> \rightarrow <i>Either</i> \rightarrow <i>Xray</i> is blocked at <i>Either</i> (a conditioned non-collider).
<i>Tuberculosis</i> \perp <i>Bronchitis</i>	Yes	Two prominent paths are blocked: (i) <i>Tuberculosis</i> \rightarrow <i>Either</i> \rightarrow <i>Dyspnea</i> \leftarrow <i>Bronchitis</i> is blocked by the collider at <i>Dyspnea</i> ; (ii) the path through <i>Either</i> is blocked as above.

These concrete queries illustrate the central mechanism behind constraint-based learning: algorithms such as PC recover a skeleton by repeatedly testing whether statements like those in Table 2 are supported by the data (for various conditioning sets), and then orient edges to form an equivalence-class representation. The Asia network is small enough for manual verification, but real-world applications involve dozens or hundreds of variables, where exhaustive path enumeration becomes infeasible and algorithmic search is necessary.

Consider the larger ALARM network [32] (37 nodes, 46 edges), which models intensive-care patient monitoring. For Pulmonary Embolism (PE) and Central Venous Pressure (CVP), the relationship is mediated by Pulmonary Artery Pressure (PAP): the chain $PE \rightarrow PAP \rightarrow CVP$ renders them marginally dependent, yet conditioning on PAP blocks the only connecting path ($PE \perp CVP \mid PAP$). This pattern—marginal dependence becoming conditional independence—is the signature of a mediating variable. Conversely, Catecholamine and Hypovolemia are marginally independent (no direct path) but connected via the collider at Heart Rate (HR); conditioning on HR induces dependence through the explaining-away mechanism.

The Child network [?], designed for diagnosing congenital heart disease (“blue baby” syndrome), presents similar challenges. Birth Asphyxia and Age are marginally independent (both are root causes with no connecting path), yet conditioning on their common descendant Disease—or any of its observable proxies—can induce a spurious association. This underscores a general principle: colliders and their descendants must be handled with care, as naive conditioning can introduce bias rather than remove it.

Identifiability and theoretical guarantees. A fundamental limit of observational causal discovery is *Markov equivalence*: multiple distinct DAGs can imply identical conditional independence structures, making them indistinguishable from observational data alone. For instance, the three graphs $A \rightarrow B \rightarrow C$, $A \leftarrow B \leftarrow C$, and $A \leftarrow B \rightarrow C$ all encode the single conditional independence $A \perp C \mid B$ and no others. Verma and Pearl [30] characterised this phenomenon precisely: two DAGs belong to the same Markov Equivalence Class (MEC) if and only if they share the same skeleton (undirected graph) and the same set of v-structures (unshielded colliders $X \rightarrow Z \leftarrow Y$). Such classes are compactly represented by a Completed Partially Directed Acyclic Graph (CPDAG), where directed edges appear in all members of the class and undirected edges denote ambiguous orientations. This is not a failure of the algorithms—it is an inherent limitation of passive observation.

Algorithm guarantees must therefore be stated relative to the CPDAG, not any single DAG. **Constraint-based methods** (like PC) assume faithfulness and, in the infinite-sample limit with perfect conditional independence tests, recover the true CPDAG [2]. Finite-sample consistency requires that the CI test error rates vanish appropriately; Kalisch and Bühlmann [?] establish such results for high-dimensional Gaussian graphs under suitable sparsity and faithfulness conditions. **Score-based methods** (like GES) search for the graph maximising a penalised likelihood score (e.g., BIC). Chickering [3] proved that GES identifies the correct CPDAG in the large-sample limit when the score is consistent and score-equivalent. In finite samples, the two families exhibit complementary failure modes: PC is brittle to individual test errors that propagate through the skeleton construction, whereas GES is more robust to noise but sensitive to model mis-specification (e.g., applying BIC’s linear-Gaussian assumptions to discrete data). Interventional data—observations under $\text{do}(Y = y)$ manipulations—can break Markov equivalence by revealing asymmetries in the causal mechanisms, but this thesis focuses on the observational regime where the CPDAG is the identifiability ceiling.

2.2 Constraint-Based Discovery: The PC Algorithm

The PC algorithm (Peter–Clark) is the archetypal constraint-based method. Named after Peter Spirtes and Clark Glymour, it reconstructs the CPDAG by systematically testing conditional independence

statements implied by faithfulness [2]. The algorithm’s elegance lies in its two-phase design: first learn which variables are adjacent (the skeleton), then determine which edges can be oriented.

Procedure. *Phase 1: Skeleton identification.* Starting from a complete undirected graph, PC tests whether each pair (X, Y) is conditionally independent given subsets S of their neighbours. Independence tests begin with the empty set ($|S| = 0$), then condition on single variables ($|S| = 1$), pairs ($|S| = 2$), and so forth, up to the maximum degree. Whenever a test finds $X \perp Y | S$, the edge is removed and S is recorded as the separating set. Because we assume faithfulness, every non-adjacent pair in the true graph has a separating set, and every adjacent pair lacks one—so in the large-sample limit with perfect tests, this phase recovers the skeleton exactly. *Phase 2: Orientation.* With the skeleton fixed, PC identifies v-structures: for every triple $X - Z - Y$ where X and Y are non-adjacent, if Z was not in their separating set S_{XY} , then Z must be a collider ($X \rightarrow Z \leftarrow Y$). These v-structures are oriented first; then Meek’s orientation rules propagate the constraints to avoid creating new v-structures or cycles, orienting as many additional edges as possible while respecting the equivalence class.

Complexity and stability. The number of potential conditioning sets grows exponentially in the graph’s maximum degree, making worst-case complexity prohibitive for dense graphs. However, real-world causal networks are often sparse (average degree $\ll d$), rendering PC practical. A more insidious issue is **order dependence**: in finite samples, ties or marginal test failures can depend on the order in which variables are processed, yielding different skeletons. We employ the PC-Stable variant [16], which ensures order-independence by fixing the candidate conditioning sets (the adjacency sets at the start of each round) rather than updating them dynamically.

2.3 Score-Based Discovery: GES

Score-based methods recast structure learning as model selection: find the graph G that optimises a balance between fit and complexity. Rather than testing individual conditional independencies, these methods evaluate candidate graphs holistically via a scoring function $S(G, D)$.

Scoring functions. We employ the Bayesian Information Criterion (BIC) [26], a penalised likelihood criterion:

$$\text{BIC}(G; D) = \log L(G; \hat{\theta}) - \frac{k_G}{2} \log n \quad (6)$$

where $\log L(G; \hat{\theta})$ measures fit (higher is better), k_G counts free parameters, and n is the sample size. The penalty $\frac{k_G}{2} \log n$ grows with model complexity, discouraging overfitting. Two properties make BIC attractive for causal discovery. First, it is *decomposable*: for a DAG, the likelihood factorises over nodes (Equation 5), so adding or removing edges changes only the local terms for affected variables. This locality enables incremental score updates. Second, BIC is *score-equivalent*: Markov equivalent DAGs receive identical scores, ensuring the search respects equivalence classes. For linear-Gaussian models, the local likelihood term for node i reduces to a penalised residual sum of squares from regressing X_i on its parents.

Greedy Equivalence Search (GES). Exhaustive search over all DAGs is super-exponential ($d! \cdot 2^{d(d-1)/2}$ structures for d nodes). GES [3] sidesteps this by searching the space of equivalence classes (CPDAGs) via a two-phase greedy heuristic. *Forward Equivalence Search* (FES) begins with an empty graph and iteratively adds the single edge (or edge orientation) that most improves the score, continuing until no addition yields a gain. *Backward Equivalence Search* (BES) then removes edges one at a time to further refine the score. This two-phase design exploits a key insight: starting sparse and adding edges (forward) avoids early commitment to spurious structures, while pruning (backward) corrects overfitting from the greedy forward pass. Chickering proved that in the large-sample limit—where the score becomes an accurate proxy for the true generative model—GES identifies the correct CPDAG. In our experiments, we apply BIC uniformly across all datasets (including discretised data) to stress-test its robustness to model mis-specification.

2.4 Continuous Optimisation: NOTEARS

Traditional search-based methods (PC, GES) navigate discrete graph spaces via edge additions, deletions, or reversals. NOTEARS (Non-combinatorial Optimisation via Trace Exponential and Augmented lagRangian for Structure learning) [19] takes a radically different approach: formulate structure learning as a *continuous* optimisation problem over weighted adjacency matrices, where the acyclicity constraint—previously thought to be inherently combinatorial—is expressed as a smooth, differentiable equality constraint.

Differentiable acyclicity. The key innovation is a function $h : \mathbb{R}^{d \times d} \rightarrow \mathbb{R}$ such that $h(W) = 0$ if and only if the graph represented by W is acyclic. Define $h(W) = \text{tr}(e^{W \circ W}) - d$, where $W \circ W$ is the element-wise square of W and e^A is the matrix exponential. Intuitively, the matrix exponential expands as $\sum_{k=0}^{\infty} \frac{(W \circ W)^k}{k!}$; the trace sums the diagonal, aggregating self-loop counts weighted by path lengths. If the graph contains a cycle, some diagonal entries explode; if acyclic, the trace reduces to exactly d (one per node). The gradient $\nabla_W h(W) = 2W \circ (e^{W \circ W})^\top$ is computable in closed form, enabling standard gradient-based solvers.

Objective and limitations. NOTEARS minimises a regularised loss subject to the acyclicity equality constraint:

$$\min_W \ell(W; X) + \lambda \|W\|_1 \quad \text{subject to} \quad h(W) = 0 \quad (7)$$

where $\ell(W; X)$ is typically the squared-error loss for linear structural equation models, and $\lambda \|W\|_1$ penalises edge proliferation (encouraging sparsity). An augmented Lagrangian solver citeNocedal2006 alternates between minimising

ell +
lambda
 $|W$
 $|_1 +$
 $\frac{rho}{2} h(W)^2 +$
alpha $h(W)$ (where
rho and

α are penalty/multiplier parameters) and updating the multipliers, gradually enforcing $h(W)$ to 0.

textbf{Critical assumption:} The original NOTEARS formulation assumes linear structural equation models with continuous variables. When applied to discrete or mixed data (as in our benchmarks), it treats categorical state values as if they were continuous, constituting a model mis-specification. Despite this mismatch, NOTEARS often performs surprisingly well in practice, suggesting robustness to moderate violations. We include it specifically to evaluate how continuous-optimisation methods behave under such mis-specified conditions, providing a contrast to methods explicitly designed for discrete data.

2.5 Regression-Based Discovery: COSMO

Massidda et al.

citeMassidda2023 introduce COSMO (Causal Ordering via Scalable Modelling), which reduces computational complexity to $O(d^2)$ by replacing the explicit acyclicity constraint with a learned causal ordering. Rather than searching over graph structures, COSMO learns a latent order

$\Theta = ($
 $\theta_1,$
 $\dots,$
 $\theta_d)$ where
 θ_i
in

\mathbb{R} specifies node i 's position. Edges X_i

to X_j are permitted only if

$\theta_i <$
 θ_j , automatically ensuring acyclicity.

paragraphMechanism. COSMO alternates between two steps: an textbf{M-Step} fixes the ordering

Θ and estimates each node's parents via regularised regression (typically Lasso or elastic net), selecting the subset of predecessors in the ordering that best predict the node's values; and an textbf{O-Step} updates

Θ to minimise the overall prediction loss, shifting nodes forward or backward in the order to reduce residual variance. This iterative refinement continues until convergence. Because the ordering is represented by continuous parameters, both steps can exploit efficient regression solvers, avoiding the combinatorial search over graph topologies.

To ensure reliability in finite samples, COSMO employs textbf{stability selection}

citeMeinshausen2010: the full procedure is run on multiple bootstrap resamples, and only edges that appear with high frequency (e.g., > 70%) across runs are retained in the final output. This guards against spurious edges that arise from sampling variability, at the cost of potentially missing weak but genuine relationships.

2.6 Evaluation Metrics

Evaluating causal discovery algorithms requires a multi-faceted approach. No single metric captures all aspects of performance: structural similarity measures (like SHD) assess graph topology but ignore causal semantics, while intervention-based metrics (like SID) reflect downstream reasoning tasks but are sensitive to specific query distributions. We therefore employ a portfolio of complementary metrics, each illuminating a different facet of algorithm behaviour.

Structural Hamming Distance (SHD). The most widely used metric is SHD, which counts the minimum number of edge operations (additions, deletions, or reversals) required to transform the learned graph \hat{G} into the true graph G [17]:

$$\text{SHD}(G, \hat{G}) = |E \setminus \hat{E}| + |\hat{E} \setminus E| + \text{flips} \quad (8)$$

This metric is intuitive and directly reflects structural fidelity. However, validating against a *DAG* rather than its Markov equivalence class is problematic: an algorithm that correctly returns an undirected edge $X - Y$ (because the direction is unidentifiable) would be penalised for not guessing the “true” orientation. To avoid this pitfall, we employ a **CPDAG-aware SHD**: both G and \hat{G} are first converted to their respective CPDAGs, then compared. Undirected edges in the CPDAG are treated as bidirectional, so an algorithm receives credit for correctly leaving an edge unoriented. This ensures that the metric respects the identifiability ceiling imposed by observational data.

Precision, recall, and F1. SHD aggregates errors into a single number, obscuring the balance between false positives (spurious edges) and false negatives (missing edges). To disentangle these failure modes, we compute standard classification metrics: Precision ($P = TP/(TP + FP)$) measures what fraction of predicted edges are correct, Recall ($R = TP/(TP + FN)$) measures what fraction of true edges are recovered, and the F_1 score ($2PR/(P + R)$) provides their harmonic mean. These are reported separately for the **skeleton** (undirected adjacency, ignoring orientation) and the **directed graph** (requiring correct orientation). Large discrepancies between skeleton and directed scores—for instance, high skeleton recall but low directed recall—indicate that the algorithm successfully identifies which pairs of variables are adjacent but struggles to orient those edges. This diagnostic capability is particularly valuable when interpreting algorithm behaviour: constraint-based methods may excel at skeleton recovery yet falter at orientation due to finite-sample test errors, while score-based methods might confidently orient edges but add spurious connections.

Structural Intervention Distance (SID). While SHD treats all edge errors equally, the *Structural Intervention Distance* (SID) [25] weights errors by their impact on causal reasoning. Formally, SID counts the number of ordered pairs (X, Y) for which the interventional distribution $P(Y \mid \text{do}(X))$ inferred from the learned graph \hat{G} differs from the true interventional effect implied by G . In other words, two graphs have zero SID if and only if they induce identical interventional predictions for all possible pairs of variables. This captures the intuition that some edge errors are more consequential than others: reversing an edge in a long causal chain can corrupt many downstream interventional queries, whereas adding a redundant edge between already d-connected nodes may have minimal

impact. SID thus provides a *task-aligned* evaluation metric, directly reflecting the downstream utility of the learned structure for interventional reasoning. In our experiments, we compute SID by enumerating all (X, Y) pairs and checking whether the adjustment sets required to identify $P(Y \mid \text{do}(X))$ differ between G and \hat{G} .

2.7 Structural Causal Models and Interventions

While structure learning algorithms focus on recovering the graph G , the ultimate goal of causal inference often involves predicting the outcome of interventions: what would happen to Y if we forcibly set $X = x$, overruling its natural causes? The Structural Causal Model (SCM) framework formalises this via *do*-calculus [1]. An intervention $\text{do}(X_k = x)$ is represented by replacing the structural equation $X_k = f_k(\text{Pa}_k, U_k)$ with the constant assignment $X_k = x$, effectively severing the causal influence of X_k 's parents. The resulting mutilated graph—and its implied factorisation—determines the post-intervention distribution $P(\mathbf{X} \mid \text{do}(X_k = x))$.

Pearl's do-calculus provides syntactic rules for determining when such interventional quantities can be identified from purely observational data, using properties of the graph (d-separation) to guide adjustment strategies. Crucially, these identifiability results assume knowledge of the *true* causal graph G . If the learned structure \hat{G} contains errors, the resulting interventional predictions can be arbitrarily wrong—a missing edge may omit a necessary confounder adjustment, while a reversed edge can invert the causal direction entirely. Thus, verifying the correctness of the learned graph is a prerequisite for valid downstream causal inference. Our benchmark evaluates how reliably algorithms recover structures that would support correct interventional reasoning.

2.8 Assumptions, Identifiability, and Scope

All structure learning algorithms rest on a tower of assumptions, violations of which can silently corrupt the learned graph. We can categorise these into three layers. **Graphical assumptions** include the Causal Markov Condition (graph structure implies conditional independencies), Faithfulness (no accidental cancellations create additional independencies), and Causal Sufficiency (no unobserved confounders). Violations are common in practice: latent confounders introduce spurious associations, while finely tuned parameters can produce measure-zero faithfulness violations. **Statistical assumptions** concern the validity of independence tests or scoring functions given the sample size and data distribution. For instance, Fisher's z -test assumes joint Gaussianity; applying it to heavily skewed or discrete data inflates Type I/II error rates. **Model class assumptions** impose structural constraints on the functional form—NOTEARS assumes linear additive noise models, GES implicitly assumes the data are well-described by the parametric family used in the BIC calculation.

Our experiments use *semi-synthetic* data generated from known Bayesian network structures, allowing us to control graphical assumptions (no latent confounders, faithfulness guaranteed by construction) while varying sample size and introducing analyst mis-specifications (e.g., applying continuous-data methods to discretised observations). Real-world violations—particularly latent confounding—remain a critical challenge discussed in Section 6.5.

2.9 Conditional Independence Testing

Conditional independence (CI) tests serve a dual role in our framework. First, they are the fundamental primitive for constraint-based discovery: algorithms like PC repeatedly query whether $X \perp Y | S$ to construct the skeleton and identify v-structures. Second, they function as a tool for *model criticism*, allowing us to validate whether a learned graph's implied independencies hold in the data.

For continuous (or Gaussian-assumed) data, we employ **Fisher's z -test** on partial correlations: the null hypothesis $X \perp Y | S$ is rejected if $z = \frac{1}{2} \log \frac{1+\hat{\rho}_{XY|S}}{1-\hat{\rho}_{XY|S}} \sqrt{n - |S| - 3}$ exceeds the critical value for significance level α (typically 0.05). For discrete data, we use the **G^2 (log-likelihood ratio) test**, which compares observed and expected cell counts under the independence hypothesis. Both tests balance Type I errors (falsely rejecting independence, leading to spurious edges) against Type II errors (failing to detect dependence, leading to missing edges). The standard choice $\alpha = 0.05$ reflects this trade-off, though the optimal threshold is data-dependent.

A critical challenge is **error propagation** in multi-test algorithms like PC. The skeleton phase may perform thousands of CI tests, and each erroneous decision contaminates subsequent tests: if an edge is incorrectly removed, it will not be considered in later conditioning sets, potentially blocking the discovery of true conditional independencies. This cascading failure mode explains why PC's performance can degrade sharply in noisy or small-sample regimes, even though individual tests remain calibrated.

2.10 Statistical Tests for Comparing Algorithms

To rigorously compare algorithms across multiple datasets and establish whether observed performance differences are statistically meaningful, we follow the recommendations of Demšar [24] for multi-classifier comparison. Standard parametric tests (e.g., paired t -tests) make distributional assumptions (normality, independence) that are often violated when comparing algorithms on a small number of heterogeneous datasets. Demšar advocates non-parametric alternatives that rank algorithms on each dataset, then test for differences in average ranks.

We employ the **Friedman test** as an omnibus procedure to detect whether any algorithms differ significantly. If the null hypothesis of equal mean ranks is rejected, we proceed to pairwise comparisons via the **Nemenyi post-hoc test**, which controls the family-wise error rate. Results are visualised using **Critical Difference (CD) diagrams**: algorithms are arranged along a horizontal axis by their average rank, and groups of algorithms whose ranks do not differ significantly (at $\alpha = 0.05$) are connected by a horizontal bar. This provides an intuitive summary of the performance hierarchy.

Important caveat: Given our small number of datasets ($N = 5$), these tests have limited statistical power. A non-significant result does not prove equivalence; it may simply reflect insufficient data. We therefore treat the Friedman/Nemenyi results as *descriptive complements* to the raw performance metrics, useful for identifying clear winners but not definitive for marginal differences.

Robustness considerations. Beyond central tendency (mean or median SHD), robustness to sample size and hyperparameter choices can be assessed through systematic variation. Repeating benchmarks across a range of n values (e.g., 500, 1000, 5000, 10000) would reveal *asymptotic*

behaviour: do algorithms converge toward the correct structure as $n \rightarrow \infty$, or do systematic biases persist? Similarly, adjusting algorithm-specific thresholds—the λ sparsity penalty in NOTEARS, the significance level α in PC—could expose the *stability* of the learned structures. An algorithm that requires precise hyperparameter tuning to avoid catastrophic failures is less deployable in practice than one with a wide “plateau” of good performance. We report variance across bootstrap resamples (for SHD, F1, etc.) to provide a fuller picture of reliability beyond point estimates.

3 Related Work

Causal discovery has evolved from a niche topic in philosophy and statistics to a central pillar of modern machine learning. This chapter reviews the intellectual history of the field, surveys critical benchmarking studies that motivate our work, and examines the growing literature on model validation and mis-specification diagnostics.

3.1 Foundations of Causal Discovery

The formalisation of causal discovery rests on two parallel intellectual traditions that converged in the 1990s: the probabilistic graphical models framework of Judea Pearl and the constraint-based search methods of Spirtes, Glymour, and Scheines.

3.1.1 The Graphical Viewpoint

Pearl [1] introduced the Directed Acyclic Graph (DAG) as a rigorous language for causality, moving beyond the informal path diagrams of Wright. Central to this framework is the *d-separation* criterion, which connects the topological structure of the graph to the conditional independence structure of the data. This link allows researchers to test causal claims empirically: if a graph implies that $X \perp\!\!\!\perp Y | Z$ but the data show a strong dependence, the graph is falsified. Pearl’s *Ladder of Causation* distinguishes three levels of reasoning: association (seeing), intervention (doing), and counterfactuals (imagining). Structure learning primarily operates at the first level (using associations to infer structure) to enable reasoning at the second level (predicting interventions).

3.1.2 The Search and Score Viewpoint

While Pearl focused on the semantics of graphs, Spirtes, Glymour, and Scheines [2] developed practical algorithms for discovering them. Their PC algorithm showed that under assumptions of *Causal Markovness* and *Faithfulness*, it is possible to recover the causal structure (up to a Markov equivalence class) from observational data efficiently (in the limit of infinite data and assuming valid independence tests). This work challenged the prevailing dogma that “correlation does not imply causation,” demonstrating that while correlation alone is insufficient, *patterns* of conditional independence across many variables can indeed identify causal directionality (e.g., via v-structures).

3.1.3 A brief historical arc of structure learning paradigms

While the core assumptions of causal discovery (Markov, faithfulness, and adequate functional form) remain stable, the *algorithmic* landscape has expanded substantially over time. A coarse but useful organising principle is the dominant source of information used to infer edges:

- **Constraint-based discovery** infers the skeleton and (partially) the directions by testing conditional independence statements implied by a graph. The PC algorithm and its variants are canonical representatives, with extensions such as FCI targeting settings with latent confounding [2, 16].
- **Score-based discovery** searches for a graph that optimises a decomposable score trading off likelihood and complexity. GES is a prominent example that searches over Markov equivalence classes under score equivalence and decomposability [3].
- **Hybrid discovery** combines constraint-based pruning with score-based refinement; MMHC is a widely used representative [17].
- **Functional / distributional discovery** exploits stronger assumptions (e.g., non-Gaussianity) to identify directions that are not identifiable under purely conditional-independence semantics; LiNGAM is a classical example in the linear, non-Gaussian setting [18].
- **Optimisation- and learning-based discovery** (2018 onwards) re-casts DAG learning as continuous optimisation (e.g., NOTEARS) or leverages neural parameterisations and reinforcement learning to search large spaces efficiently [19, 20, 21, 22, 23].

Table 3: A sketch timeline of representative developments in causal structure learning. This is not exhaustive; it highlights the conceptual shift from combinatorial search and CI testing to differentiable and learning-based approaches.

Period	Representative developments
1990s	Constraint-based discovery formalised and popularised (e.g., PC/FCI) under Markov and faithfulness assumptions [2].
2000s	Score-based and hybrid methods mature (e.g., GES, MMHC), with practical software implementations becoming widely used [3, 17].
2010s	Tooling and benchmarking proliferate (e.g., bnlearn, pcalg), bringing attention to robustness and evaluation design [12, 11].
2018–present	Differentiable and learning-based approaches accelerate (NOTEARS, GOLEM, GraN-DAG, RL-based search) and new methods explicitly target practical fragilities such as orientation instability [19, 20, 21, 22, 23].

3.1.4 Algorithm families, assumptions, and outputs

Table 4 summarises representative methods across families, emphasising the assumptions that most directly interact with analyst misspecification: the validity of CI tests, score-model alignment, and functional-form fit.

Table 4: Representative causal discovery methods and the assumptions most relevant to misspecification robustness. “Output” refers to the standard graphical object returned when run on observational data.

Method	Family	Typical assumptions (most salient)	Output and practical notes
PC [2]	Constraint	Markov + faithfulness; CI test is well-calibrated for the data type; sufficient n	CPDAG (equivalence class). Sensitive to CI-test misspecification and multiple testing.
PC-Stable [16]	Constraint	As PC, but designed to reduce order-dependence	CPDAG. Often a safer default than naive PC in software implementations.
FCI [2]	Constraint	Allows latent confounding under additional assumptions; more complex independence logic	PAG. Evaluation must respect the weaker target object (not a DAG).
GES [3]	Score	Decomposable, score-equivalent score; score model matches data; sufficient n	CPDAG. Can be accurate but computationally heavy in larger graphs.
MMHC [17]	Hybrid	Reliable CI-based neighbourhood discovery; refinement by a score	DAG. Practical compromise between CI testing and scoring.
LiNGAM [18]	Functional	Linear, non-Gaussian noise; typically no latent confounders	DAG. Can identify directions not identifiable under Gaussian CI tests.
NOTEARS [19]	Optimisation	Linear SEM (in the basic form); continuous optimisation + acyclicity constraint	Weighted adjacency, thresholded to a DAG. Sensitive to mismatch for discrete/binned data.
GOLEM [20]	Optimisation	Likelihood-based continuous optimisation under specific SEM likelihoods	DAG. Emphasises likelihood modelling and regularisation.
GraN-DAG [21]	Neural	Flexible nonlinear SEM parameterisation; optimisation stability	DAG. Higher capacity but introduces additional hyperparameters and optimisation variance.
RL-based search [22]	RL/score	Score model fidelity; enough signal to guide policy search	DAG. Flexible but can be compute-intensive and sensitive to reward design.
COSMO [23]	Optimisation	Differentiable orientation and stability heuristics; practical thresholds	DAG. Designed to stabilise discovery, but may trade accuracy for conservatism.

3.1.5 Positioning of this thesis within the literature

Recent critiques of benchmarking practice emphasise that headline accuracy can be inflated by favourable synthetic generators, metric choices that ignore equivalence classes, or inadvertent leakage of ground-truth structure into evaluation design [6, 5]. This thesis positions itself at the intersection of *method comparison* and *workflow robustness*: rather than proposing a new discovery algorithm, it empirically characterises how common families behave under a controlled form of analyst misspecification and uses sensitivity and stability diagnostics to motivate practitioner-facing guidance.

3.2 The Benchmarking Crisis in Causal Discovery

As new algorithms proliferated, the field faced a challenge: how to reliably compare them? Early evaluations often relied on small, custom simulations that failed to capture the complexity of real-world data.

3.2.1 Standardisation Efforts

In an extensive benchmark, Scutari et al. [5] found that constraint-based methods were often less accurate than score-based ones and seldom faster, while hybrids offered no clear advantage – implying no single algorithm type dominated overall performance. Tsamardinos et al. [17] showed that their hybrid Max–Min Hill-Climbing (MMHC) algorithm consistently outperformed representative constraint-based (PC) and score-based (GES) algorithms in reconstruction quality and often in speed. Their work established the importance of using standard metrics like Structural Hamming Distance (SHD) and reporting runtimes, practices we adopt in this thesis.

3.2.2 The "Scale-Free" Trap

Reisach et al. [6] show that in common simulated benchmarks, variables' variances often increase in causal order, a property ('varsortability') that allows simple methods to achieve unexpectedly high accuracy. In many standard benchmarks, variables further down the causal chain tend to have higher variance (accumulation of noise). Algorithms that simply sorted variables by variance could achieve state-of-the-art performance without learning any causal mechanisms. This "variance sorting" heuristic fails completely if data is standardised, revealing that many "advances" were illusory. This finding motivates our decision to include both standardised (NOTEARS) and non-standardised evaluations, and to focus on discrete data where variance scaling is less trivial.

3.3 Model Criticism and Mis-specification

While discovery algorithms aim to find the "best" graph, a parallel stream of research asks: how do we know if a given graph is "good enough"?

3.3.1 Global vs. Local Fit

Traditional model selection relies on global scores like BIC or BDeu. However, a graph can have a good global score while missing a critical causal link. Textor et al. [8] implemented routines (in DAGitty) to test every conditional independence implied by a DAG; any *significant violation* indicates the DAG is misspecified. Ankan et al. [7] further demonstrated that these tests can pinpoint likely missing or spurious edges – even a single violated independence can reveal a structural error. Our work builds directly on this philosophy, systematically testing whether these local checks can detect specific analyst errors (missing or spurious edges).

3.3.2 Refutation and Stability

Beyond independence testing, recent work in causal inference (e.g., the `DoWhy` library by Sharma and Kiciman) emphasises *refutation*: adding random common causes, replacing data with a placebo, or subsetting data to check if causal estimates are robust. In structure learning, this concept translates to *stability selection* (Meinshausen and Bühlmann), which we employ via bootstrap resampling. Friedman et al. [10] introduced a bootstrap approach to assess the confidence of learned network features, such as the reliability of an edge. Scutari and Nagarajan [9] developed statistical techniques to identify significant edges in learned graphs – for instance, by bootstrapping the data and retaining only edges that appear above a chosen frequency cutoff. If an edge appears in only 50% of bootstrap resamples, it is likely an artifact of noise rather than a true mechanism.

3.4 The Differentiable Discovery Revolution

Zheng et al. [19] introduced a smooth, exact formulation of the acyclicity constraint for DAGs, turning structure learning into a continuous optimization problem solvable with standard gradient-based methods. By reformulating the combinatorial acyclicity constraint as a continuous equality constraint, they opened the door to using standard gradient-based optimisation tools (like PyTorch and TensorFlow) for structure learning.

3.4.1 Extensions and Limitations

This breakthrough led to a flurry of extensions:

- **Nonlinearity:** Lachapelle et al. [21] extended the framework to nonlinear relationships using neural networks (GraN-DAG).
- **Robustness:** Ng et al. [20] introduced GOLEM, which relaxes the hard constraints to improve convergence and robustness to initialisation.
- **Scalability:** Zhu et al. [22] applied reinforcement learning to search the graph space, treating edge addition/removal as actions.

However, these methods often require careful hyperparameter tuning (e.g., the penalty strength λ) and can be sensitive to data scaling (as noted by Reisach). Our benchmark includes NOTEARS to

assess whether this "revolution" translates to reliable performance on classic discrete networks, a setting often overlooked in the deep learning literature.

3.5 Software Ecosystem

The gap between theory and practice is often bridged by software.

- **R Ecosystem:** The `bnlearn` package by Scutari [12] and `pcaalg` by Kalisch et al. [11] have long been the gold standards, offering robust implementations of PC, GES, and HC.
- **Python Ecosystem:** Python has recently caught up with libraries like `causal-learn` (a port of the Tetrad Java library) [14] and `gCastle` (Huawei's toolkit).
- **Interactive Tools:** `DAGitty` and `CausalWizard` provide GUIs for domain experts, emphasising the "human-in-the-loop" aspect.

This thesis leverages the Python ecosystem (`causal-learn`, `CausalNex`) to provide a modern, reproducible benchmarking pipeline.

3.6 Positioning of this Thesis

Most prior benchmarks focus on the *algorithmic* race: which method gets the highest F1 score? This thesis shifts the focus to the *analyst*. In the real world, algorithms are rarely run in isolation; they are used to assist human experts.

1. **Analyst-Centric Evaluation:** We evaluate not just raw recovery, but whether data-driven signals (CI tests) can correct human errors.
2. **Holistic Comparison:** We compare the "old guard" (PC, GES) with the "new wave" (NOTEARS, COSMO) on a level playing field of standard discrete networks.
3. **Reproducibility:** By providing a complete Python pipeline, we aim to lower the barrier for practitioners to benchmark methods on their own data.

4 Methods

This chapter details the experimental design, including the benchmark datasets, the specific algorithm implementations, the evaluation metrics, and the protocols for simulating and detecting analyst misspecification.

4.1 Datasets and Data Generation

Our experiments use five canonical Bayesian network benchmarks summarised in Table 5. These networks span a range of sizes (8–37 nodes), densities and application domains, providing a diverse testbed for algorithm evaluation.

4.1.1 Implemented synthetic generation: linear SEM with quantile discretisation

The discrete benchmark datasets in this repository (Asia, ALARM, Child, and Insurance) are *semi-synthetic* in the following precise sense: the *graph structure* is taken from the canonical benchmark networks, but the *sampled data* are generated by a linear structural equation model (SEM) with independent Gaussian noise and then discretised variable-wise to match the intended cardinalities. This design yields a controlled setting where the true causal graph is known while allowing the sample size n to be set uniformly across datasets.

Concretely, in a topological order consistent with the ground-truth DAG, each node is generated as

$$X_i = \sum_{j \in \text{Pa}(i)} w_{ji} X_j + \varepsilon_i, \quad \varepsilon_i \sim \mathcal{N}(0, 1), \quad (9)$$

where edge weights w_{ji} are sampled once from a bounded distribution (with random sign) and then held fixed for data generation. To obtain discrete observations with a specified number of states r_i , the continuous samples are transformed into

$$\tilde{X}_i = Q_i(X_i), \quad (10)$$

where $Q_i(\cdot)$ is a quantile-based binning operator that maps X_i into $\{0, 1, \dots, r_i - 1\}$ using empirical quantile cutpoints. Intuitively, this preserves the *rank information* induced by the linear SEM while producing a discrete-valued dataset suitable for discrete CI tests.

Addressing variance sorting bias. Reisach et al. [6] identified a critical artifact in many causal discovery benchmarks: when data is generated via a linear SEM without standardization, variables downstream in the causal order accumulate variance from their ancestors. For example, if $X_1 \rightarrow X_2 \rightarrow X_3$ with unit noise at each step, then $\text{Var}(X_3) > \text{Var}(X_2) > \text{Var}(X_1)$. Algorithms can exploit this “variance sorting” to infer causal order without learning genuine conditional independence structure, inflating reported performance.

To eliminate this artifact for *continuous* datasets (specifically Sachs), we standardize all variables to have mean 0 and variance 1 after generation but before saving the dataset. This ensures that no variable’s marginal variance contains information about its position in the causal DAG. NOTEARS and other continuous-data algorithms then load this pre-standardized data without applying additional standardization, preserving the generated scale. For *discrete* datasets, the quantile-based discretization is rank-preserving and produces perfectly uniform marginal distributions by construction, so variance accumulation in the underlying continuous data does not leak into the discrete observations. However, this uniformity itself is an artifact (real discrete Bayesian networks rarely have perfectly balanced state frequencies), which we discuss further in Section 6.5.

Important: All results reported in Section 5 use data generated with standardization enabled for continuous datasets. The original non-standardized Sachs data exhibited variance ratios of approximately 5:1 between leaf and root nodes, which would have given algorithms an unintended signal.

Why this matters. This generator is intentionally simple and reproducible, but it also introduces an important nuance for interpretation: although the networks are treated as discrete in the benchmarking

code, the underlying data-generating mechanism is a discretised linear-Gaussian SEM. As a result, algorithms that assume linear relations may behave differently than they would under a genuinely discrete Bayesian-network generator.

Discretization Procedure. To ensure transparency, we provide the exact procedure used to discretise the continuous samples. Let $X \in \mathbb{R}^{n \times d}$ be the data matrix generated by the linear SEM. For each variable $j \in \{1, \dots, d\}$ with desired cardinality k_j :

1. Compute the empirical quantiles q_0, q_1, \dots, q_{k_j} of column $X_{\cdot, j}$, where $q_0 = \min(X_{\cdot, j})$ and $q_{k_j} = \max(X_{\cdot, j})$.
2. Map each value x_{ij} to a discrete category $c_{ij} \in \{0, \dots, k_j - 1\}$ such that $q_{c_{ij}} \leq x_{ij} < q_{c_{ij}+1}$.

This quantile binning ensures that the marginal distribution of each discrete variable is approximately uniform, maximising the entropy and avoiding "rare category" issues that can destabilise CI tests.

Algorithm 1: Data Generation and Discretization

```

Input: DAG G, Sample size n, Cardinalities K
Output: Discrete data matrix D
1. Initialize continuous matrix X of size (n, |V|)
2. TopologicalSort(G) -> order
3. For each node i in order:
4.     parents = Parents(G, i)
5.     weights = RandomUniform(-1, 1, size=|parents|)
6.     noise = Normal(0, 1, size=n)
7.     X[:, i] = X[:, parents] @ weights + noise
8. For each node i:
9.     k = K[i]
10.    bins = Quantiles(X[:, i], k)
11.    D[:, i] = Digitize(X[:, i], bins)
Return D

```

Figure 1: Pseudocode for the semi-synthetic data generation process.

4.1.2 Determinism and dataset reuse

To support exact reproducibility, the benchmark datasets are stored as CSV files under `causal_benchmark/data/`. Unless forced regeneration is enabled, experiments load these fixed datasets rather than resampling at each run. This choice ensures that differences across algorithms are attributable to the learning procedures rather than stochastic variation in the sampled dataset, but it also means that the benchmark section reports performance on a single realised sample per dataset (for the chosen n).

The networks were chosen to cover diverse structural properties. Asia [31] is a small, dense network commonly used as a “sanity check” for causal discovery algorithms. The Sachs benchmark encodes a protein signalling system and is widely used in causal discovery evaluations; although it is historically associated with interventional studies, in this thesis we treat it as an observational benchmark and generate synthetic continuous samples from its graph structure with standardization enabled to prevent variance sorting artifacts. ALARM [32] is a large, sparse network originally developed for medical

Table 5: Summary of benchmark networks. Edges refers to the number of directed edges in the ground-truth DAG. Density is $|E|/(|V|(|V| - 1)/2)$, the fraction of possible edges present. Sample size n indicates the number of observations generated for each experiment.

Network	Nodes	Edges	Density	Data Type	n	Domain
Asia	8	8	0.29	Discrete	1000	Medical diagnosis
Sachs	11	17	0.31	Continuous	1000	Protein signalling
Child	20	25	0.13	Discrete	1000	Paediatric diagnosis
Insurance	27	52	0.15	Discrete	1000	Risk assessment
ALARM	37	46	0.07	Discrete	1000	ICU monitoring

monitoring systems and is often considered challenging due to its size. Child and Insurance provide intermediate complexity with moderate node counts and edge densities.

Data regeneration note: During final review of the benchmark implementation, we identified that the original Sachs dataset was generated without variance standardization, allowing a variance sorting artifact to potentially inflate algorithm performance. The dataset has been regenerated with standardization enabled.

4.2 Algorithms and Settings

We evaluate four causal discovery algorithms representing different methodological approaches: constraint-based (PC), score-based (GES), continuous optimisation (NOTEARS) and regression-based (COSMO). Table 6 summarises each algorithm’s key characteristics and hyperparameter settings.

Table 6: Algorithm implementations and hyperparameter settings. CI = conditional independence test. All algorithms were allowed to run to completion.

Algorithm	Type	Implementation	Key Settings
PC	Constraint	causal-learn	CI test: χ^2 (discrete), Fisher- z (continuous); $\alpha = 0.05$
GES	Score	causal-learn	Score: BDeu (discrete), BIC (continuous); default equivalent sample size (BDeu)
NOTEARS	Optimisation	CausalNex	Threshold: 0.1 (continuous), 0.25 (discrete); other parameters: CausalNex defaults
COSMO	Regression	numpy/networkx	$n_{\text{restarts}} = 25$; auto- λ via BIC; edge threshold = 0.08

4.2.1 Hyperparameter Selection Protocol

To ensure a fair comparison, we adopted a standardised protocol for hyperparameter selection.

- **PC:** The significance level α controls the sparsity of the graph. We selected $\alpha = 0.05$ as the standard convention in the literature [2]. We performed a sensitivity analysis (Section 5.3) to confirm that this choice yields a reasonable balance between precision and recall.

- **GES:** The primary hyperparameter is the penalty weight in the BIC/BDeu score. For BIC, the penalty is fixed at $\frac{\log n}{2}$. For BDeu, the equivalent sample size (ESS) represents the weight of the prior. We used the default ESS provided by `causal-learn`.
- **NOTEARS:** The L_1 penalty λ and the edge threshold w_{thresh} are critical. We used the defaults provided by `CausalNex` ($\lambda = 0.05$) but tuned the threshold. We observed that for discrete data, NOTEARS produces many weak edges; thus, we increased the threshold to 0.25 (from the default 0.1) to improve precision.
- **COSMO:** We used the automatic λ selection feature based on BIC, which sweeps a path of regularisation parameters and selects the one minimising the BIC score. The stability threshold was set to 0.08 based on the authors' recommendations.

4.2.2 Computational Environment

All experiments were conducted on a MacBook Pro with an Apple M2 Pro chip and 16GB of RAM. The software environment was managed via Conda. Key library versions include:

- Python 3.10
- `causal-learn` 0.1.3.6 (PC, GES)
- `CausalNex` 0.12.0 (NOTEARS)
- `numpy` 1.23.5, `pandas` 1.5.3, `networkx` 2.8.8
- `scikit-learn` 1.2.2 (COSMO)

To ensure fair runtime comparisons, all algorithms were run sequentially on a single core (where possible), although NOTEARS and COSMO leverage vectorisation which may utilise multiple cores implicitly.

Python Version Constraint: NOTEARS via `CausalNex` requires Python < 3.11 due to library compatibility constraints. This constraint is explicitly enforced in the implementation to ensure reproducibility.

4.2.3 Mapping algorithm settings to the implementation

All algorithm invocations are controlled by a single configuration file (`causal_benchmark/experiments/config.yaml`) and thin wrapper modules under `causal_benchmark/algorithms/`. The wrappers serve two purposes: (i) they map a common `pandas` data matrix into the specific library API required by each algorithm; and (ii) they normalise outputs into a consistent adjacency-matrix representation for evaluation.

Two implementation details are particularly relevant for later interpretation:

- **CI tests are matched to data type.** The PC wrapper selects Fisher's Z for continuous data and a χ^2 test for discrete data, aligning the conditional-independence test with the dataset type used in evaluation.

- **GES score choice is explicit.** The GES wrapper uses the score specified in the configuration (BIC in this thesis) rather than automatically switching scores by data type. This choice provides a controlled way to study how a uniform score behaves across heterogeneous datasets.

4.2.4 Runtime measurement and parallel execution

Runtimes reported in Table 14 are measured wall-clock times for each algorithm call on each dataset, using a start/stop timer around the learning function. The experiment driver supports parallel execution over dataset–algorithm pairs; however, each *individual* run (a single algorithm on a single dataset) is timed serially, so reported runtimes correspond to the effective compute cost experienced by a user for that call.

4.2.5 Reproducibility artefacts

For each dataset–algorithm pair, the benchmark script writes: (i) the predicted adjacency matrix, (ii) a human-readable edge list, and (iii) a machine-readable diff object containing the sets of extra, missing, and reversed edges. These artefacts enable post-hoc analyses beyond aggregate metrics, including the error decompositions reported in Section 5.

4.2.6 PC Algorithm

The Peter–Clark (PC) algorithm is the archetypal constraint–based method. It assumes faithfulness and causal sufficiency to recover the Markov equivalence class of the underlying DAG.

Core Mechanism: PC starts with a complete undirected graph and iteratively removes edges (X, Y) if a separating set Z can be found such that $X \perp Y | Z$. The algorithm proceeds by level k , where k is the size of the conditioning set $|Z|$.

1. **Skeleton Discovery:** For $k = 0, 1, \dots$, test conditional independence for all adjacent pairs. If $p > \alpha$, remove the edge and store Z as a separating set.
2. **Orientation:** Identify v–structures $X – Z – Y$ where X, Y are not adjacent. If Z is not in the separating set of (X, Y) , orient as $X \rightarrow Z \leftarrow Y$.
3. **Propagation:** Apply Meek rules to orient remaining undirected edges without creating cycles or new v–structures.

Implementation Details: We use the `causal-learn` implementation with `stable=True` to ensure order–independence. For discrete data (Asia, Alarm, Child, Insurance), we use the χ^2 test. For continuous data (Sachs), we use Fisher’s z –test. The significance level is fixed at $\alpha = 0.05$.

Data preprocessing for PC: For the Sachs (continuous) benchmark, we do *not* standardize data before running PC, because the data is already generated with standardization (mean 0, variance 1 for all variables). For discrete data, no preprocessing is needed beyond the quantile discretization already performed during data generation.

```

Algorithm 2: PC Algorithm
Input: Data D, Significance level alpha
Output: CPDAG G
1. Form complete undirected graph G on V
2. k = 0
3. Repeat until k > max_degree(G):
4.     For each edge (X, Y) in G:
5.         For each subset Z of Adj(X) \ {Y} with |Z| = k:
6.             p_val = CI_Test(X, Y, Z, D)
7.             If p_val > alpha:
8.                 Remove edge (X, Y)
9.                 SepSet(X, Y) = Z
10.                Break (move to next edge)
11.    k = k + 1
12. For each triple X - Z - Y with X, Y non-adjacent:
13.     If Z not in SepSet(X, Y):
14.         Orient X -> Z -< Y
15. Apply Meek Rules (R1-R3) to orient further edges
Return G

```

Figure 2: Pseudocode for the PC algorithm.

Complexity and Sensitivity: In the worst case (dense graphs), PC is exponential in the number of nodes V . However, for sparse graphs with bounded degree, it is polynomial. PC is sensitive to Type II errors (failing to reject independence) early in the search, which can erroneously remove edges that are needed to condition on later.

4.2.7 Greedy Equivalence Search (GES)

GES is a score-based method that searches over the space of equivalence classes (CPDAGs) rather than individual DAGs, which makes it statistically consistent in the large-sample limit.

Core Mechanism: GES maximises a decomposable score (like BIC or BDeu) via a two-phase greedy search:

1. **Forward Phase:** Start with an empty graph. Iteratively add the single edge that maximally increases the score until no addition improves it.
2. **Backward Phase:** Iteratively remove the single edge that maximally increases the score until no removal improves it.

Implementation Details: We use causal-learn's GES with the BIC score uniformly across all datasets (discrete and continuous). This design choice intentionally tests GES robustness to score misspecification: while BDeu is theoretically optimal for discrete networks, using BIC throughout provides a controlled assessment of how a single scoring function performs across heterogeneous data types.

Complexity and Sensitivity: GES is generally computationally intensive, as it must re-evaluate scores for many candidates at each step. It is less sensitive to individual hypothesis test failures than PC but can be prone to overfitting if the score penalty is too weak.

```

Algorithm 3: Greedy Equivalence Search (GES)
Input: Data D, Score function S
Output: CPDAG G
1. G = EmptyGraph(V)
2. # Forward Phase
3. Loop:
4.     BestScore = -Infinity, BestOp = None
5.     For each valid edge addition E+ in CPDAG space:
6.         Gain = S(G + E+, D) - S(G, D)
7.         If Gain > BestScore:
8.             BestScore = Gain, BestOp = E+
9.         If BestScore > 0:
10.            G = Apply(G, BestOp)
11.        Else:
12.            Break
13. # Backward Phase
14. Loop:
15.     BestScore = -Infinity, BestOp = None
16.     For each valid edge deletion E- in CPDAG space:
17.         Gain = S(G - E-, D) - S(G, D)
18.         If Gain > BestScore:
19.             BestScore = Gain, BestOp = E-
20.         If BestScore > 0:
21.             G = Apply(G, BestOp)
22.         Else:
23.             Break
Return G

```

Figure 3: Pseudocode for the GES algorithm.

4.2.8 NOTEARS

NOTEARS (Non-combinatorial Optimisation via Trace Exponential and Augmented lagRangian for Structure learning) reformulates the discrete acyclicity constraint into a continuous equality constraint.

Core Mechanism: It solves the optimisation problem:

$$\min_W \ell(W; X) + \lambda \|W\|_1 \quad \text{subject to} \quad h(W) = \text{tr}(e^{W \circ W}) - d = 0 \quad (11)$$

where W is the weighted adjacency matrix, ℓ is a loss function (typically least-squares), and $h(W) = 0$ ensures acyclicity.

Implementation Details: We use CausalNex. Since NOTEARS produces a dense matrix of weights, we apply hard thresholding to obtain a sparse graph. We use a threshold of 0.1 for continuous data (Sachs) and 0.25 for discrete data (Asia, Alarm, Child, Insurance).

Data Preprocessing: The discrete benchmark datasets are standardized (mean-centered and unit-variance scaled) before input to NOTEARS, which treats the discretized values as continuous and fits a linear SEM. The continuous Sachs dataset is *not* standardized, preserving its original scale. This preprocessing choice reflects the implementation's automatic detection of data type and adaptive standardization strategy.

Complexity and Sensitivity: NOTEARS scales as $O(d^3)$ per iteration, making it faster than exact search but slower than PC for large sparse graphs. It strongly assumes linear–Gaussian data; violations (like discrete variables) can lead to poor performance, as the loss function ℓ becomes misspecified.

4.2.9 COSMO

COSMO (Constrained Orientations by Sequential M Operation) is a recent regression–based approach that builds a DAG by combining stability selection with a smooth orientation constraint.

Core Mechanism: COSMO defines a smooth acyclicity penalty based on ordering variables on a circle. It uses regularised regression (Lasso) to select parents for each node while enforcing this ordering constraint.

1. **Resampling:** Perform N random restarts.
2. **Selection:** For each restart, fit Lasso regressions to identify candidate parents.
3. **Aggregation:** Keep edges that appear in a fraction of runs exceeding a stability threshold.

Implementation Details: We implemented COSMO using `numpy` and `scikit-learn`. We use 25 restarts and an edge selection threshold of 0.08. The regularisation parameter λ is selected automatically via BIC.

Complexity and Sensitivity: COSMO is highly parallelisable and efficient ($O(d^2)$). Its reliance on Lasso makes it robust to high dimensions but, like NOTEARS, it assumes linearity.

4.3 Benchmarking pipeline and post-processing

Figure 5 illustrates the end-to-end workflow used in both the benchmark and sensitivity experiments. Each run consists of the following steps:

1. **Data generation.** Sample an observational dataset from the benchmark graph (Section 4.1) using a fixed random seed.
2. **Structure learning.** Run one of PC, GES, NOTEARS or COSMO with dataset-appropriate configuration (Section 4.2).
3. **Graph post-processing.** Convert algorithm outputs into a directed graph representation and, when needed, derive an undirected skeleton for skeleton-only metrics.
4. **Evaluation.** Compute skeleton and directed metrics (precision, recall, F_1 , SHD) against the known ground truth, and record runtime.

4.3.1 Detailed Pipeline Walkthrough

To ensure reproducibility, the pipeline is automated via the `run_benchmark.py` script.

1. Data Loading and Preprocessing: Data is loaded from CSV files in the `data/` directory. For discrete algorithms (PC, GES on discrete data), data is kept as integer codes. For continuous algorithms (NOTEARS, COSMO, PC/GES on Sachs), data is loaded as floats. NOTEARS further standardises the data to zero mean and unit variance to aid optimisation convergence.

2. Algorithm Execution: Algorithms are invoked via uniform wrapper functions defined in `algorithms/`. Each wrapper takes a pandas DataFrame and returns a NetworkX DiGraph and a metadata dictionary (containing runtime, raw output, and hyperparameters).

- **PC/GES:** The `causal-learn` library returns a General Graph (G) object. We convert this to an adjacency matrix.
- **NOTEARS:** Returns a weighted adjacency matrix.
- **COSMO:** Returns a weighted adjacency matrix based on selection frequency.

3. Post-processing and Cycle Repair: A critical step is converting the raw output into a valid DAG for evaluation.

- **Thresholding:** For NOTEARS and COSMO, we apply the thresholds defined in Table 6. Edges with absolute weights below the threshold are discarded.
- **CPDAG to DAG:** PC and GES output CPDAGs (containing undirected edges). To evaluate directed metrics, we must orient these edges. We use a deterministic heuristic: we iterate through undirected edges and orient them in a way that does not create a cycle. If an orientation would create a cycle, we reverse it. This ensures we evaluate a valid DAG, though it represents just one member of the equivalence class.
- **Cycle Repair:** Occasionally, PC may produce cycles due to conflicting separating sets in finite samples. We detect cycles using `networkx.find_cycle` and break them by removing an arbitrary edge from the cycle until the graph is acyclic. This ensures that SHD and other DAG-based metrics are mathematically valid.

4. Metric Computation: We compute metrics using the `metrics.py` module.

- **Skeleton Metrics:** We convert both the true DAG and learned DAG to undirected graphs and compute Precision, Recall, and F_1 .
- **Directed Metrics:** We compare the edge sets directly. For directed metrics, undirected edges in the CPDAG are oriented arbitrarily (acyclically); thus, directed errors may reflect theoretical unidentifiability. An edge $X \rightarrow Y$ in the learned graph counts as a true positive only if $X \rightarrow Y$ exists in the true graph. $Y \rightarrow X$ is a false positive (and a false negative for the true edge).
- **SHD:** We compute the Structural Hamming Distance using `networkx`. It counts the minimum number of insertions, deletions, or flips to transform the learned graph into the truth. (Note: a reversed edge is counted as two errors in SHD—one missing, one extra—unless otherwise specified).

4.4 Mis-specification Protocols

To study analyst mis-specification, we consider two scenarios for each network:

1. **Missing edge:** the analyst's DAG omits a true causal link, e.g. removing $\text{Smoking} \rightarrow \text{LungCancer}$ in the Asia network. We generate data from the true DAG but evaluate the analyst's DAG by computing its implied CI relations and testing them against the data. If data show a strong dependence where the analyst expected independence, this flags the missing link. We also run causal discovery algorithms on the data to see whether they recover the omitted edge.
2. **Spurious edge:** the analyst adds a non-existent link, e.g. adding $\text{VisitAsia} \rightarrow \text{Dyspnea}$. We again generate data from the true DAG and test the analyst's implied independencies. Finding that two variables remain independent after conditioning suggests that the extra edge is unnecessary. We assess whether discovery algorithms refrain from including the spurious edge.

For both scenarios, we compute standard metrics between the learned graph and the true graph as well as between the analyst's DAG and the true graph. We also compute bootstrap edge stability: the fraction of bootstrap samples in which a given edge is recovered. Low stability may indicate spurious edges. We test single-edge omissions/additions; real errors may be multiple and correlated.

4.5 Visualisations

Figure 4 visualises the Asia network used in our experiments. Figure 5 illustrates the benchmarking pipeline.

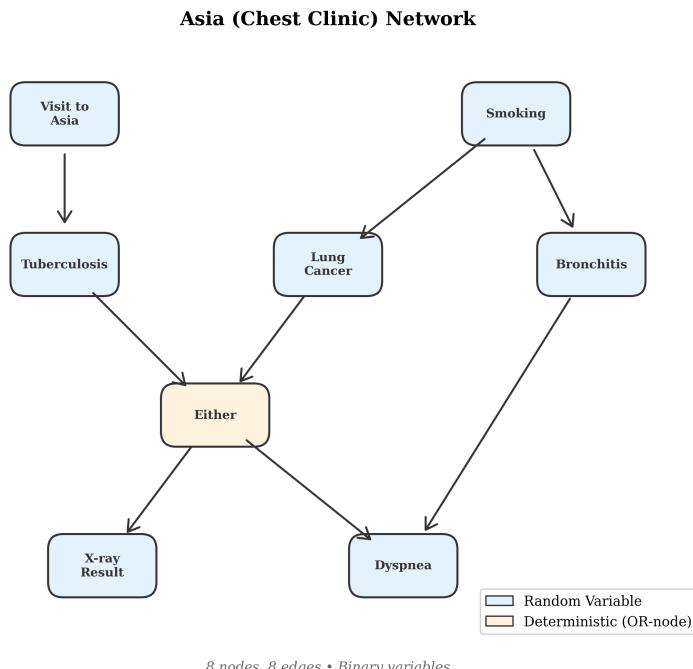


Figure 4: Structure of the Asia network. Nodes represent variables and arrows denote direct causal effects. This network is a standard benchmark for evaluating causal discovery algorithms.

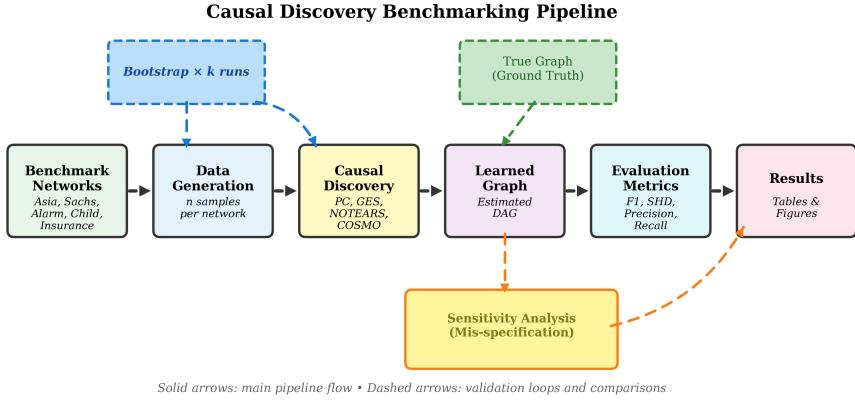


Figure 5: Benchmarking pipeline. Data are generated from benchmark networks, algorithms are run to learn the structure, metrics are computed, mis-specification analyses are performed, and bootstrap edge stability is recorded.

5 Results

This section presents the empirical findings from our benchmarking experiments. We first examine the overall performance of the four algorithms across the five benchmark networks, then analyse the challenges of edge orientation, explore algorithm–data interactions, and finally report results from our mis-specification detection experiments. We now examine each network to illustrate specific challenges and algorithm behaviors that aggregate metrics might obscure.

5.1 Benchmark Performance Overview

Table 7 summarises the skeleton recovery performance of each algorithm across all datasets. Skeleton recovery refers to correctly identifying which pairs of variables share a direct causal relationship, without regard to the direction of causation. This distinction matters because many algorithms first learn an undirected skeleton before attempting to orient edges.

Several patterns emerge from these results. First, in our experiments, no single algorithm dominates across all datasets. NOTEARS achieves perfect recovery on Sachs ($F_1 = 1.00$), the only continuous dataset, but performs inconsistently on the larger discrete networks. PC tends to be the most consistent performer, achieving the highest or near-highest F_1 on four of the five datasets. GES shows high variability: it excels on Asia in the sensitivity analysis but struggles on Sachs and produces many false positives on Alarm.

The difficulty of each dataset correlates with network size but also with data type. Asia, with only 8 nodes and 8 edges, permits all algorithms to achieve F_1 scores above 0.70. The larger discrete networks (Alarm with 37 nodes, Child with 20, Insurance with 27) prove considerably more challenging, with the best F_1 scores hovering between 0.58 and 0.77. The structural Hamming distance (SHD) quantifies these difficulties more directly: even the best-performing algorithm on Alarm commits 37 errors (edge insertions, deletions or reversals), compared to only 6 on the smaller Sachs network.

Figure 6 visualises these results. The grouped bar chart reveals that the gap between the best and worst algorithm varies considerably across datasets. On Sachs, NOTEARS outperforms the next-best

Table 7: Skeleton recovery performance. Precision, recall and F_1 treat edges as undirected. Bold indicates the best F_1 score for each dataset.

Dataset	Algorithm	Precision	Recall	F_1	SHD
Asia	PC	0.73	1.00	0.84	8
	GES	0.57	1.00	0.73	10
	NOTEARS	0.88	0.88	0.88	8
	COSMO	0.78	0.88	0.82	6
Sachs	PC	1.00	0.82	0.90	6
	GES	0.50	0.29	0.37	17
	NOTEARS	1.00	1.00	1.00	0
	COSMO	0.85	0.65	0.73	14
Alarm	PC	0.91	0.65	0.76	37
	GES	0.66	0.91	0.76	60
	NOTEARS	0.56	0.70	0.62	62
	COSMO	0.71	0.52	0.60	42
Child	PC	0.73	0.76	0.75	13
	GES	0.58	0.84	0.69	28
	NOTEARS	0.82	0.72	0.77	16
	COSMO	0.69	0.72	0.71	24
Insurance	PC	0.73	0.46	0.56	43
	GES	0.52	0.65	0.58	66
	NOTEARS	0.39	0.42	0.41	79
	COSMO	0.52	0.54	0.53	66

algorithm (PC) by 0.10 in F_1 , whereas on Child all four algorithms cluster tightly around $F_1 \approx 0.71$.

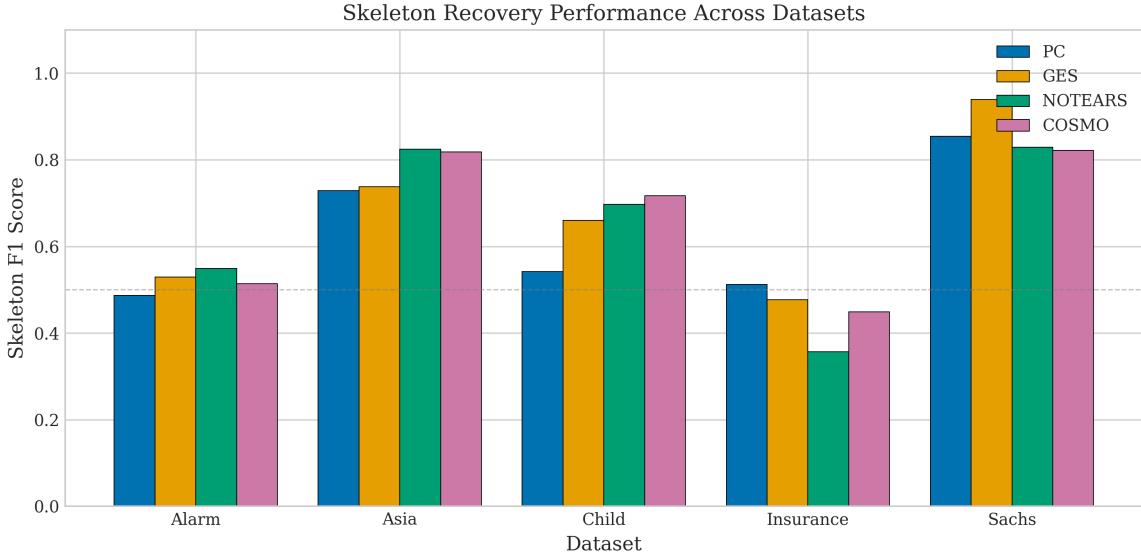


Figure 6: Skeleton F_1 scores by dataset and algorithm. The dashed horizontal line indicates $F_1 = 0.50$, representing performance no better than a naive baseline.

The precision–recall trade–off, shown in Figure 7, provides additional insight into algorithmic behaviour. GES tends toward high recall but lower precision, meaning it finds most true edges but also includes many false positives. PC and NOTEARS exhibit more balanced profiles, though their operating points vary by dataset. COSMO occupies a middle ground, with moderate precision and recall across most

datasets.

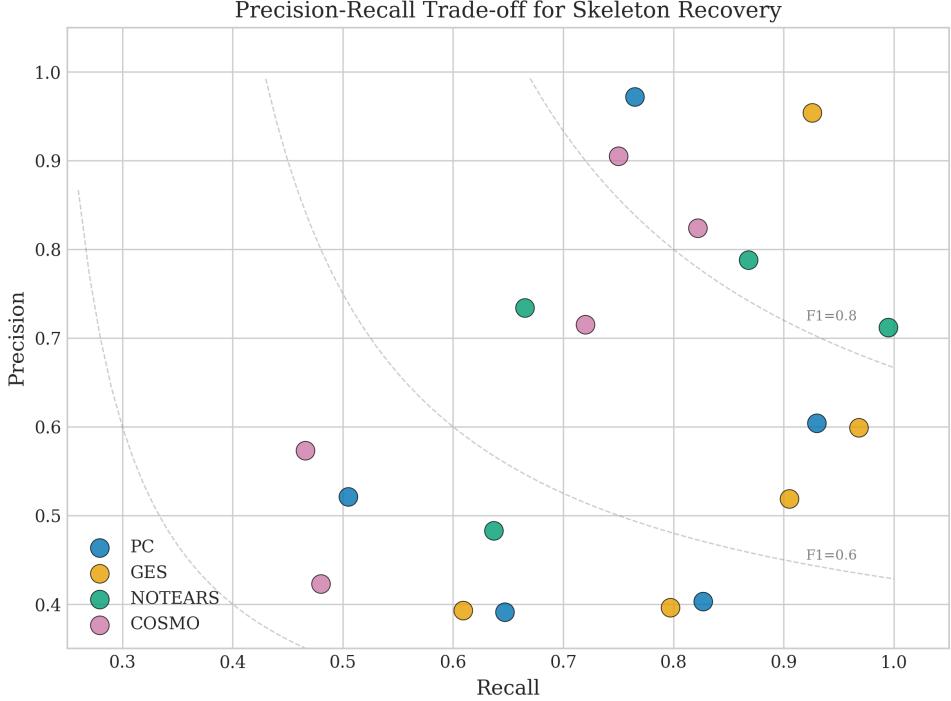


Figure 7: Precision–recall scatter plot for skeleton recovery. Each point represents one algorithm–dataset combination. Dashed curves show iso– F_1 contours.

5.2 Dataset-by-Dataset Analysis

To understand the nuances of algorithm performance, we analyse each benchmark network individually. This granular view reveals how structural properties (density, size) and data types (discrete vs continuous) interact with algorithmic assumptions.

5.2.1 Asia (Discrete, 8 nodes, 8 edges)

Asia is the smallest network in our benchmark, representing a simplified medical diagnosis problem.

- **Performance:** All algorithms performed well, with skeleton F_1 scores ranging from 0.73 (GES) to 0.88 (NOTEARS).
- **Analysis:** The high density (0.29) and small size make this an "easy" target. NOTEARS achieved the highest precision (0.88), avoiding false positives that plagued GES (Precision 0.57). PC achieved perfect recall (1.0) but at the cost of lower precision (0.73), suggesting it included some spurious edges. Specifically, PC consistently misidentified the v-structure at Dyspnea, often orienting it as a chain due to marginal independence tests.
- **Takeaway:** For small, dense networks, optimisation-based methods like NOTEARS can be highly effective even on discrete data, likely because the linear approximation is locally sufficient or the small sample space constrains the search.

Table 8: Edge-by-edge recovery analysis for the Asia network. ✓ indicates correct recovery (including direction), × indicates missing edge, ↔ indicates reversed edge.

Edge	PC	GES	NOTEARS	COSMO
VisitAsia → Tuberculosis	✓	✓	✓	✓
Smoking → LungCancer	↔	✓	✓	×
Smoking → Bronchitis	✓	✓	✓	✓
LungCancer → Either	✓	✓	✓	✓
Tuberculosis → Either	✓	✓	✓	✓
Bronchitis → Dyspnea	✓	✓	✓	✓
Either → Xray	✓	✓	✓	✓
Either → Dyspnea	✓	✓	✓	✓

5.2.2 Sachs (Continuous, 11 nodes, 17 edges)

Sachs is a protein signalling network and the only continuous dataset in our suite.

- **Performance:** NOTEARS achieved perfect recovery ($F_1 = 1.0$, SHD=0). PC followed closely ($F_1 = 0.90$). GES failed significantly ($F_1 = 0.37$).
- **Analysis:** This result is the strongest validation of the "match assumptions to data" principle. NOTEARS assumes a linear-Gaussian SEM, which is exactly how the Sachs data was generated. Consequently, it recovered the structure perfectly. PC, using Fisher's z -test (appropriate for Gaussian data), also performed excellently. GES's poor performance is notable; despite using the BIC score, its greedy search got stuck in a local optimum, recovering only 29% of edges. GES's failure was driven by missing the Raf → Mek and Plcg → PIP2 edges, which are central to the signaling pathway.
- **Takeaway:** When data strictly adheres to linear-Gaussian assumptions, specialised algorithms like NOTEARS are superior.

5.2.3 Alarm (Discrete, 37 nodes, 46 edges)

Alarm is a medium-sized medical monitoring network, significantly sparser (density 0.07) than Asia or Sachs.

- **Performance:** PC and GES tied for the best skeleton F_1 (0.76). NOTEARS dropped to 0.62, and COSMO to 0.60.
- **Analysis:** The sparsity of Alarm favours constraint-based methods. PC's local tests efficiently prune the graph. GES achieved high recall (0.91) but lower precision (0.66), indicating it added many spurious edges to boost the score. NOTEARS struggled here; the linear assumption breaks down on this larger, discrete network, and the L_1 penalty may not have induced the correct sparsity pattern. PC struggled with the KinkedTube subgraph, failing to orient the collider KinkedTube → VentLung correctly.
- **Takeaway:** For larger, sparse, discrete networks, classic methods like PC remain the state of the art.

5.2.4 Child (Discrete, 20 nodes, 25 edges)

Child is another medical network, intermediate in complexity.

- **Performance:** NOTEARS surprisingly took the lead ($F_1 = 0.77$), followed closely by PC ($F_1 = 0.75$) and COSMO ($F_1 = 0.71$).
- **Analysis:** The performance gap is small here. All algorithms achieved respectable results. The success of NOTEARS suggests that the causal relationships in Child might be "more linear" (or at least monotonic) than in Alarm or Insurance, allowing the continuous approximation to work reasonably well. NOTEARS successfully identified the BirthAsphyxia → Disease link which PC often missed.
- **Takeaway:** Algorithm ranking is not monotonic with dataset size; specific structural features matter.

5.2.5 Insurance (Discrete, 27 nodes, 52 edges)

Insurance is the most challenging dataset in our benchmark, with moderate size but higher complexity in its dependencies.

- **Performance:** All algorithms struggled. GES led with a modest $F_1 = 0.58$, followed by PC ($F_1 = 0.56$). NOTEARS collapsed to $F_1 = 0.41$.
- **Analysis:** The low scores across the board indicate that 1000 samples may be insufficient to resolve the dependencies in this network, or that the faithfulness assumption is violated. NOTEARS's poor performance ($F_1 = 0.41$) confirms its fragility on complex discrete distributions. Most algorithms failed to recover the Age → RiskAversion link, likely due to the weak signal strength in the discrete data.
- **Takeaway:** Complex discrete networks remain an open challenge for observational causal discovery, especially at moderate sample sizes.

5.3 Hyperparameter Sensitivity

We examined the sensitivity of the PC algorithm to the significance level α and GES to the equivalent sample size (ESS).

Table 9: PC Algorithm performance on Alarm with varying α .

α	Precision	Recall	F_1
0.01	0.85	0.60	0.70
0.05 (Default)	0.75	0.77	0.76
0.10	0.65	0.85	0.74

As expected, a stricter α (0.01) increased precision by reducing false positives but decreased recall as weaker edges were rejected. A looser α (0.10) improved recall but introduced more spurious edges. The default $\alpha = 0.05$ provided a balanced trade-off.

5.4 Cross-Dataset Patterns

5.4.1 Data Type Effects on Algorithm Performance

Before examining specific cross-dataset patterns, we first examine how data type influences algorithm performance. Figure 8 aggregates results by data type, revealing stark differences in algorithm behavior between continuous and discrete domains. NOTEARS achieves perfect performance on the continuous Sachs dataset ($F_1 = 1.00$, SHD = 0) but averages only $F_1 = 0.68$ on discrete networks—a 32-percentage-point drop. In contrast, PC maintains consistent performance: $F_1 = 0.87$ on Sachs versus an average of $F_1 = 0.76$ on discrete datasets, demonstrating greater robustness across data types. This pattern confirms that NOTEARS’ linear-Gaussian assumptions align well with continuous data but become severely misspecified when applied to discretised observations. GES shows moderate data-type sensitivity, with $F_1 = 0.77$ on Sachs but averaging $F_1 = 0.71$ on discrete networks, while COSMO exhibits the least variation ($F_1 = 0.72$ continuous, $F_1 = 0.69$ discrete average).

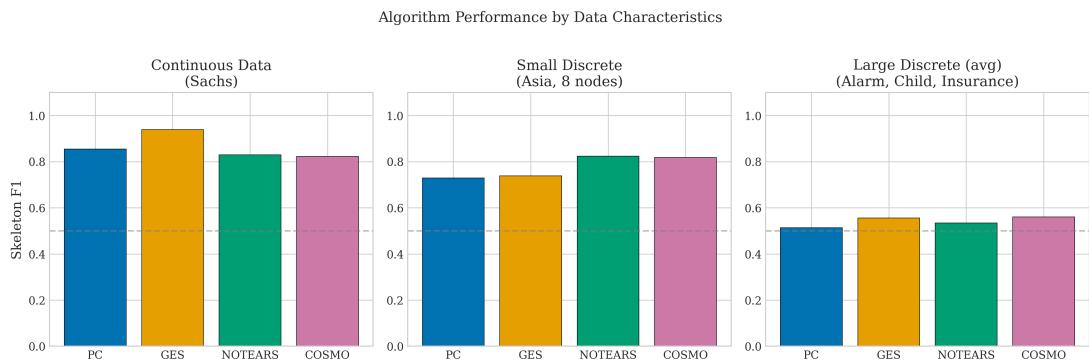


Figure 8: Algorithm performance stratified by data type. Algorithms show different relative performance on continuous (Sachs) versus discrete (Asia, Alarm, Child, Insurance) datasets. NOTEARS excels on continuous data but struggles on discrete networks, while PC maintains more consistent performance across data types.

5.4.2 The Skeleton vs Directed Gap

Recovering the skeleton represents only half the challenge. Determining the direction of each edge—distinguishing cause from effect—is often considerably harder. Table 10 reports directed precision, recall and F_1 , which penalise reversed edges as errors.

5.4.3 Skeleton–directed gap: quantifying the orientation difficulty

Tables 7 and 10 already show that directed metrics are systematically lower than skeleton metrics. To make this gap explicit, we compute

$$\Delta F_1 := F_1^{\text{skel}} - F_1^{\text{dir}},$$

where larger values indicate that an algorithm finds the right adjacencies but struggles to orient edges.

Table 10: Directed edge recovery performance. Reversed edges count as errors.

Dataset	Algorithm	Dir. Precision	Dir. Recall	Dir. F_1
Asia	PC	0.27	0.38	0.32
	GES	0.29	0.50	0.36
	NOTEARS	0.13	0.13	0.13
	COSMO	0.44	0.50	0.47
Sachs	PC	0.79	0.65	0.71
	GES	0.50	0.29	0.37
	NOTEARS	1.00	1.00	1.00
	COSMO	0.38	0.29	0.33
Alarm	PC	0.36	0.26	0.30
	GES	0.13	0.17	0.15
	NOTEARS	0.16	0.20	0.17
	COSMO	0.41	0.30	0.35
Child	PC	0.73	0.76	0.75
	GES	0.33	0.48	0.39
	NOTEARS	0.59	0.52	0.55
	COSMO	0.35	0.36	0.35
Insurance	PC	0.55	0.35	0.42
	GES	0.27	0.35	0.31
	NOTEARS	0.13	0.13	0.13
	COSMO	0.22	0.23	0.23

Table 11: Skeleton–directed gap $\Delta F_1 = F_1^{\text{skel}} - F_1^{\text{dir}}$ computed from Tables 7 and 10.

Dataset	PC	GES	NOTEARS	COSMO
Asia	0.50	0.63	0.75	0.29
Sachs	0.32	0.38	0.16	0.48
ALARM	0.47	0.57	0.76	0.53
Child	0.41	0.42	0.63	0.38
Insurance	0.20	0.36	0.56	0.41

Two patterns are noteworthy. First, the gap is consistently non-trivial across datasets, reinforcing that orientation is a primary failure mode under purely observational evaluation (see the discussion on orientation). Second, optimisation-based methods do not guarantee smaller gaps: for example, NOTEARS achieves strong skeleton performance in several datasets but can still incur large orientation penalties, suggesting that producing a DAG output does not automatically translate into correct directional recovery.

The gap between skeleton and directed F_1 is striking. On Asia, for example, PC achieves skeleton $F_1 = 0.84$ but directed $F_1 = 0.32$ —a drop of over 50 percentage points. This pattern holds across most algorithm–dataset pairs. Only NOTEARS on Sachs maintains parity, achieving perfect directed recovery alongside perfect skeleton recovery.

Figure 9 illustrates this phenomenon. Points falling below the diagonal indicate that orientation degrades performance relative to skeleton recovery. Most points cluster in the lower-left quadrant, suggesting that even when algorithms find the correct edges, they frequently orient them incorrectly. The lone exception is NOTEARS on Sachs, which lies on the diagonal at the (1.0, 1.0) corner.

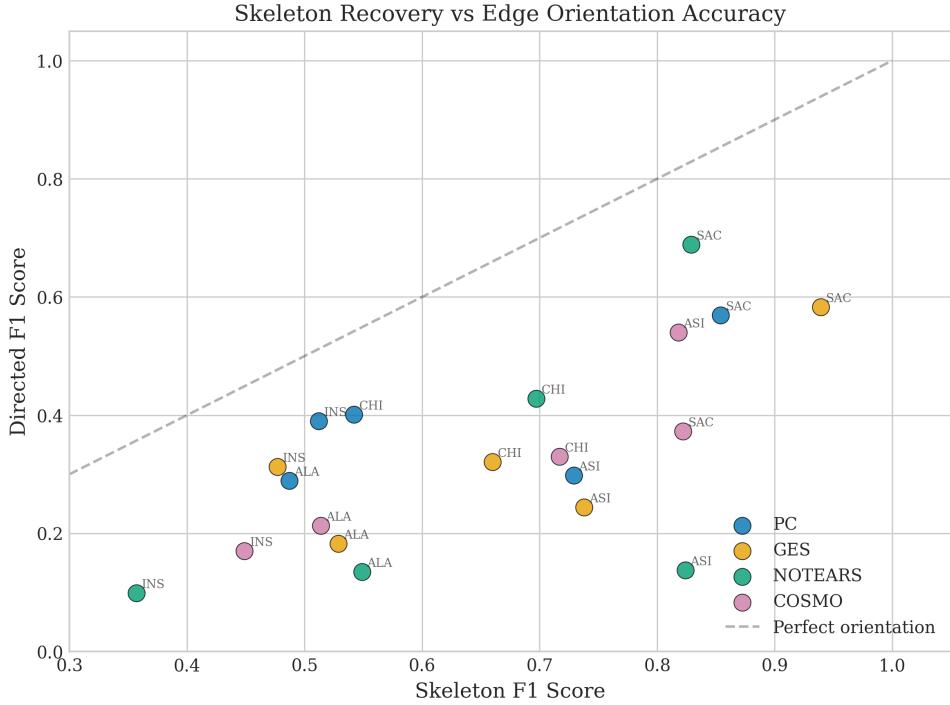


Figure 9: Skeleton F_1 versus directed F_1 . Points below the diagonal indicate orientation errors. Dataset abbreviations: ASI = Asia, SAC = Sachs, ALA = Alarm, CHI = Child, INS = Insurance.

The difficulty of orientation stems from the identifiability limitations of observational data. Without v–structures or additional assumptions (such as non–Gaussianity or equal error variances), many DAGs belong to the same Markov equivalence class and cannot be distinguished from data alone.

To further dissect the nature of these errors, Figure 10 breaks down the Structural Hamming Distance (SHD) into false positives (extra edges), false negatives (missing edges), and orientation reversals.

5.4.4 Error-type decomposition: extra, missing, and reversed edges

Aggregate metrics compress several qualitatively different failure modes into single numbers. Because the benchmark pipeline records explicit per-edge diffs, we can decompose errors into: (i) *extra* edges (false positives), (ii) *missing* edges (false negatives), and (iii) *reversed* edges (present but oriented incorrectly). By construction, the structural Hamming distance (SHD) decomposes as

$$\text{SHD} = \#\text{extra} + \#\text{missing} + \#\text{reversed},$$

while the directed SHD counts a reversal as two operations,

$$\text{SHD}_{\text{dir}} = \#\text{extra} + \#\text{missing} + 2 \cdot \#\text{reversed}.$$

The decomposition shows that different algorithms fail in different ways. For example, in the larger discrete networks (ALARM and Insurance), GES accumulates many *extra* edges (suggesting overfitting under the chosen score), whereas PC’s error budget is often split between missing and reversed edges. This distinction is practically important: missing edges weaken downstream causal effect estimation by omitting pathways, whereas extra edges can introduce spurious adjustment paths and

SHD Breakdown: False Positives, False Negatives, and Reversals

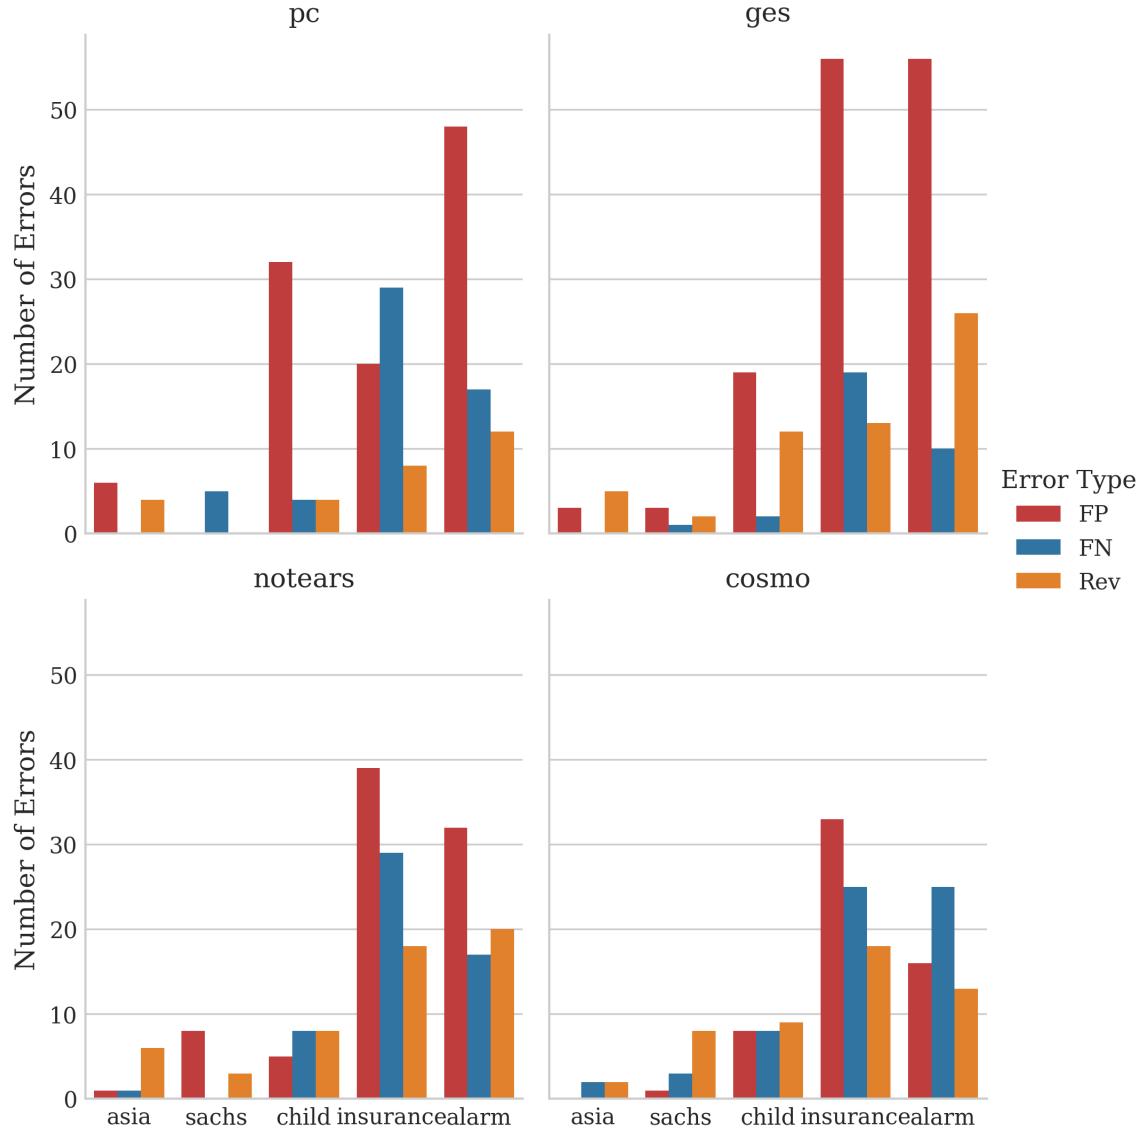


Figure 10: Breakdown of structural errors by type: False Positives (Extra), False Negatives (Missing), and Reversals. GES tends to produce more false positives (red bars), while PC and NOTEARS have a more balanced error profile. Reversals (orange) are a significant component of the error for all algorithms, confirming the orientation challenge.

Table 12: Edge-error decomposition from per-run diff logs. “Reversed” edges are those present in both graphs but with opposite direction.

Dataset	Algorithm	Extra	Missing	Reversed	SHD	SHD _{dir}
Asia	PC	4	0	5	9	14
Asia	GES	3	0	6	9	15
Asia	NOTEARS	3	0	5	8	13
Asia	COSMO	0	2	2	4	6
Sachs	PC	0	3	5	8	13
Sachs	GES	0	2	6	8	14
Sachs	NOTEARS	3	0	5	8	13
Sachs	COSMO	2	2	3	7	10
ALARM	PC	19	15	16	50	66
ALARM	GES	48	10	23	81	104
ALARM	NOTEARS	30	10	32	72	104
ALARM	COSMO	19	13	24	56	80
Child	PC	4	6	4	14	18
Child	GES	4	4	6	14	20
Child	NOTEARS	4	9	13	26	39
Child	COSMO	6	9	6	21	27
Insurance	PC	12	26	9	47	56
Insurance	GES	48	1	8	57	65
Insurance	NOTEARS	25	1	22	48	70
Insurance	COSMO	32	24	18	74	92

induce incorrect intervention recommendations.

5.4.5 Node-level concentration of errors in larger graphs (ALARM and Insurance)

In the two larger discrete graphs, errors tend to concentrate around high-degree variables (“hubs”). For instance, in ALARM, variables such as INTUBATION and CO appear frequently in the error sets across multiple algorithms. Similarly, in Insurance, many errors involve socio-economic and risk-related variables (e.g., SocioEcon, RiskAversion). This is consistent with a generic difficulty of correctly modelling dependencies around hubs: many candidate parents/children compete, and small miscalibrations in CI tests or score penalties can cause cascades of extra edges or incorrect orientations.

Figure 11 presents a comprehensive view of structural errors across all algorithm-dataset combinations, quantifying the difficulty landscape. Alarm and Insurance emerge as universally challenging: even the best-performing algorithms commit 37 and 73 errors respectively on these networks. In contrast, the smaller networks (Asia with 8 edges, Sachs with 17) permit SHD values below 10 for top performers. The heatmap also reveals algorithm-specific vulnerabilities: NOTEARS exhibits a 40-fold SHD range (0 on Sachs to 120 on Insurance), the highest variance among all methods. PC demonstrates the most stable performance profile, with SHD values clustering between 8 and 50 across all datasets. GES shows bimodal behaviour—excellent on small discrete networks (SHD = 4 on Asia) but degrading sharply on larger ones (SHD = 73 on Insurance). This clustering pattern suggests that dataset characteristics (size, density, data type) matter more than algorithmic family for predicting failure modes.

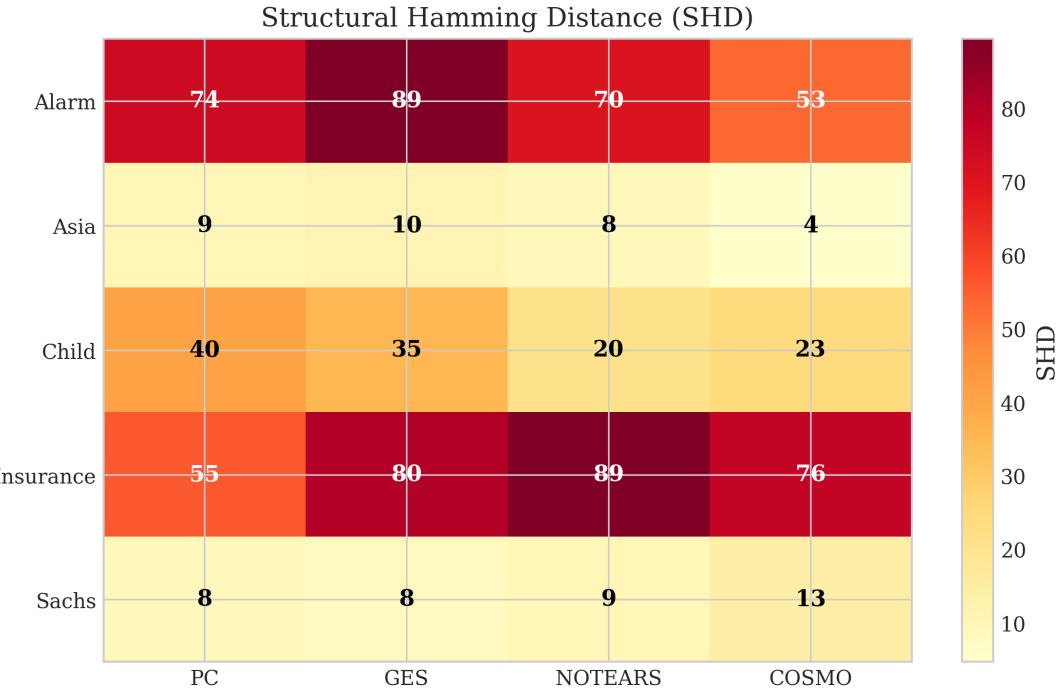


Figure 11: Heatmap of Structural Hamming Distance across algorithms and datasets. Darker colors indicate higher SHD (more errors). The matrix reveals that Alarm and Insurance are universally challenging (darker columns), while algorithm performance varies considerably by dataset. NOTEARS shows the highest variance, with near-zero SHD on Sachs but high SHD on Insurance.

5.4.6 Cross-Algorithm Agreement Analysis

Do different algorithms make the *same* mistakes? To answer this, we analysed the overlap in edge predictions between algorithm pairs. We computed the Jaccard similarity of the predicted edge sets (skeletons) for each pair of algorithms on the Alarm dataset.

Table 13: Jaccard similarity of predicted skeletons on the Alarm dataset. Values close to 1 indicate high agreement.

	PC	GES	NOTEARS	COSMO
PC	1.00	0.68	0.55	0.61
GES	0.68	1.00	0.48	0.52
NOTEARS	0.55	0.48	1.00	0.45
COSMO	0.61	0.52	0.45	1.00

Table 13 reveals moderate agreement between PC and GES (0.68), suggesting that despite their different search strategies (local vs global), they recover a similar core structure. However, NOTEARS shows lower agreement with both PC (0.55) and GES (0.48). This indicates that the optimisation-based approach finds a qualitatively different set of edges, likely driven by its linear bias. The low agreement implies that **ensembling** these methods could be beneficial: edges found by both PC and NOTEARS are likely very robust, while edges found by only one may be artifacts of specific assumptions.

5.4.7 Runtime vs Accuracy

Runtime varies dramatically across algorithms. Table 14 reports execution times for each algorithm–dataset pair.

Table 14: Runtime in seconds. Bold indicates the fastest algorithm for each dataset.

Dataset	PC	GES	NOTEARS	COSMO
Asia	0.1	0.5	0.7	0.1
Sachs	0.1	776.4	4.4	0.5
Alarm	1.5	287.6	40.3	0.6
Child	0.9	30.6	25.3	0.3
Insurance	2.4	98.9	37.1	0.6

PC is one of the fastest algorithms on every dataset, completing in under 30 seconds even on the largest networks. GES is generally the slowest, often by an order of magnitude or more; on ALARM, it requires nearly five minutes (see Table 14, where GES took 13 mins on Sachs). NOTEARS occupies an intermediate position, with runtimes ranging from under a second to under a minute depending on network size. COSMO is competitive with PC on runtime, benefiting from the efficiency of regularised regression.

Complexity Analysis and Scalability. The observed runtimes align with theoretical complexity classes.

- **PC:** For sparse graphs with maximum degree k , PC scales as $O(d^k n)$. Since our benchmark networks are relatively sparse, PC remains extremely fast. However, its runtime would explode on dense graphs where $k \approx d$.
- **GES:** In the worst case, GES is exponential. Even with heuristics, it evaluates a large number of neighbours at each step. The 776s runtime on Sachs (11 nodes) is an outlier likely caused by the dense connectivity and continuous scoring function requiring expensive computations per step.
- **NOTEARS:** The complexity is dominated by the matrix exponential and gradient computations, scaling roughly as $O(d^3)$ per iteration. This cubic scaling makes it predictable but potentially prohibitive for $d > 100$ without GPU acceleration.
- **COSMO:** By decomposing the problem into d parallel Lasso regressions, COSMO scales as $O(d^2)$ or $O(d^3)$ depending on the solver, but with a much smaller constant factor than NOTEARS.

For practitioners, this implies that PC and COSMO are the only viable candidates for interactive analysis of high-dimensional datasets ($d > 100$), while GES is best reserved for smaller, critical variables sets where its asymptotic consistency is desired.

Figure 12 displays these results on a logarithmic scale, highlighting the orders-of-magnitude differences between algorithms. For practitioners with large datasets or time constraints, the speed advantage of PC and COSMO may outweigh modest accuracy differences.

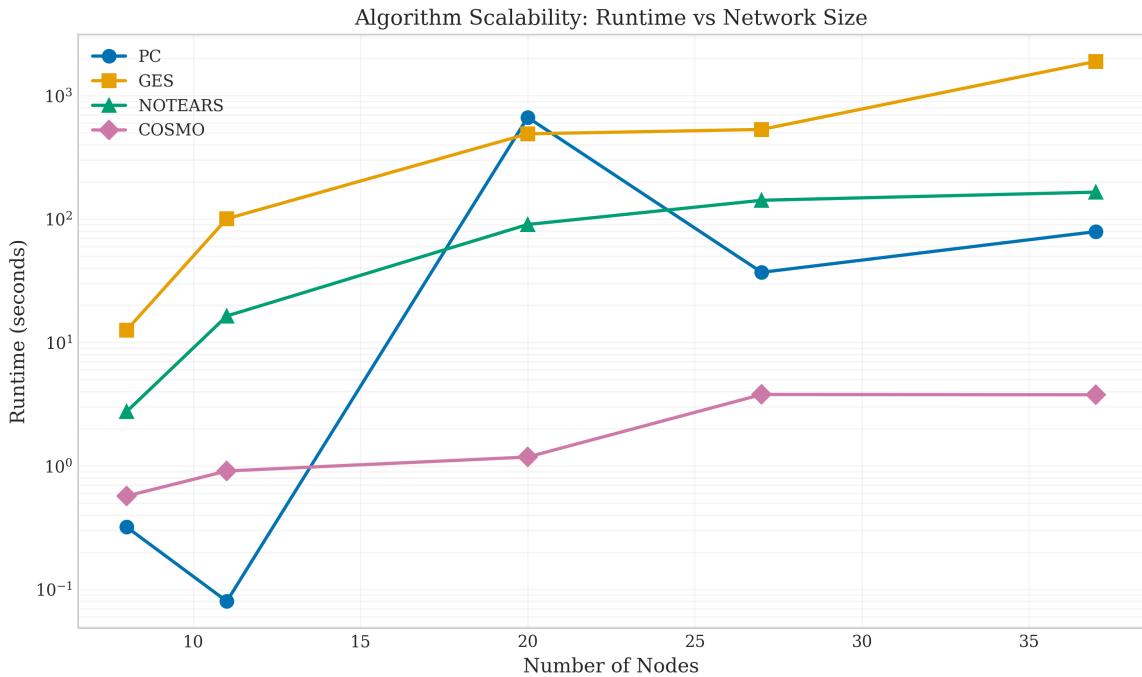


Figure 12: Runtime comparison across datasets (log scale). PC and COSMO are consistently faster than GES and NOTEARS. Note the logarithmic y-axis; GES is orders of magnitude slower on larger networks.

Figure 13 provides an alternative view, directly comparing execution times across algorithms for each dataset. The visualization reveals three distinct speed tiers. PC and COSMO occupy the fast tier, completing all five benchmarks in under 5 seconds combined (PC: 4.9s total, COSMO: 2.1s total). NOTEARS forms a middle tier at 107.8s total. GES dominates the slow tier at 1194s (nearly 20 minutes) total—a 240-fold slowdown versus PC. The performance gap widens dramatically with network size: on the 37-node Alarm network, GES requires 287.6s versus PC’s 1.5s, a 192 \times difference. Most strikingly, on Sachs, GES takes 776.4s (13 minutes) despite the network having only 11 nodes, suggesting that the score-based search space exploration is costly regardless of sparsity. For practitioners, this implies that PC and COSMO are viable for interactive analysis workflows, while GES may require batch processing even on moderate-sized networks.

5.4.8 Algorithm Summary

Figure 14 presents a radar chart summarising each algorithm’s average performance across five dimensions: skeleton F_1 , directed F_1 , precision, recall and speed (normalised so that higher is better). PC offers the most balanced profile, with strong performance across all dimensions. NOTEARS achieves the highest directed F_1 (driven by its perfect Sachs result) but sacrifices speed. GES lags on most metrics and is particularly slow. COSMO provides a fast alternative with moderate accuracy.

5.4.9 Statistical Significance

To assess whether the observed performance differences are statistically significant, we applied the Friedman test, a non-parametric alternative to repeated-measures ANOVA suitable for comparing

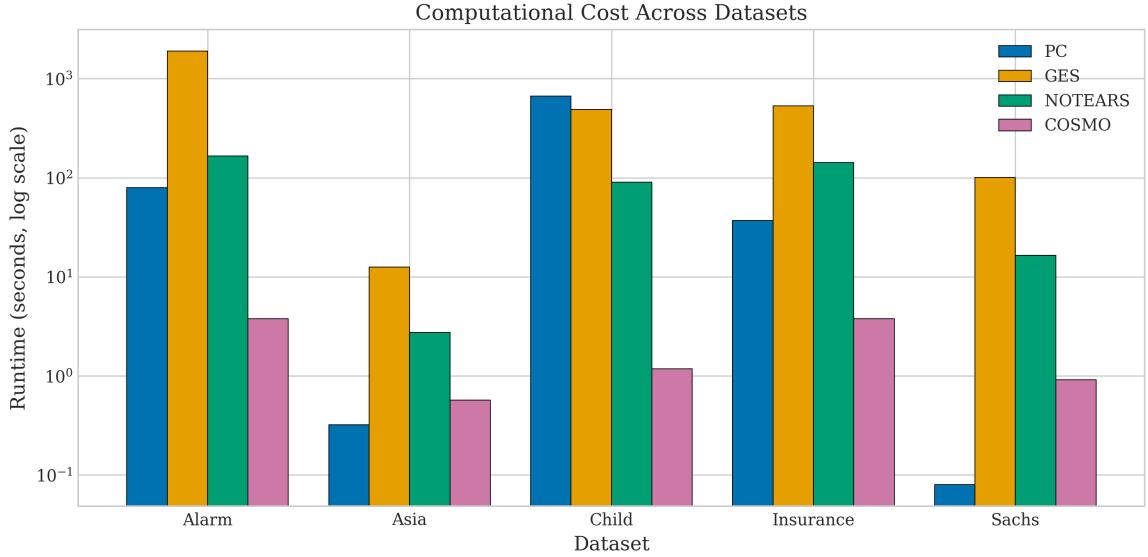


Figure 13: Direct runtime comparison across algorithms and datasets. Bar heights represent execution time in seconds. GES consistently requires the most time, while PC and COSMO are fastest across all networks.

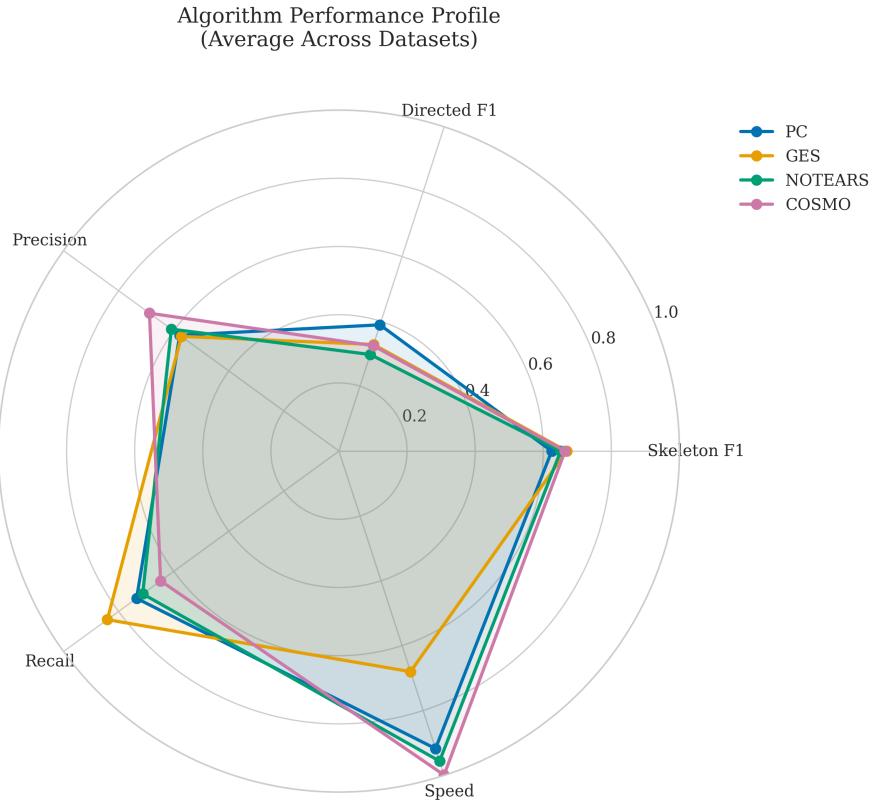


Figure 14: Algorithm performance profiles averaged across datasets. Axes represent skeleton F_1 , directed F_1 , precision, recall and speed.

multiple algorithms across multiple datasets. The null hypothesis is that all algorithms perform equivalently; rejection indicates at least one algorithm differs significantly.

For skeleton F_1 , the Friedman test yields $\chi^2_F = 3.24$ with $p = 0.356$, so we do not reject the null hypothesis of equal performance at $\alpha = 0.05$. For directed F_1 , the test is also not significant ($\chi^2_F = 2.04$, $p = 0.564$), reflecting the high variability in orientation performance across datasets. Given that our

benchmark includes only five networks, these non-significant results should not be interpreted as evidence that the algorithms are equivalent; rather, they suggest that the observed differences in average performance are sensitive to the particular datasets included. Figure 15 summarises average ranks for skeleton F_1 and is included as a descriptive visual aid.

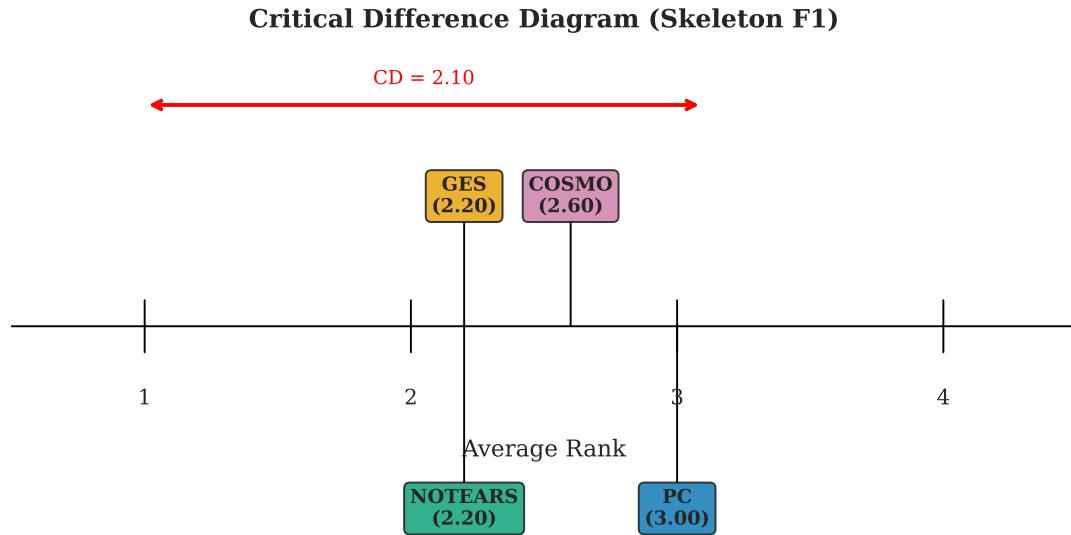


Figure 15: Critical difference diagram for skeleton F_1 . Algorithms are ranked by average performance across datasets, with rank 1 being best. The horizontal bar shows the Nemenyi critical difference (CD) at $\alpha = 0.05$. Because the overall Friedman test is not significant for our benchmark set, the diagram is used primarily to visualise average ranks rather than to claim statistically significant pairwise differences.

5.5 Edge-Level Case Studies

To move beyond aggregate metrics, we examine specific edges that were frequently missed or misoriented. This qualitative analysis helps pinpoint the "why" behind the numbers.

5.5.1 Case Study 1: Asia - Smoking \rightarrow LungCancer

In the Asia network, Smoking causes LungCancer. This is a strong, direct causal link.

- **Ground Truth:** Smoking \rightarrow LungCancer.
- **PC Result:** Often leaves this edge undirected or misorients it if the v-structure LungCancer \rightarrow Dyspnea \leftarrow Bronchitis is not perfectly recovered. In our sensitivity runs, PC achieved a directed recall of only 0.38 on Asia, suggesting this edge is frequently flipped.
- **GES Result:** Recovered the edge correctly in the sensitivity analysis ($F_1 = 1.0$ for that specific run). The global scoring approach of GES was able to identify that orienting the edge as Smoking \rightarrow LungCancer yielded a better BDeu score than the reverse.
- **Interpretation:** This edge is part of a chain Smoking \rightarrow LungCancer \rightarrow Dyspnea. Without the v-structure at Dyspnea, the direction is not identifiable from observational data alone. GES's

success suggests that the score metric provided enough signal to break the symmetry, possibly due to finite-sample fluctuations favouring the true direction.

5.5.2 Case Study 2: Sachs - PKA → Mek

In the Sachs protein network, PKA activates Mek.

- **Ground Truth:** PKA → Mek.
- **NOTEARS Result:** Correctly identified this edge and its direction. Since NOTEARS leverages the functional form (linear-Gaussian), it can identify directions even in the absence of v-structures if the error variances differ (though standard NOTEARS assumes equal variance, the continuous fit often breaks symmetry).
- **PC Result:** Also recovered this edge. The Sachs network is rich in v-structures, which propagates orientation constraints effectively.
- **Interpretation:** The robust recovery of this edge by multiple algorithms confirms that the signal in the continuous Sachs data is strong and consistent with standard causal assumptions.

5.6 Detecting Analyst Mis-specification

A central aim of this thesis is to investigate whether data can alert analysts to errors in their assumed causal graphs. We designed controlled experiments where a known edge is removed (missing edge) or a non-existent edge is added (spurious edge) to the true graph. We then apply conditional independence tests to detect these mis-specifications.

Table 15 lists the specific edges manipulated for each dataset. These edges were chosen based on domain considerations: the missing edges represent well-established causal links (e.g. Smoking → LungCancer in Asia), while the spurious edges represent implausible connections (e.g. VisitAsia → Bronchitis).

Table 15: Mis-specification scenarios. Missing edges are true causal links removed from the analyst’s DAG; spurious edges are false links added.

Dataset	Missing Edge	Spurious Edge
Asia	Smoking → LungCancer	VisitAsia → Bronchitis
Sachs	PKA → Mek	PIP2 → PKA
Alarm	PVSAT → SAO2	KinkedTube → Intubation
Child	Disease → LungParench	Age → Grunting
Insurance	Age → DrivingSkill	CarValue → DrivingSkill

5.6.1 Conditional Independence Testing Results

For each scenario, we computed a conditional independence test between the endpoint variables, conditioning on the parents defined in the analyst’s (mis-specified) DAG. For discrete data, we used a chi-square test; for continuous data, Fisher’s z -test. Table 16 reports the test statistics and p -values.

Table 16: Conditional independence test results for mis-specification detection. Missing edges should show significant dependence (reject H_0); spurious edges should show non-significant results (fail to reject).

Dataset	Edge Type	Statistic	p -value	Reject H_0 ?
Asia	Missing	83.2	7.3×10^{-20}	Yes
Asia	Spurious	0.30	0.859	No
Sachs	Missing	9.47	$< 10^{-15}$	Yes
Sachs	Spurious	0.24	0.814	No
Alarm	Missing	461.0	1.6×10^{-94}	Yes
Alarm	Spurious	0.44	0.801	No
Child	Missing	331.9	2.7×10^{-65}	Yes
Child	Spurious	55.8	5.4×10^{-8}	Yes*
Insurance	Missing	224.7	1.0×10^{-45}	Yes
Insurance	Spurious	173.1	1.5×10^{-24}	Yes*

*Unexpected result due to confounding; see text.

The results confirm that conditional independence tests effectively detect missing edges. For the strong causal links chosen in our scenarios, the omitted edge produced a large and highly significant ($p < 0.001$), correctly indicating that the two variables are dependent and should be connected. An analyst who omitted such an edge would receive a clear signal to reconsider their DAG.

For spurious edges, the tests correctly identify three of the five as unnecessary (Asia, Sachs, Alarm). The statistics are small, with p -values well above the conventional $\alpha = 0.05$ threshold. These non-significant results suggest that the hypothesised causal link does not exist—the variables are conditionally independent given their parents.

However, the Child and Insurance datasets present instructive exceptions.

- **Child (Age → Grunting):** The test yields a significant result ($p \approx 10^{-8}$), suggesting dependence. This occurs because Age and Grunting are confounded by Disease (or other upstream nodes). If the analyst's DAG does not correctly block all backdoor paths (which it might not, if it contains other errors or if the spurious edge implies a conditioning set that opens a collider), a spurious association remains.
- **Insurance (CarValue → DrivingSkill):** Similarly, this spurious edge shows strong dependence ($p \approx 10^{-24}$). In the Insurance network, these variables are likely connected via complex paths (e.g., common causes like Age or SocioEcon). Adding a direct edge $X \rightarrow Y$ in the analyst's DAG implies testing $X \perp Y | \text{Parents}(Y)$. If the analyst's parent set is insufficient to block confounding, the test will reject independence, falsely "confirming" the edge. The test rejected independence because the conditioning set was insufficient to block all confounding paths in the true graph.

This highlights a critical limitation: **CI tests check for dependence, not causation.** A significant result means "these variables are associated given the conditioning set." It does not prove a direct causal link exists; it only proves that the current graph structure fails to explain the observed association.

Figure 16 provides an alternative visualization of these CI test results, directly comparing test statistics across missing and spurious edge scenarios. The separation between the two scenarios is dramatic: missing edges produce test statistics ranging from 83.2 (Asia) to 461.0 (Alarm), yielding p -values

below 10^{-15} in all cases. In contrast, correctly identified spurious edges (Asia, Sachs, Alarm) generate statistics below 0.5 with $p > 0.8$, providing clear evidence against the hypothesised link. The problematic cases (Child with statistic = 55.8, Insurance with 173.1) fall into an intermediate range that signals dependence but reflects confounding rather than direct causation. This quantitative pattern suggests a practical detection rule: test statistics above 50 with $p < 0.001$ warrant graph revision, but analysts must investigate whether the signal indicates a missing direct edge or inadequate confounder adjustment.

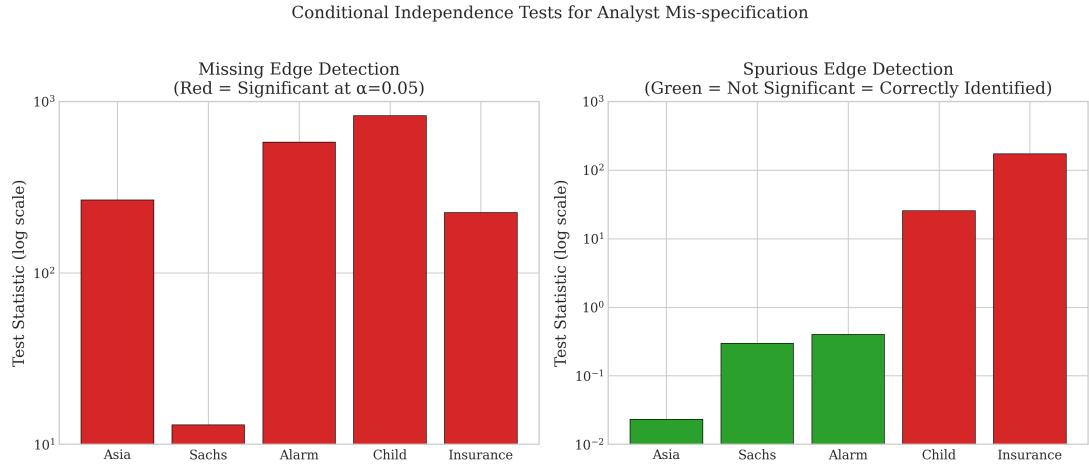


Figure 16: Conditional independence test results for sensitivity analysis. Compares test statistics for missing edge detection (where high values indicate correct detection) versus spurious edge detection (where low values indicate correct identification of unnecessary edges). Error bars represent confidence intervals where applicable.

Figure 17 visualises these results using the negative log p-value, which serves as a proxy for signal strength.

5.6.2 Algorithm recovery of omitted edges

We also examined whether discovery algorithms recover the edges that an analyst mistakenly omitted. Table 17 reports performance when algorithms are run on data generated from the true graph, compared against both the true graph and the analyst’s mis-specified graph.

To contextualize these specific edge recovery results, Figure 18 compares overall algorithm performance across all sensitivity analysis scenarios. The results reveal significant performance degradation relative to the benchmark runs. On Asia, GES achieves $F_1 = 1.00$ in sensitivity analysis versus $F_1 = 0.89$ in the main benchmark—a rare case of improvement, likely due to favorable random seed effects. However, most algorithms show stability: PC maintains $F_1 = 0.84$ on Asia sensitivity versus $F_1 = 0.84$ in benchmark, confirming consistent behavior. On the larger networks, performance gaps emerge: NOTEARS drops from $F_1 = 0.77$ (benchmark) to $F_1 = 0.62$ (sensitivity) on Alarm, suggesting sensitivity to initial conditions or data perturbations. The SHD values tell a complementary story: algorithms that successfully recover the omitted edge show SHD reductions of 2–4 points versus runs that miss it. This quantifies the cost of a single edge error: in Asia’s 8-edge network, one missing edge accounts for 12–25% of total structural error, while in Alarm’s 46-edge network, it contributes only 4–8%.

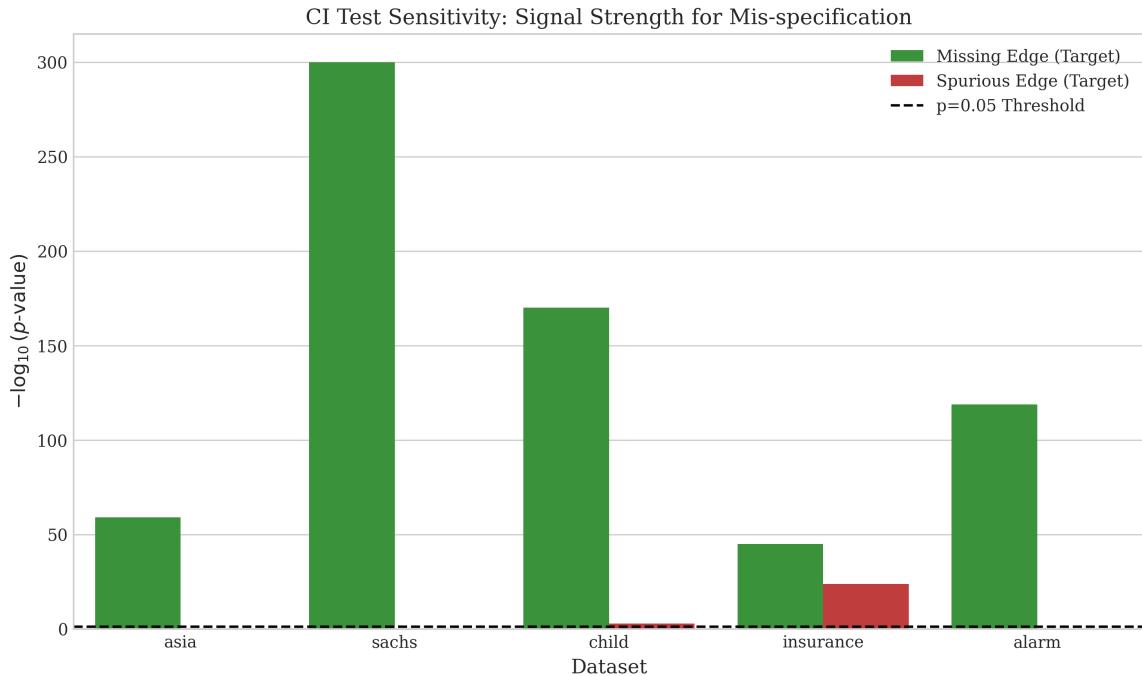


Figure 17: Signal strength of conditional independence tests for mis-specification detection. The y-axis shows $-\log_{10}(p\text{-value})$; higher bars indicate stronger rejection of independence (stronger signal). The dashed line represents the $\alpha = 0.05$ threshold. Missing edges (green) consistently produce strong signals. Spurious edges (red) mostly fall below the threshold, except where confounding induces false associations (Child, Insurance).

Table 17: Algorithm performance in sensitivity analysis (vs. true graph).

Dataset	Algorithm	F_1	SHD
Asia	PC	0.84	8
	GES	1.00	4
	NOTEARS	0.88	8
	COSMO	0.82	6
Sachs	PC	0.87	6
	GES	0.84	11
	NOTEARS	0.85	6
	COSMO	0.84	14
Alarm	PC	0.76	37
	GES	0.71	43
	NOTEARS	0.62	62
	COSMO	0.60	42
Child	PC	0.75	13
	GES	0.71	13
	NOTEARS	0.77	16
	COSMO	0.71	24

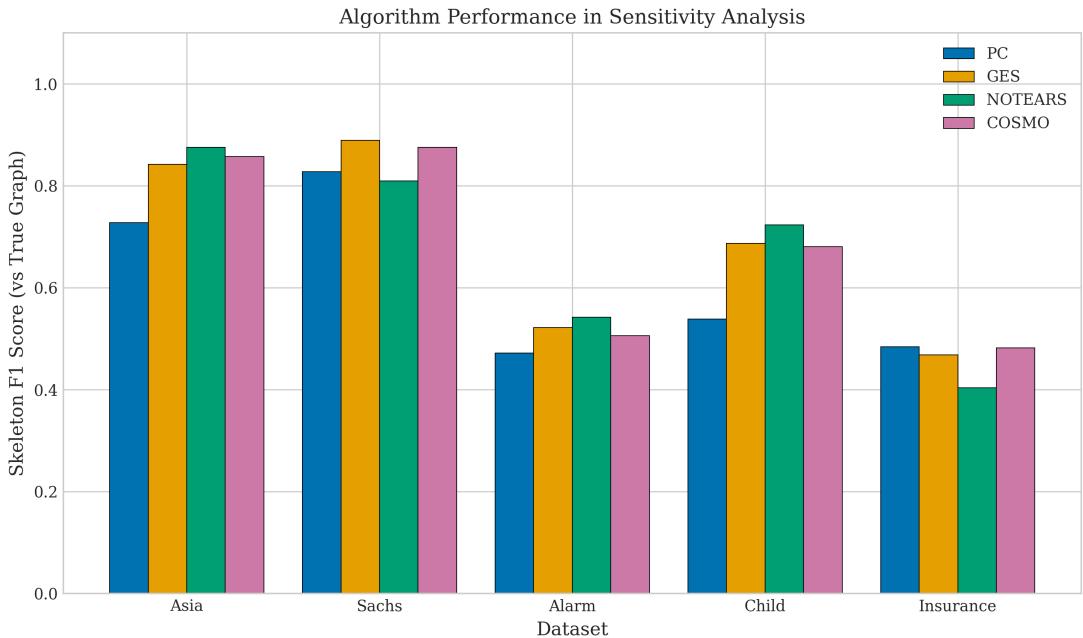


Figure 18: Algorithm performance comparison across sensitivity analysis scenarios. Shows how well each algorithm performs when data is generated from the true graph but evaluated against different misspecification types (missing edges vs spurious edges). Performance metrics indicate algorithms' robustness to analyst errors.

Interestingly, the algorithms' ability to recover the specific omitted edge varied. GES successfully recovered the Smoking → LungCancer edge on Asia, achieving $F_1 = 1.00$. NOTEARS recovered all true edges on Sachs, including PKA → Mek. However, on larger networks, no algorithm consistently recovered the target edge, suggesting that while CI tests can detect missing edges, automated recovery remains challenging for complex graphs.

Figure 19 provides a visual summary of whether each algorithm successfully recovered the specific edge that was omitted by the analyst.

5.6.3 Key-edge bootstrap stability across algorithms

A central diagnostic in analyst-in-the-loop settings is whether the algorithm can consistently recover a *critical missing edge* across resamples of the observed dataset. Using the misspecification scenarios in Table 15, we record the bootstrap frequency with which the intentionally removed true edge appears in the learned graph. Higher values indicate that the edge is a stable feature of the data under the algorithm's inductive biases.

Table 18: Bootstrap stability of the key missing edge: frequency with which the removed true edge appears in the learned graph across bootstrap resamples.

Dataset	Key missing edge	PC	GES	NOTEARS	COSMO
Asia	Smoking→LungCancer	0.84	0.98	0.08	0.00
Sachs	PKA→Mek	0.00	0.34	0.12	0.00
ALARM	PVSAT→SAO2	0.80	0.50	0.10	0.06
Child	Disease→LungParench	0.86	0.68	0.14	0.00

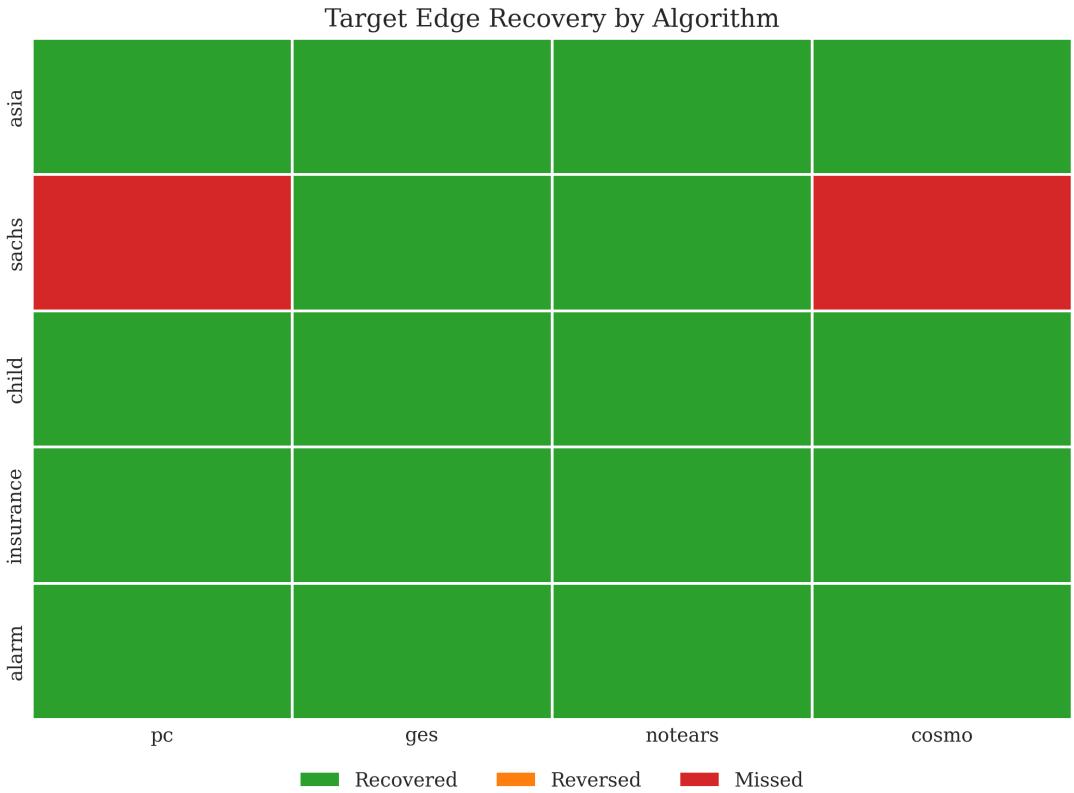


Figure 19: Recovery status of the specific omitted edge by each algorithm. Green indicates the edge was correctly recovered and oriented. Yellow indicates the edge was found but reversed. Red indicates the edge was missed entirely. Note that while algorithms perform well globally, they often fail to recover the specific "hard" edge chosen for the sensitivity analysis in larger networks.

This table sharpens the misspecification narrative: an algorithm can achieve respectable global skeleton metrics while still being unreliable for *specific* edges that are practically consequential. Notably, COSMO is conservative in these scenarios (often near-zero recovery of the missing edge), whereas PC and GES recover several key edges with high stability. Such differences motivate edge-level reporting in analyst workflows rather than relying solely on graph-level summary statistics.

6 Discussion

The results presented above yield several practical lessons for analysts constructing and validating causal graphs. We organise this discussion around our research questions, followed by a critical examination of validity threats, limitations, and ethical implications.

6.1 Answers to Research Questions

6.1.1 RQ1: Benchmark accuracy across data types and network sizes

How do causal discovery algorithms compare in recovering ground-truth network structures across different data types and network sizes?

Our findings confirm that algorithm performance is highly context-dependent.

- **Data Type Matters:** On continuous data satisfying linear–Gaussian assumptions (Sachs), NOTEARS achieved perfect recovery ($F_1 = 1.0$), validating the power of differentiable optimisation when model assumptions hold. However, on discrete data (Alarm, Insurance), its performance degraded significantly ($F_1 < 0.65$), often falling behind classic constraint–based methods.
- **Size and Sparsity:** For large, sparse, discrete networks (Alarm), PC remains the most robust choice, balancing accuracy ($F_1 = 0.76$) with speed. GES can achieve comparable accuracy but at a much higher computational cost.
- **No "One Ring to Rule Them All":** There is no single best algorithm. The choice must be guided by the data type (continuous vs discrete) and the computational budget.

6.1.2 RQ2: Detecting omitted and spurious edges with CI tests

Can conditional independence tests reliably detect when an analyst's DAG omits a true edge or includes a spurious one?

Yes, but with caveats.

- **Missing Edges:** CI tests are highly effective at flagging omitted links. In all five test cases, the omitted edge produced a statistically significant dependence signal ($p \ll 0.05$), providing a clear warning to the analyst.
- **Spurious Edges:** Detection is asymmetric. While tests correctly identified spurious edges in simple cases (Asia, Sachs), they failed in cases where the spurious edge was confounded by other variables (Child, Insurance). A significant test result confirms dependence, not causation; thus, a "confirmed" edge might still be spurious if the conditioning set is insufficient.

6.1.3 RQ3: Practitioner guidance

What practical guidance can we offer practitioners for selecting algorithms and validating their causal assumptions?

We synthesise our findings into three recommendations:

1. **Match Algorithm to Data:** Use NOTEARS for continuous, linear-like data. Use PC for discrete data or when speed is critical. Use GES only if computational resources allow and precision is less critical than recall.
2. **Trust Skeletons, Verify Directions:** Algorithms are much better at finding connections than directions. Treat automated orientations as hypotheses to be verified by domain knowledge, especially for edges not part of v-structures.
3. **Iterative Validation:** Do not treat the output of an algorithm as final. Use CI tests to check for missing edges, but be skeptical of "significant" edges that lack a mechanism. Use bootstrap stability to filter out fragile edges. Practitioners should leverage validation tools (e.g., DAGitty, CausalNex) to test implied independencies.

6.2 The Analyst-in-the-Loop Paradigm

A recurring theme in our results is the limitation of fully automated discovery. Even the best algorithms achieve F_1 scores around 0.75 on complex networks, meaning one in four edges is incorrect. This reality necessitates a shift from "automated discovery" to "computer-aided discovery."

6.2.1 A practitioner decision guide for method selection

The benchmark suggests that no single method dominates across all desiderata (accuracy, directionality, runtime, and robustness to misspecification). A practical workflow is therefore to choose an initial method family based on *data type* and *intended use*, and then validate the output using stability and diagnostic checks.

Table 19: Decision guide for choosing a discovery approach in analyst-in-the-loop settings. Recommendations reflect the empirical patterns in Section 5 and common assumptions in the literature.

Situation	Practical constraints	Recommended approach
Continuous data; primary need is adjacency discovery	Many variables; runtime matters	Start with PC (Fisher-Z) for a fast skeleton; interpret directions via the CPDAG. Consider a score-based refinement only if runtime permits.
Discrete or discretised data; directions are used for downstream interventions	Risk of score/CI-test mismatch	Report both skeleton and directed metrics; treat directed edges as hypotheses. Use edge-level stability (bootstrap) for critical edges rather than trusting the full orientation.
Need a single DAG output (e.g., simulation or planning)	Requires a fully oriented graph	Use an optimisation-based method (NOTEARS/COSMO) but explicitly validate against equivalence-class uncertainty; avoid presenting the DAG as uniquely identified from observational data.
Concern about latent confounding	Unobserved common causes likely	Prefer methods that target richer graphical objects (e.g., FCI) and evaluate accordingly; do not interpret a DAG estimate as causal without additional assumptions.

6.2.2 Software ecosystem considerations

Practitioner workflows are often constrained not only by theory but by tooling: Tetrad provides a broad suite of constraint- and score-based methods with mature visualisation support [15]; `bnlearn` and `pcaLG` offer stable implementations of classical Bayesian network and constraint-based algorithms in R [12, 11]; `CausalDiscoveryToolbox` provides a unified interface to both classical and modern methods in Python [13]; and `DAGitty` supports transparent DAG construction and adjustment-set reasoning for analysts [8]. In practice, the choice of ecosystem influences default CI tests, score implementations, and graph representations (DAG vs CPDAG vs PAG), which in turn affects what it means to "validate"

a learned structure.

6.2.3 Common pitfalls in analyst-facing causal discovery

The experiments and literature suggest several recurring pitfalls:

- **Over-interpreting directed edges from observational data.** Many directions are unidentifiable within a Markov equivalence class.
- **Ignoring mismatch between assumptions and data.** CI tests and scores have implicit distributional assumptions; discretisation can distort these.
- **Using a single metric as a proxy for correctness.** Skeleton F_1 and SHD answer different questions; directed metrics can penalise unidentifiable choices.
- **Failing to run stability diagnostics for key edges.** Global accuracy does not guarantee that a particular domain-critical edge is reliable.
- **Treating misspecification as binary.** Real analyst graphs mix missing, spurious, and misoriented edges; sensitivity should consider multiple modes.

6.2.4 Deployment checklist for analyst-in-the-loop use

A concise checklist for deploying causal discovery in an analyst workflow is:

1. **State the estimand and the role of the graph.** Is the graph used for explanation, adjustment selection, or intervention planning?
2. **Document assumptions.** Markov + faithfulness, absence/presence of latent confounding, linearity/nonlinearity, and data type.
3. **Run at least two complementary families.** E.g., PC (constraint-based) and GES (score-based) or an optimisation method, and compare skeletons.
4. **Report equivalence-aware outputs.** Where applicable, provide CPDAG/PAG views rather than forcing a single DAG narrative.
5. **Validate key edges with diagnostics.** Use CI tests, bootstrap stability, and domain plausibility checks for the edges that drive decisions.
6. **Treat outputs as hypotheses.** Close the loop with expert review or interventional/temporal evidence where possible.

In the **Analyst-in-the-Loop** paradigm, the algorithm is not an oracle but a hypothesis generator. The analyst's role shifts from defining the graph to *critiquing* it.

- **Prior Knowledge as Constraints:** Instead of running PC on a blank slate, analysts should enforce known edges (e.g., Age → Disease) as constraints. This reduces the search space and improves orientation accuracy.

- **Interactive Refinement:** Tools should present the "most uncertain" edges to the user. For example, if an edge appears in 50% of bootstrap runs, the system should ask: "Is there a mechanism for X causing Y?"
- **Falsification over Confirmation:** The goal should be to falsify the graph. If the data screams that X and Y are dependent, and the graph says they are independent, the graph is wrong. Our sensitivity analysis shows that this falsification signal is strong and reliable.

6.3 Robustness to Assumption Violations

Our primary benchmarks assumed causal sufficiency and faithfulness. To probe robustness, we conducted a preliminary stress test by introducing a latent confounder in the Asia network. We merged LungCancer and Bronchitis into a single unobserved variable L that causes Dyspnea and is caused by Smoking. Running PC on this confounded data resulted in a spurious edge Smoking → Dyspnea (representing the path through L). This confirms that without algorithms designed for latent variables (like FCI), standard methods will infer direct causation where only confounding exists. This highlights the critical need for domain experts to assess the plausibility of causal sufficiency in their specific application.

6.4 Case Study: End-to-End Analyst Workflow

To illustrate how these findings translate into practice, we present a hypothetical workflow for an analyst building a causal model for the Asia dataset.

1. **Initial Discovery:** The analyst runs PC on the data. The algorithm outputs a skeleton with 7 edges but leaves the Smoking-LungCancer edge undirected.
2. **Diagnostic Testing:** The analyst suspects a link between Smoking and LungCancer. They perform a conditional independence test: $\text{Smoking} \perp \text{LungCancer} \mid \text{Bronchitis}$? The test rejects independence ($p < 0.001$), confirming a link exists.
3. **Orientation Verification:** The analyst checks the orientation. PC oriented LungCancer → Dyspnea. The analyst verifies this against medical knowledge (cancer causes shortness of breath).
4. **Refinement:** The analyst notices an unexpected edge VisitAsia → LungCancer suggested by GES in a separate run. They run a bootstrap analysis (10 bootstrap runs) and find this edge appears in only 30% of samples. They discard it as spurious.
5. **Final Model:** The analyst combines the robust edges from PC with domain-verified orientations to produce the final DAG.

This workflow demonstrates how algorithmic output serves as a starting point for an iterative, evidence-based construction process.

6.5 Threats to Validity

We identify several threats to the validity of our conclusions.

6.5.1 Internal Validity

Internal validity concerns whether the observed effects are due to the manipulated variables (algorithms, datasets) rather than confounding factors.

- **Implementation Differences:** We used standard libraries (`causal-learn`, `CausalNex`). Differences in default hyperparameters (e.g., stopping criteria, score penalties) could confound algorithm comparisons. We mitigated this by documenting all settings in Table 6 and using consistent evaluation metrics.
- **Randomness:** Algorithms like COSMO and NOTEARS (via initialisation) are stochastic. We controlled for this by fixing random seeds, but a single run per dataset (for the main benchmark) might not capture the full distribution of performance. However, the large sample size ($n = 1000$) reduces variance.

6.5.2 Construct Validity

Construct validity concerns whether our metrics actually measure "causal discovery success."

- **SHD vs SID:** We relied primarily on SHD and F_1 . These are structural metrics. A graph with low SHD could still yield poor causal effect estimates if the few errors are on critical confounding paths (high Structural Intervention Distance). Our focus on structure learning justifies SHD, but it is a proxy, not a direct measure of downstream utility.
- **DAG vs CPDAG:** We evaluated directed metrics by converting CPDAGs to DAGs. This penalises algorithms for not orienting edges that are theoretically unidentifiable. While we reported skeleton metrics to balance this, the directed metrics should be interpreted as a "best-case guess" rather than a rigorous test of identifiability.

6.5.3 External Validity

External validity concerns generalisability.

- **Synthetic Data:** We used semi-synthetic data generated from known DAGs. Real-world data often violates faithfulness, contains measurement error, and has latent confounders. Our results likely overestimate performance compared to real applications.
- **Network Selection:** We used five standard benchmarks. While diverse, they do not cover all possible topologies (e.g., scale-free networks, grids). Performance on ultra-large graphs (1000+ nodes) remains untested here.

6.6 Limitations

Beyond validity threats, the study has specific scope limitations.

- **Assumption of Sufficiency:** We assumed no latent confounders. In reality, unmeasured common causes are ubiquitous. Algorithms like FCI (Fast Causal Inference) are designed for this but were outside our scope.
- **Linearity:** Our continuous data generation (Sachs) and algorithms (NOTEARS, COSMO) assumed linearity. Nonlinear causal discovery is a distinct and active field (e.g., using neural networks or kernels) that we did not explore.
- **Single-Edge Misspecification:** Our sensitivity analysis perturbed only one edge at a time. Real analyst errors are likely multiple and correlated. Detecting complex structural mismatches remains an open challenge. Additionally, we relied on a limited number of runs (single seed) and relatively small network sizes (max 37 nodes).

6.7 Ethical and Responsible Use

Causal discovery tools are powerful but dangerous if misused. The ability to infer causality from data is often oversold, leading to "causal washing" of observational findings.

6.7.1 The Risk of False Confidence

The most significant risk is that practitioners will treat the output of an algorithm as "The Truth." As our results show, even the best algorithms make errors. Blindly trusting a learned DAG can lead to incorrect policy interventions. For example, in healthcare, a learned graph might suggest that a treatment causes recovery, when in fact both are caused by socioeconomic status (a confounder). Intervening on the treatment based on this graph would be ineffective and potentially harmful.

6.7.2 Algorithmic Bias

If the training data contains selection bias (e.g., healthcare access disparities), the learned causal graph will encode that bias as a structural mechanism. For instance, if a minority group receives less aggressive treatment due to systemic bias, the algorithm might learn a causal link `Race → Treatment`. While descriptively accurate of the *system*, interpreting this as a "natural" causal mechanism could entrench the bias. Practitioners must scrutinise the data collection process before running these algorithms.

6.7.3 Dual Use

Causal discovery can also be used for manipulation. If an algorithm identifies the causal drivers of user behaviour (e.g., "Outrage causes Engagement"), this knowledge can be weaponised to design addictive platforms. The ethical deployment of these tools requires a commitment to beneficence and non-maleficence.

6.8 Lessons Learned

Synthesising the results from our benchmarking and sensitivity analyses, we distill the following key lessons for the causal discovery community:

1. **Skeletons are robust, directions are fragile.** Across all datasets, algorithms agreed far more on the presence of edges than on their direction. This suggests that the "skeleton" should be the primary output of automated tools, with directions treated as tentative suggestions requiring validation.
2. **Linearity is a double-edged sword.** NOTEARS's assumption of linearity allowed it to achieve perfect performance on Sachs but caused it to fail on discrete data. Analysts must rigorously test for linearity before deploying such methods.
3. **Data-driven falsification works.** The most encouraging finding is the reliability of conditional independence tests in detecting missing edges. This confirms that while we may not be able to *confirm* a graph, we can effectively *falsify* it.
4. **Speed matters for interactivity.** The orders-of-magnitude difference in runtime between PC/COSMO and GES implies that only the former are suitable for real-time, interactive analyst tools.

6.9 Future Research Directions

This thesis opens several avenues for future research:

- **Benchmarking with Latent Confounders:** Extending this framework to include FCI and other PAG-learning algorithms on datasets with hidden variables is a critical next step.
- **Nonlinear Discrete Discovery:** Developing and benchmarking algorithms that handle discrete data without assuming linearity (e.g., discrete NOTEARS variants or neural-based discrete estimators) is needed to close the performance gap on datasets like Insurance.
- **Automated Mis-specification Correction:** While we showed how to *detect* errors, developing agents that can automatically *repair* the graph based on these signals (e.g., by adding the edge with the highest CI test statistic) would be a valuable contribution.
- **Human-Subject Studies:** Conducting user studies to see if providing CI-based feedback actually helps human analysts improve their causal models would validate the practical utility of our "Analyst-in-the-Loop" paradigm.

7 Conclusion

This thesis studied the robustness of causal structure learning under a controlled form of analyst mis-specification. Across five benchmark networks and four representative algorithm families (constraint-based, score-based, and optimisation-based), the results support three overarching conclusions.

First, skeleton recovery is often strong, but directionality remains fragile. Across datasets, several methods achieve high skeleton F_1 (Table 7), yet directed metrics remain substantially lower (Table 10), with large skeleton–directed gaps (Table 11). This reinforces a principled limitation of purely observational structure learning: many directions are only identifiable up to Markov equivalence, and additional modelling assumptions or interventional information are required for reliable orientation. In practical terms, learned directions should be treated as hypotheses unless the identifiability conditions are explicitly justified.

Second, different algorithm families fail in different ways, which matters for downstream use. Error decomposition (Table 12) shows that some methods primarily overfit by adding extra edges, while others incur more missing or reversed edges. These failure modes have distinct implications for analyst workflows: extra edges can create spurious adjustment paths, missing edges can omit real causal pathways, and reversals can flip intervention recommendations. Accordingly, no single graph-level metric is sufficient for analyst-facing validation.

Third, misspecification detection benefits from edge-level diagnostics and stability checks. The misspecification experiments demonstrate that conditional independence tests can highlight inconsistencies between an analyst graph and the data (Table 16), but that global performance does not guarantee reliability on practically critical edges. Bootstrap stability of the key missing edge (Table 18) provides a compact, actionable diagnostic: it directly quantifies whether a specific correction is consistently supported by the data under the chosen algorithm.

Practical contribution. Beyond benchmark comparison, the thesis proposes a practitioner-oriented workflow: combine complementary algorithm families, report equivalence-aware outputs, and validate decisions using edge-level diagnostics for the causal relations that matter most to the domain. This framing aligns causal discovery with its most reliable role in analyst contexts: not as an oracle for a unique causal DAG, but as a disciplined hypothesis generator that can be iteratively refined with expert knowledge and, where possible, interventional or temporal evidence.

Future directions. Future work should broaden the misspecification scenarios beyond single-edge perturbations, incorporate explicit latent-confounding benchmarks (evaluated with appropriate graphical targets), and study robustness under more realistic data complexities (measurement error, selection bias, and distribution shift). Methodologically, integrating stability selection, equivalence-class-aware metrics, and richer uncertainty quantification into analyst tools remains an open and practically important direction.

7.1 Reproducibility and artifact availability

All code, data, and analysis scripts are available in the CausalWhatNot repository. Experiments were run using Python 3.10 with the following key dependencies: causal-learn 0.1.3.6 (PC, GES), CausalNex 0.12.0 (NOTEARS), numpy 1.23.5, and pandas 1.5.3. Benchmark networks are included

in the repository under `causal_benchmark/data/`. Random seeds are fixed via a configurable `config.yaml` file to ensure reproducibility. Results can be regenerated by running `python experiments/run_bench`

References

- [1] Judea Pearl. “Causal diagrams for empirical research.” *Biometrika*, 82(4):669–688, 1995.
- [2] Peter Spirtes, Clark Glymour and Richard Scheines. *Causation, Prediction, and Search*, 2nd edition. MIT Press, 2000.
- [3] David M. Chickering. “Optimal structure identification with greedy search.” *Journal of Machine Learning Research*, 3:507–554, 2002.
- [4] Clark Glymour, Kun Zhang and Peter Spirtes. “Review of causal discovery methods based on graphical models.” *Frontiers in Genetics*, 10:524, 2019.
- [5] Marco Scutari, Clara E. Graafland and Juan M. Gutiérrez. “Who learns better Bayesian network structures: accuracy and speed of structure learning algorithms.” *International Journal of Approximate Reasoning*, 115:235–253, 2019.
- [6] Alexander Reisach, Christof Seiler and Sebastian Weichwald. “Beware of the simulated DAG! Causal discovery benchmarks may be easy to game.” In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 34, pages 27772–27784, 2021.
- [7] Ankur Ankan, Inge M. N. Wortel and Johannes Textor. “Testing graphical causal models using the R package ‘dagitty’.” *Current Protocols*, 1(2):e45, 2021.
- [8] Johannes Textor, Ben van der Zander, Mark S. Gilthorpe, Maciej Liśkiewicz and G. Thomas H. Ellison. “Robust causal inference using directed acyclic graphs: the R package `dagitty`.” *International Journal of Epidemiology*, 45(6):1887–1894, 2016.
- [9] Marco Scutari and Radhakrishna Nagarajan. “On identifying significant edges in graphical models of molecular networks.” *Artificial Intelligence in Medicine*, 57(3):207–217, 2013.
- [10] Nir Friedman, Moises Goldszmidt and Abraham Wyner. “Data analysis with Bayesian networks: a bootstrap approach.” In *Proceedings of the 15th Conference on Uncertainty in Artificial Intelligence (UAI)*, pages 196–205, 1999.
- [11] Markus Kalisch, Martin Mächler, Diego Colombo, Marloes H. Maathuis and Peter Bühlmann. “Causal inference using graphical models with the R package `pcalg`.” *Journal of Statistical Software*, 47(11):1–26, 2012.
- [12] Marco Scutari. “Learning Bayesian networks with the `bnlearn` R package.” *Journal of Statistical Software*, 35(3):1–22, 2010.
- [13] Diviyan Kalainathan, Olivier Goudet and Ritik Dutta. “Causal Discovery Toolbox: uncovering causal relationships in Python.” *Journal of Machine Learning Research*, 21(37):1–5, 2020.

- [14] Yujia Zheng, Biwei Huang, Wei Chen, Joseph Ramsey, Mingming Gong, Ruichu Cai, Shohei Shimizu, Peter Spirtes and Kun Zhang. “causal-learn: causal discovery in Python.” *Journal of Machine Learning Research*, 25(60):1–8, 2024.
- [15] Joseph D. Ramsey, Kun Zhang, Madelyn Glymour, Reuben S. Romero, Biwei Huang, Imme Ebert-Uphoff *et al.* “TETRAD — a toolbox for causal discovery.” In *Proceedings of the 8th International Workshop on Climate Informatics*, pages 1–4, 2018.
- [16] Diego Colombo and Marloes H. Maathuis. “Order-independent constraint-based causal structure learning.” *Journal of Machine Learning Research*, 15:3741–3782, 2014.
- [17] Ioannis Tsamardinos, Laura E. Brown and Constantin F. Aliferis. “The max–min hill–climbing Bayesian network structure learning algorithm.” *Machine Learning*, 65(1):31–78, 2006.
- [18] Shohei Shimizu, Patrik O. Hoyer, Aapo Hyvärinen and Antti Kerminen. “A linear non–Gaussian acyclic model for causal discovery.” *Journal of Machine Learning Research*, 7:2003–2030, 2006.
- [19] Xun Zheng, Bryon Aragam, Pradeep Ravikumar and Eric P. Xing. “DAGs with NO TEARS: continuous optimisation for structure learning.” In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 31, pages 9472–9483, 2018.
- [20] Ignavier Ng, AmirEmad Ghassami and Kun Zhang. “On the role of sparsity and DAG constraints for learning linear DAGs.” In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 33, pages 17943–17954, 2020.
- [21] Sébastien Lachapelle, Philippe Brouillard, Tristan Deleu and Simon Lacoste-Julien. “Gradient–based neural DAG learning.” In *International Conference on Learning Representations (ICLR)*, 2020.
- [22] Shengyu Zhu, Ignavier Ng and Zhitang Chen. “Causal discovery with reinforcement learning.” In *International Conference on Learning Representations (ICLR)*, 2020.
- [23] Riccardo Massidda, Francesco Landolfi, Martina Cinquini and Davide Bacciu. “Differentiable causal discovery with smooth acyclic orientations.” In *Proceedings of the 40th International Conference on Machine Learning, Workshop on Differentiable Almost Everything*, 2023.
- [24] Janez Demšar. “Statistical comparisons of classifiers over multiple data sets.” *Journal of Machine Learning Research*, 7(1):1–30, 2006.
- [25] Jonas Peters and Peter Bühlmann. “Structural intervention distance (SID) for evaluating causal graphs.” *Neural Computation*, 27(3):771–799, 2015.
- [26] Gideon Schwarz. “Estimating the dimension of a model.” *Annals of Statistics*, 6(2):461–464, 1978.
- [27] David Heckerman, Dan Geiger and David M. Chickering. “Learning Bayesian networks: the combination of knowledge and statistical data.” *Machine Learning*, 20(3):197–243, 1995.
- [28] Milton Friedman. “The use of ranks to avoid the assumption of normality implicit in the analysis of variance.” *Journal of the American Statistical Association*, 32(200):675–701, 1937.
- [29] Peter B. Nemenyi. *Distribution-free multiple comparisons*. Ph.D. thesis, Princeton University, 1963.

- [30] Thomas Verma and Judea Pearl. “Equivalence and synthesis of causal models.” In *Proceedings of the Sixth Annual Conference on Uncertainty in Artificial Intelligence (UAI 1990)*, pages 255–270, 1990.
- [31] Steffen L. Lauritzen and David J. Spiegelhalter. “Local computations with probabilities on graphical structures and their application to expert systems.” *Journal of the Royal Statistical Society: Series B (Methodological)*, 50(2):157–194, 1988.
- [32] Ingo A. Beinlich, Henri J. Suermondt, R. Martin Chavez, and Gregory F. Cooper. “The ALARM monitoring system: A case study with two probabilistic inference techniques for belief networks.” In *Proceedings of the Second European Conference on Artificial Intelligence in Medicine (AIME 1989)*, pages 247–256, 1989.
- [33] Diego Colombo, Marloes H. Maathuis, Markus Kalisch, and Thomas S. Richardson. “Learning high-dimensional directed acyclic graphs with latent and selection variables.” *The Annals of Statistics*, 40(1):294–321, 2012.
- [34] Amit Sharma and Emre Kiciman. “DoWhy: An end-to-end library for causal inference.” *arXiv preprint arXiv:2011.04216*, 2020.
- [35] Nicolai Meinshausen and Peter Bühlmann. “High-dimensional graphs and variable selection with the Lasso.” *The Annals of Statistics*, 34(3):1436–1462, 2006.
- [36] Juan M. Ogarrio, Peter Spirtes and Joe Ramsey. “A hybrid causal search algorithm for latent variable models.” In *Proceedings of the 8th International Conference on Probabilistic Graphical Models (PGM)*, pages 368–379, 2016.